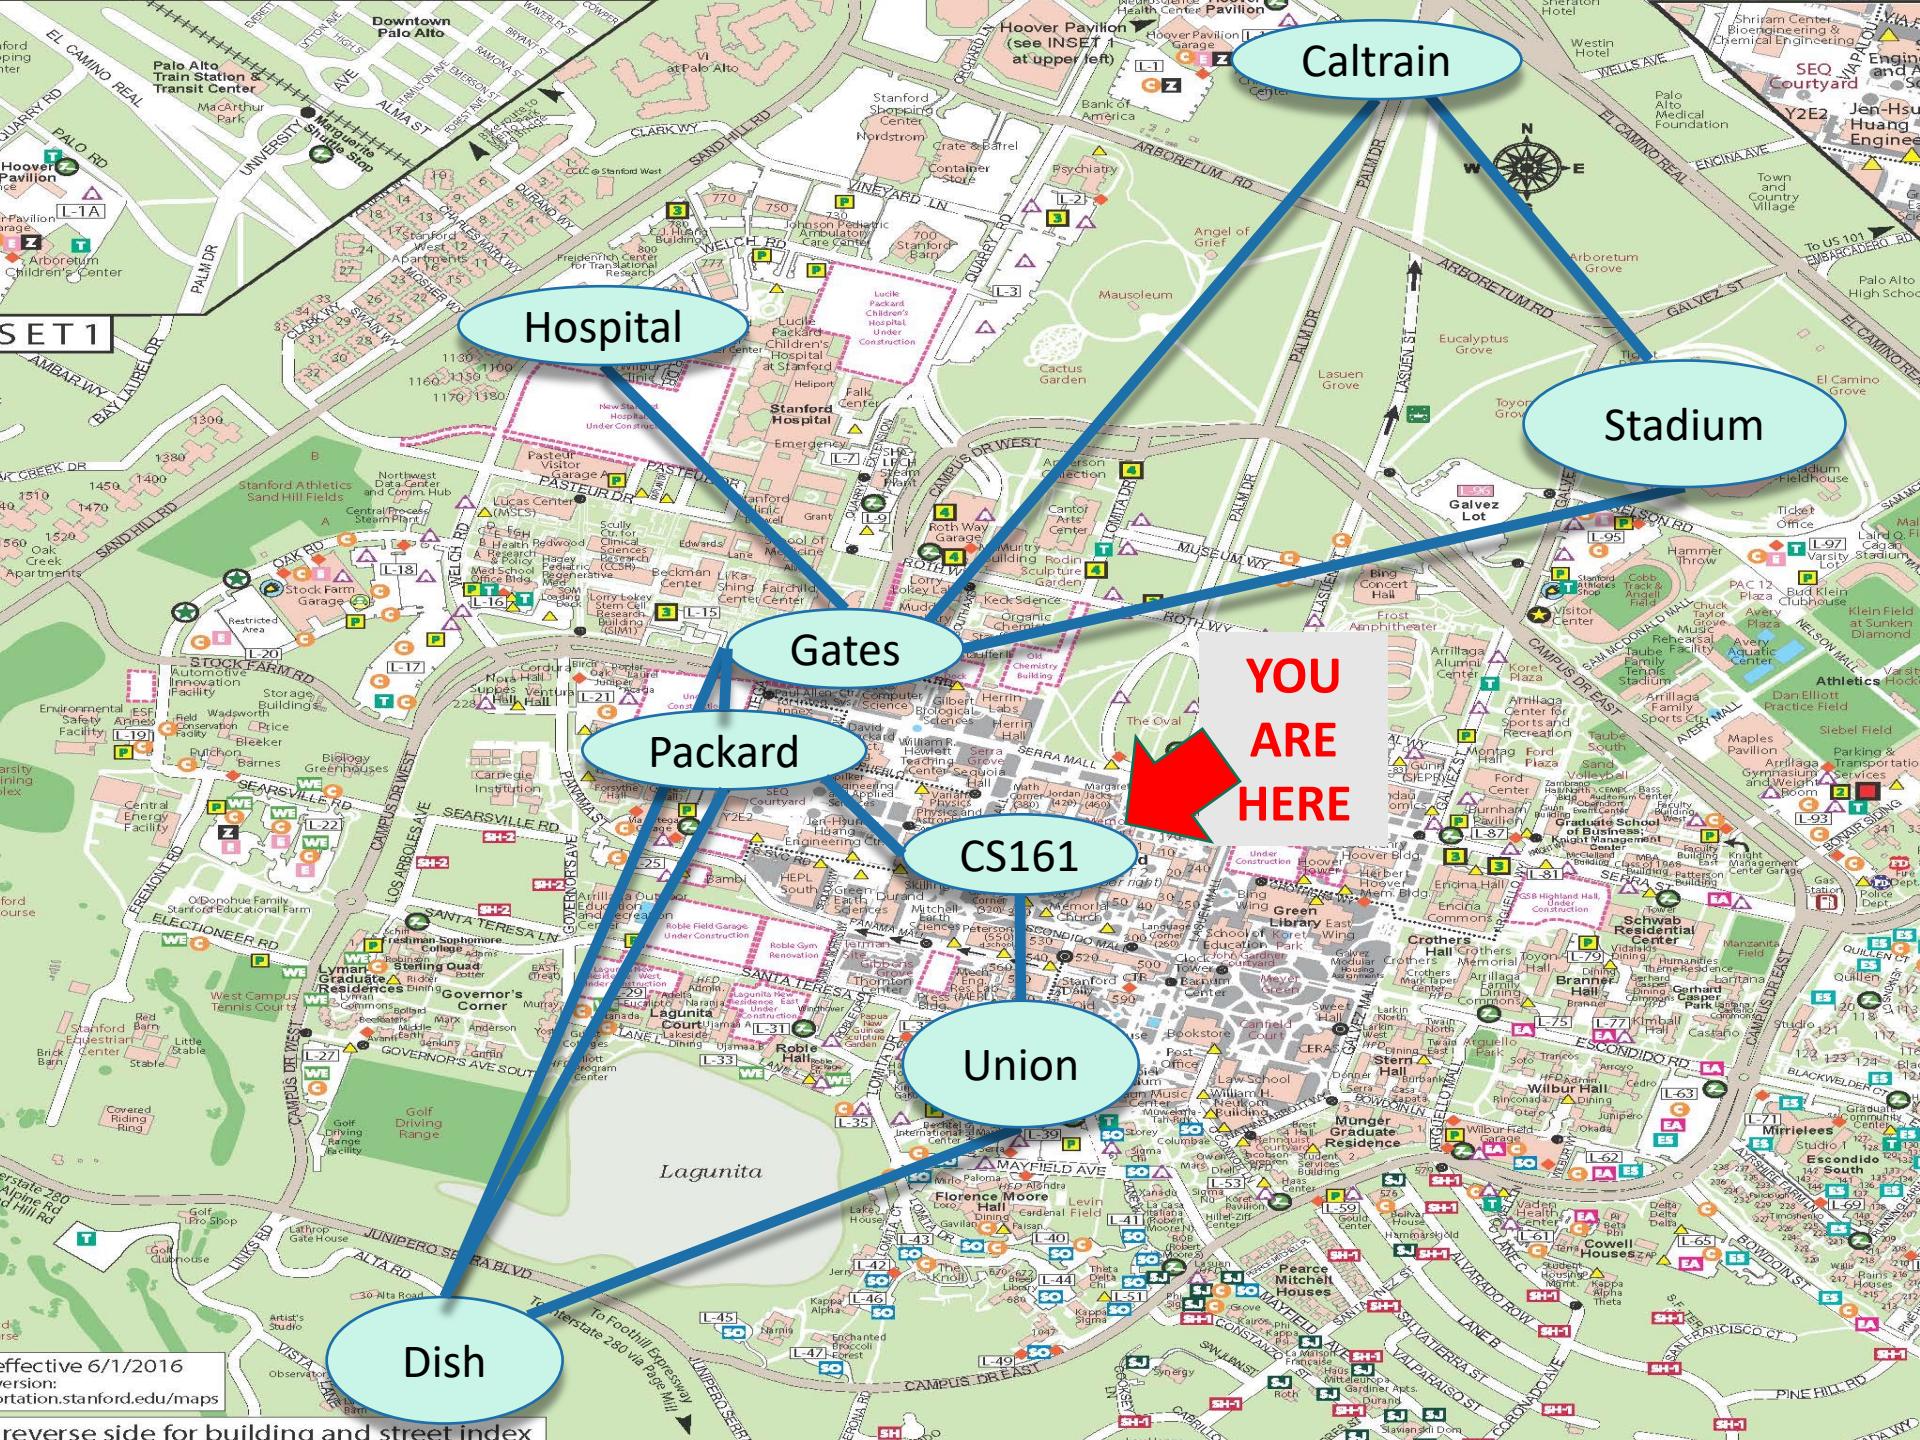


Module – Graphs II

Modified/Inspired from Stanford's CS161 by Nayyar
Zaidi

This module

- What if the graphs are **weighted**?
 - All nonnegative weights: Dijkstra!
 - If there are negative weights: Bellman-Ford!



Caltrain

Hospital

Stadium

YOU
ARE
HERE

S161

Union

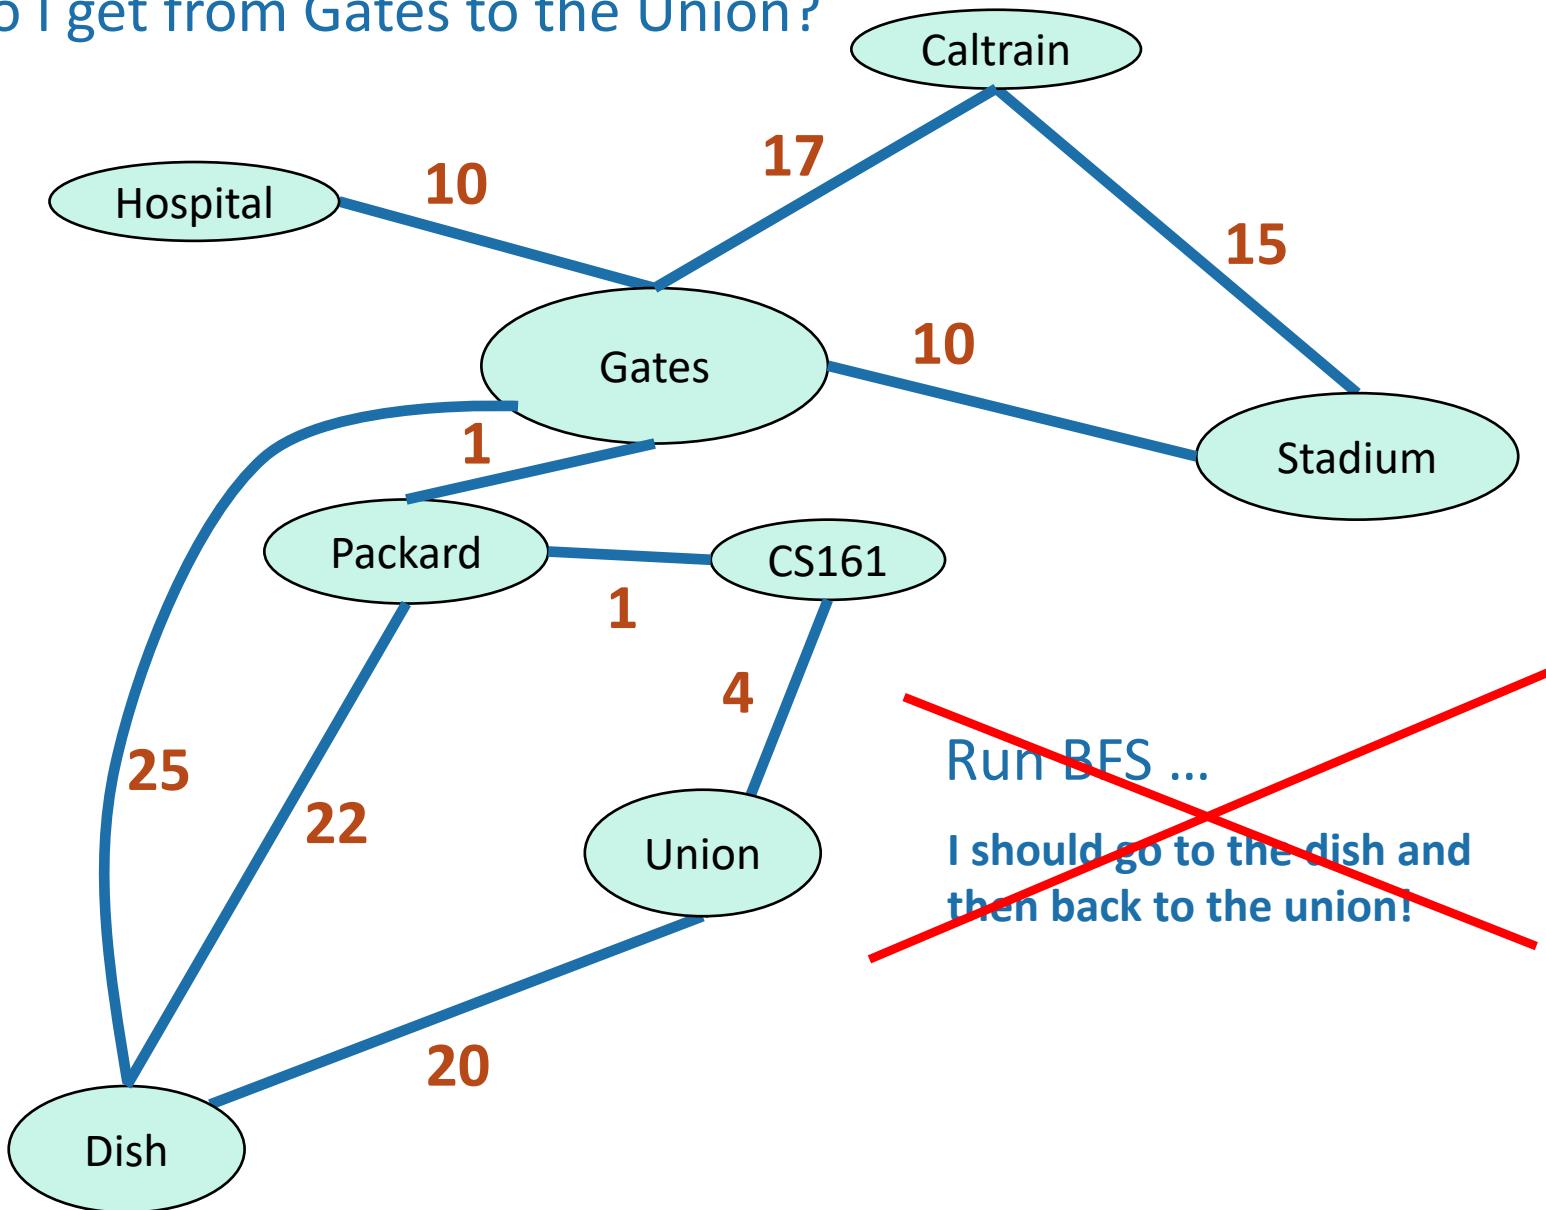
Dish

effective 6/1/2016
version:
stanford.edu/manc

reverse side for building and street index

Just the graph

How do I get from Gates to the Union?



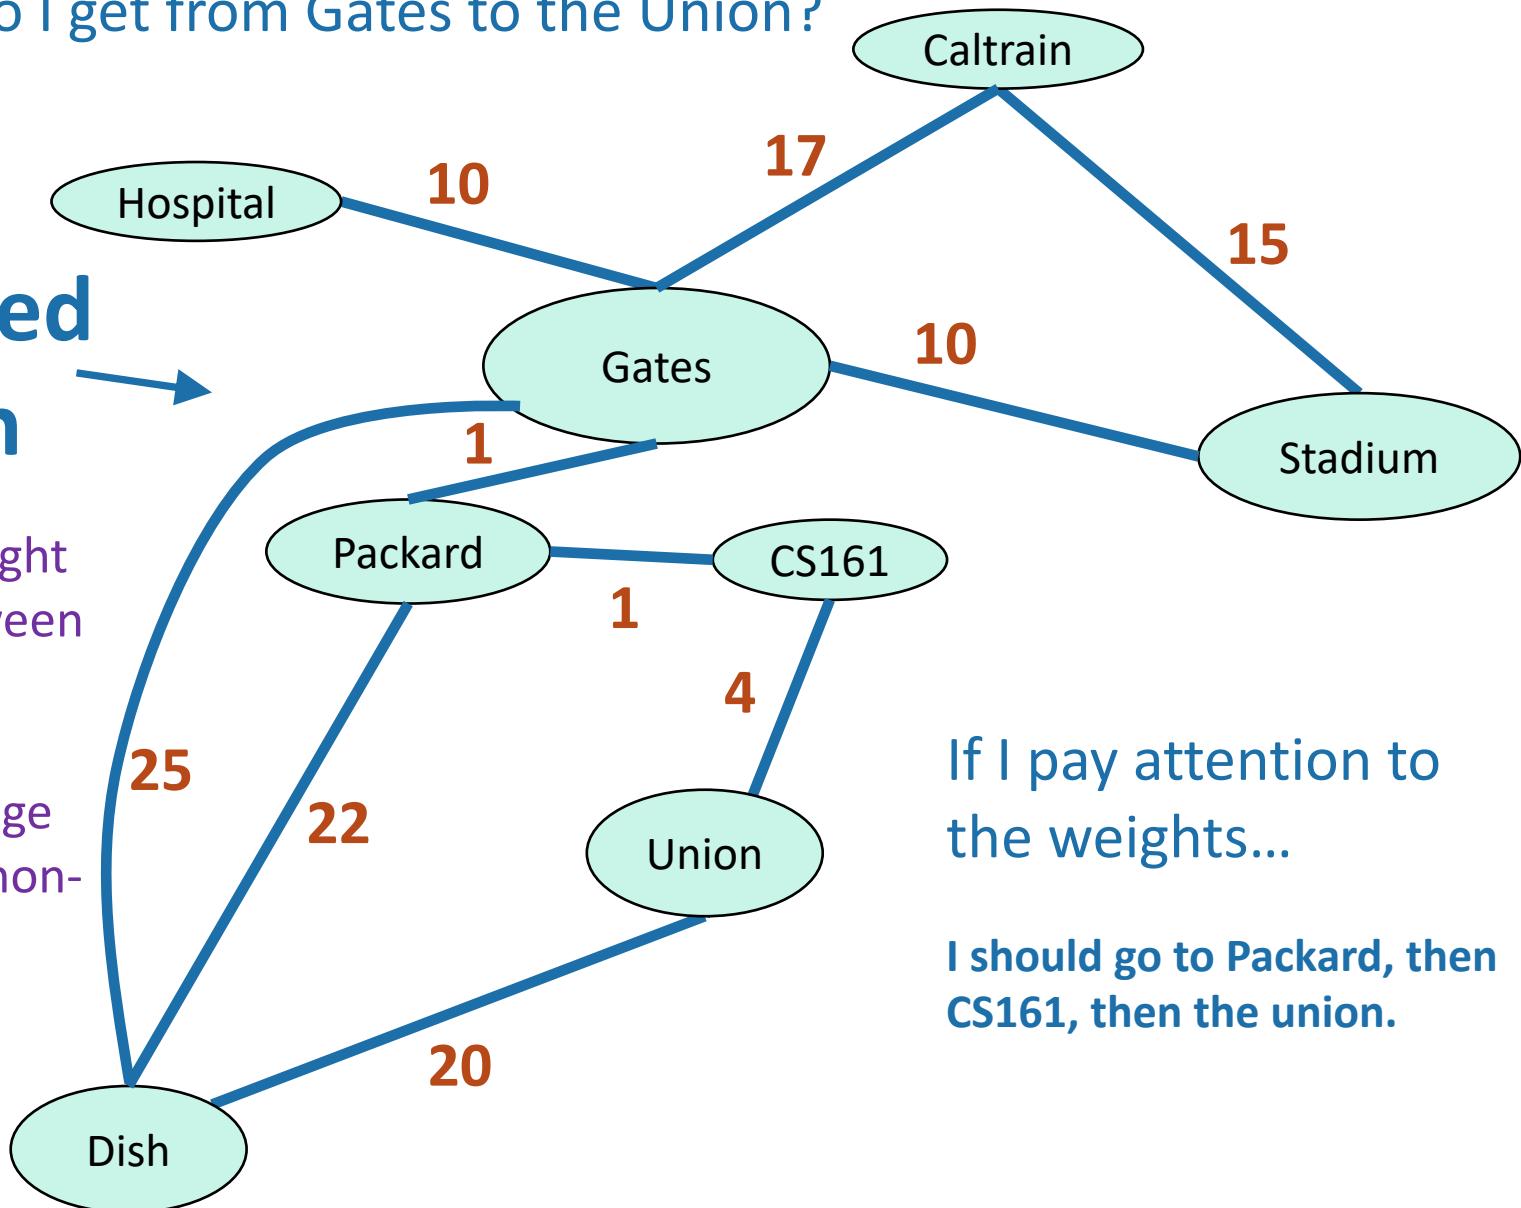
Just the graph

How do I get from Gates to the Union?

**weighted
graph**

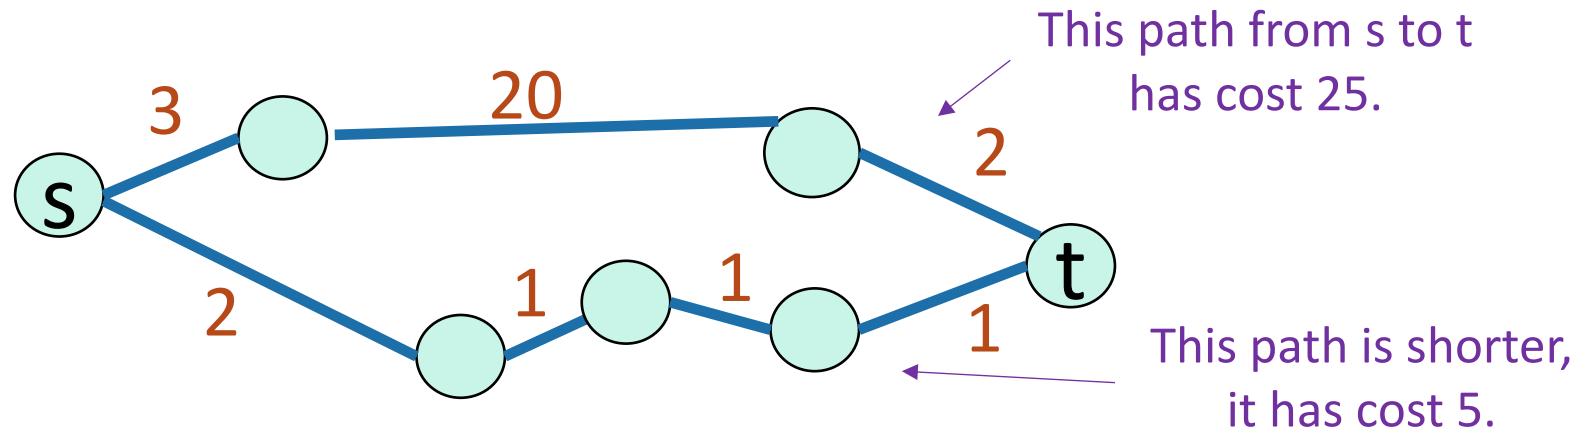
$w(u,v)$ = weight
of edge between
 u and v .

For now, edge
weights are non-
negative.



Shortest path problem

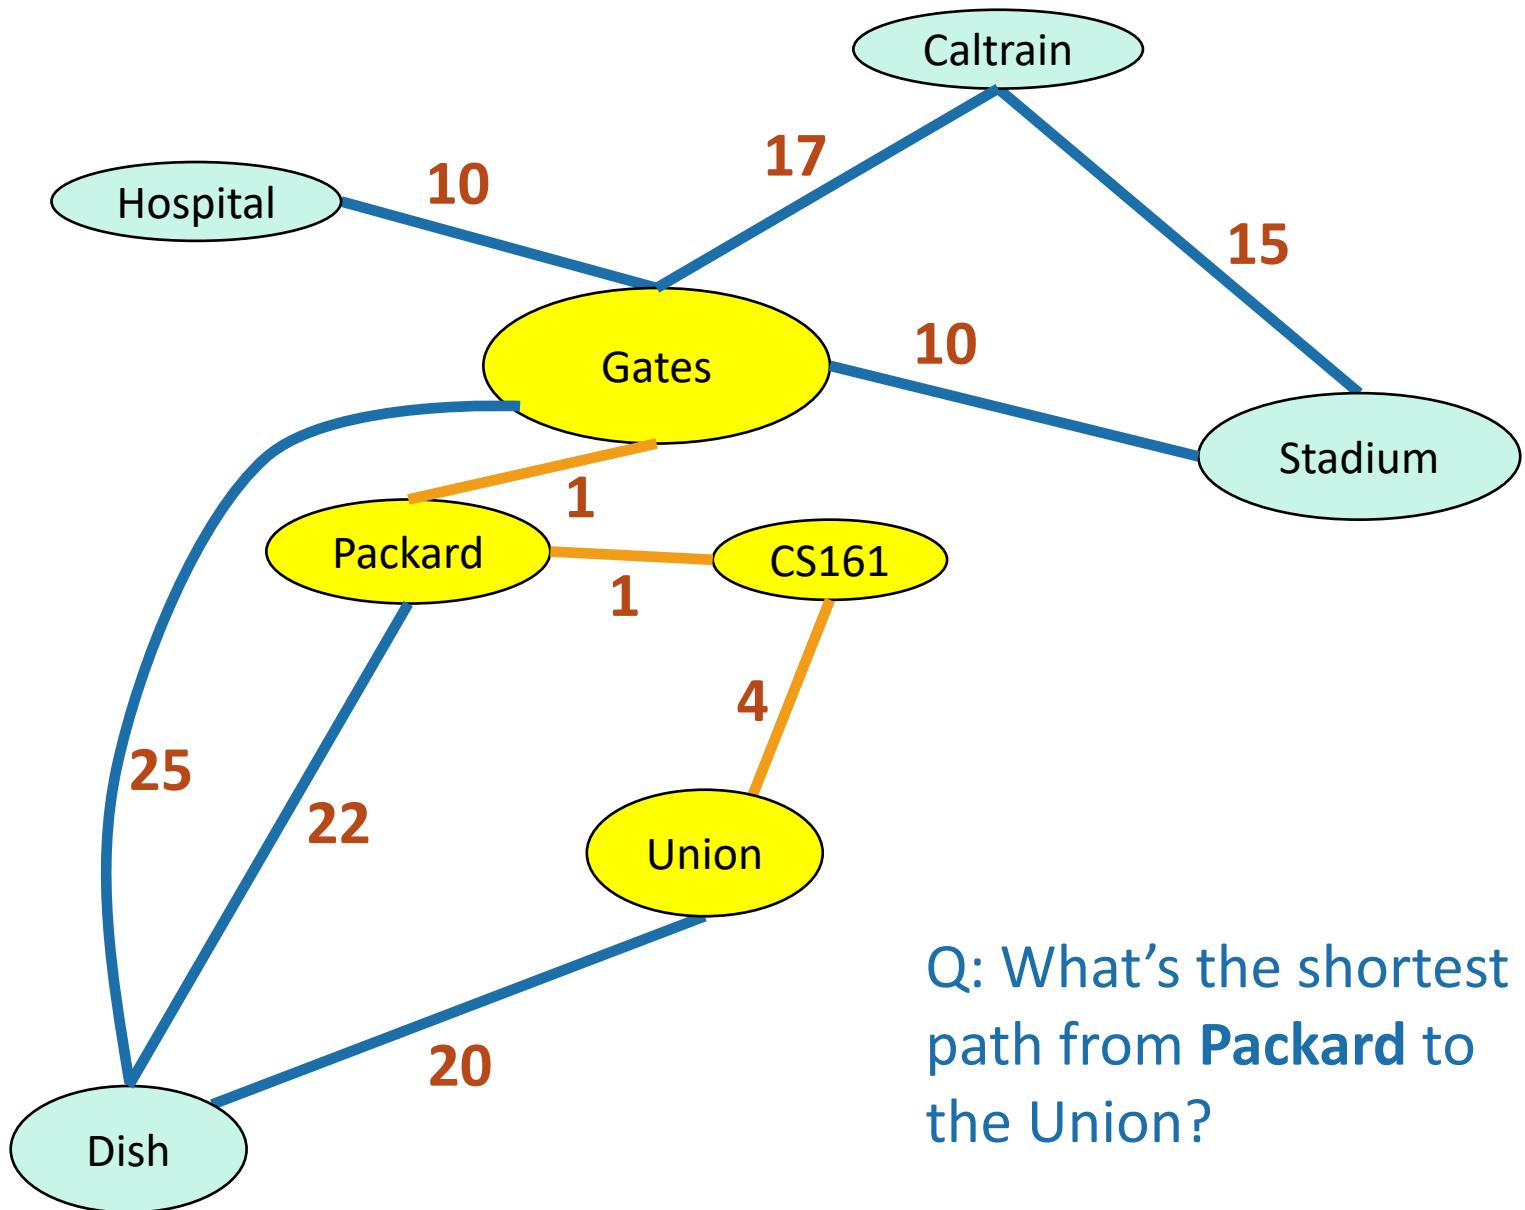
- What is the **shortest path** between u and v in a weighted graph?
 - the **cost** of a path is the sum of the weights along that path
 - The **shortest path** is the one with the minimum cost.



- The **distance** $d(u,v)$ between two vertices u and v is the cost of the the shortest path between u and v .
- For this lecture **all graphs are directed**, but to save on notation I'm just going to draw undirected edges.



Shortest paths

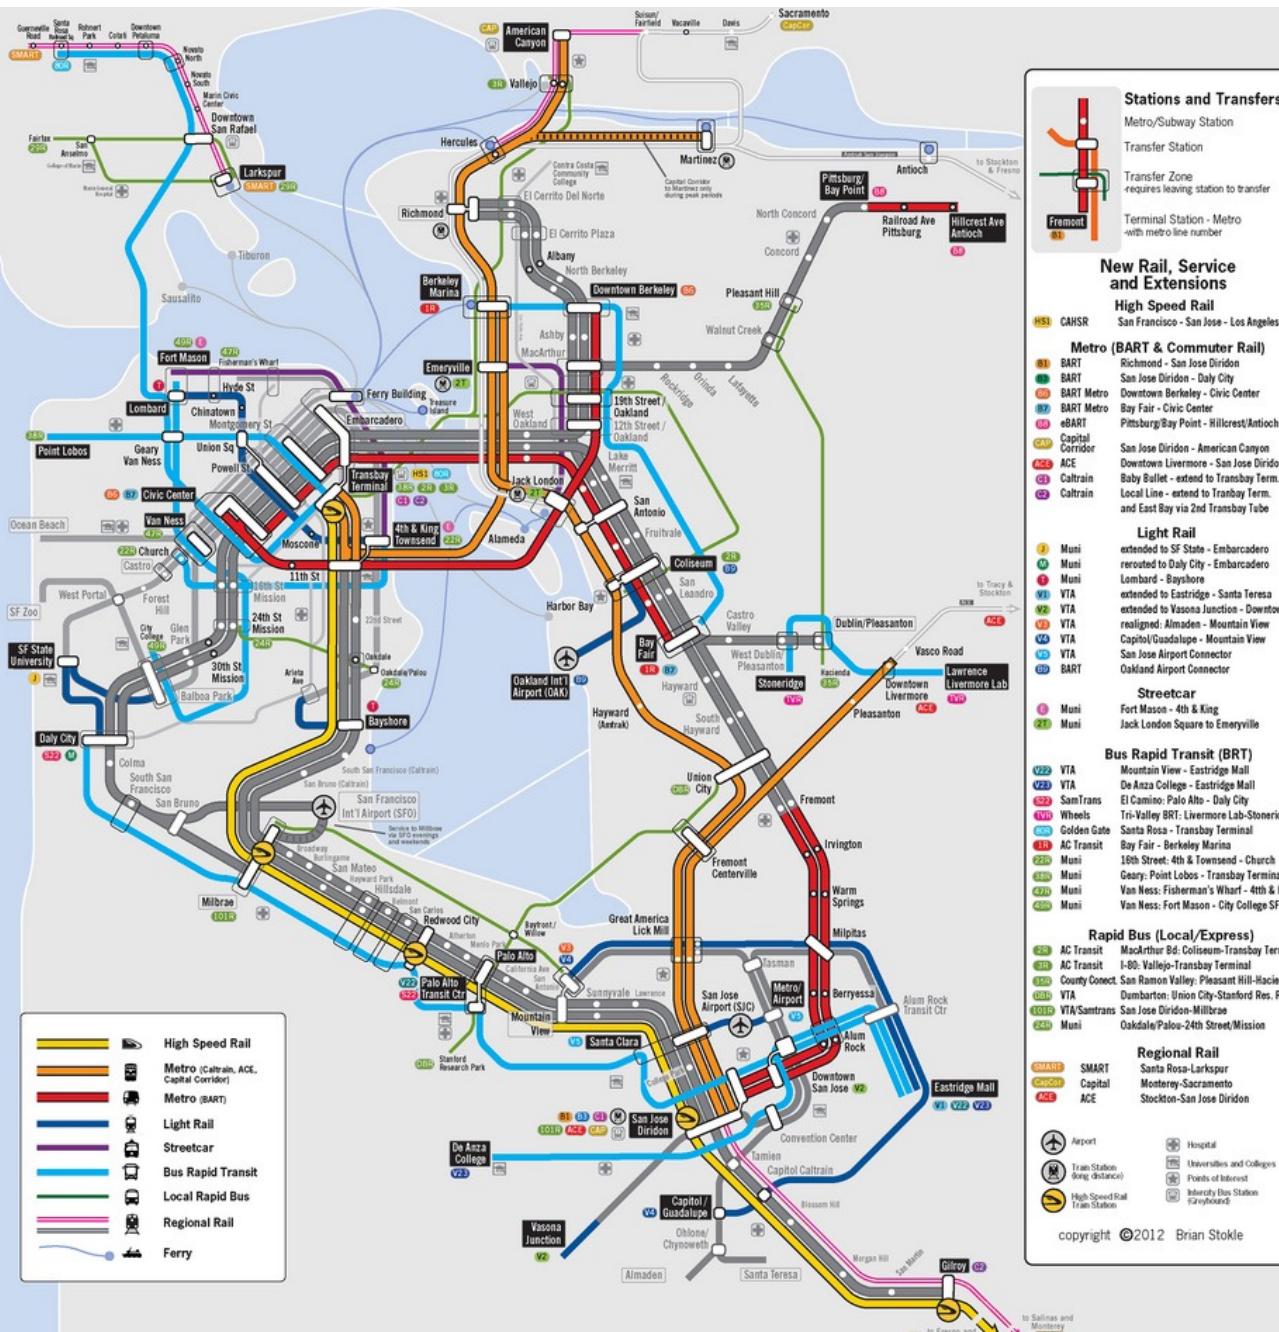


Single-source shortest-path problem

- I want to know the shortest path from one vertex (**Gates**) to all other vertices.

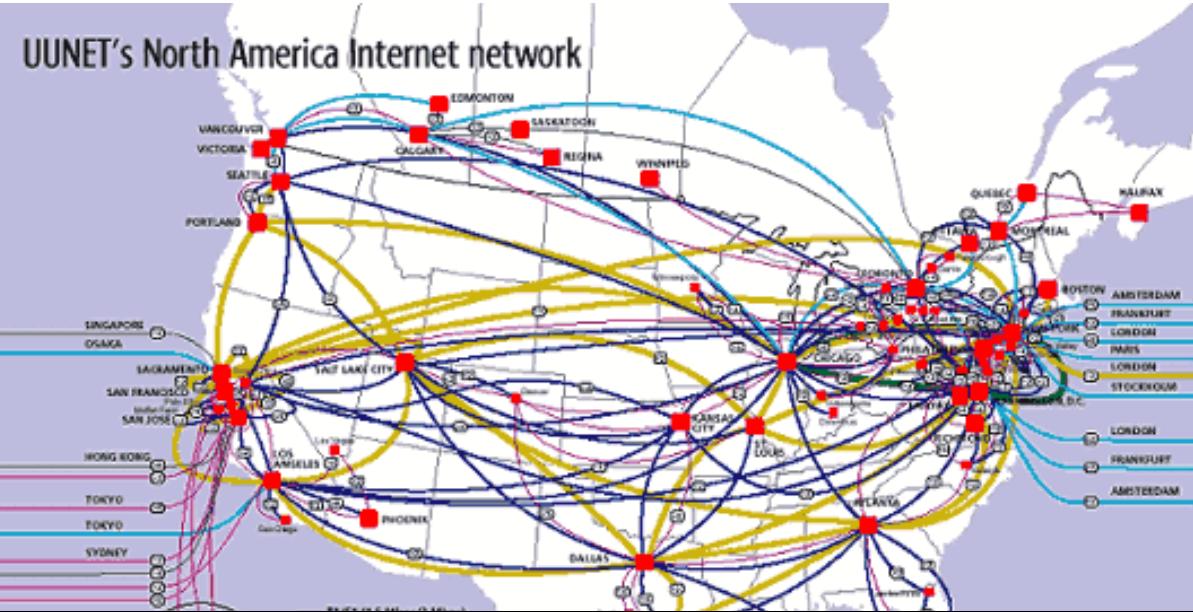
Destination	Cost	To get there
Packard	1	Packard
CS161	2	Packard-CS161
Hospital	10	Hospital
Caltrain	17	Caltrain
Union	6	Packard-CS161-Union
Stadium	10	Stadium
Dish	23	Packard-Dish

Example



Example

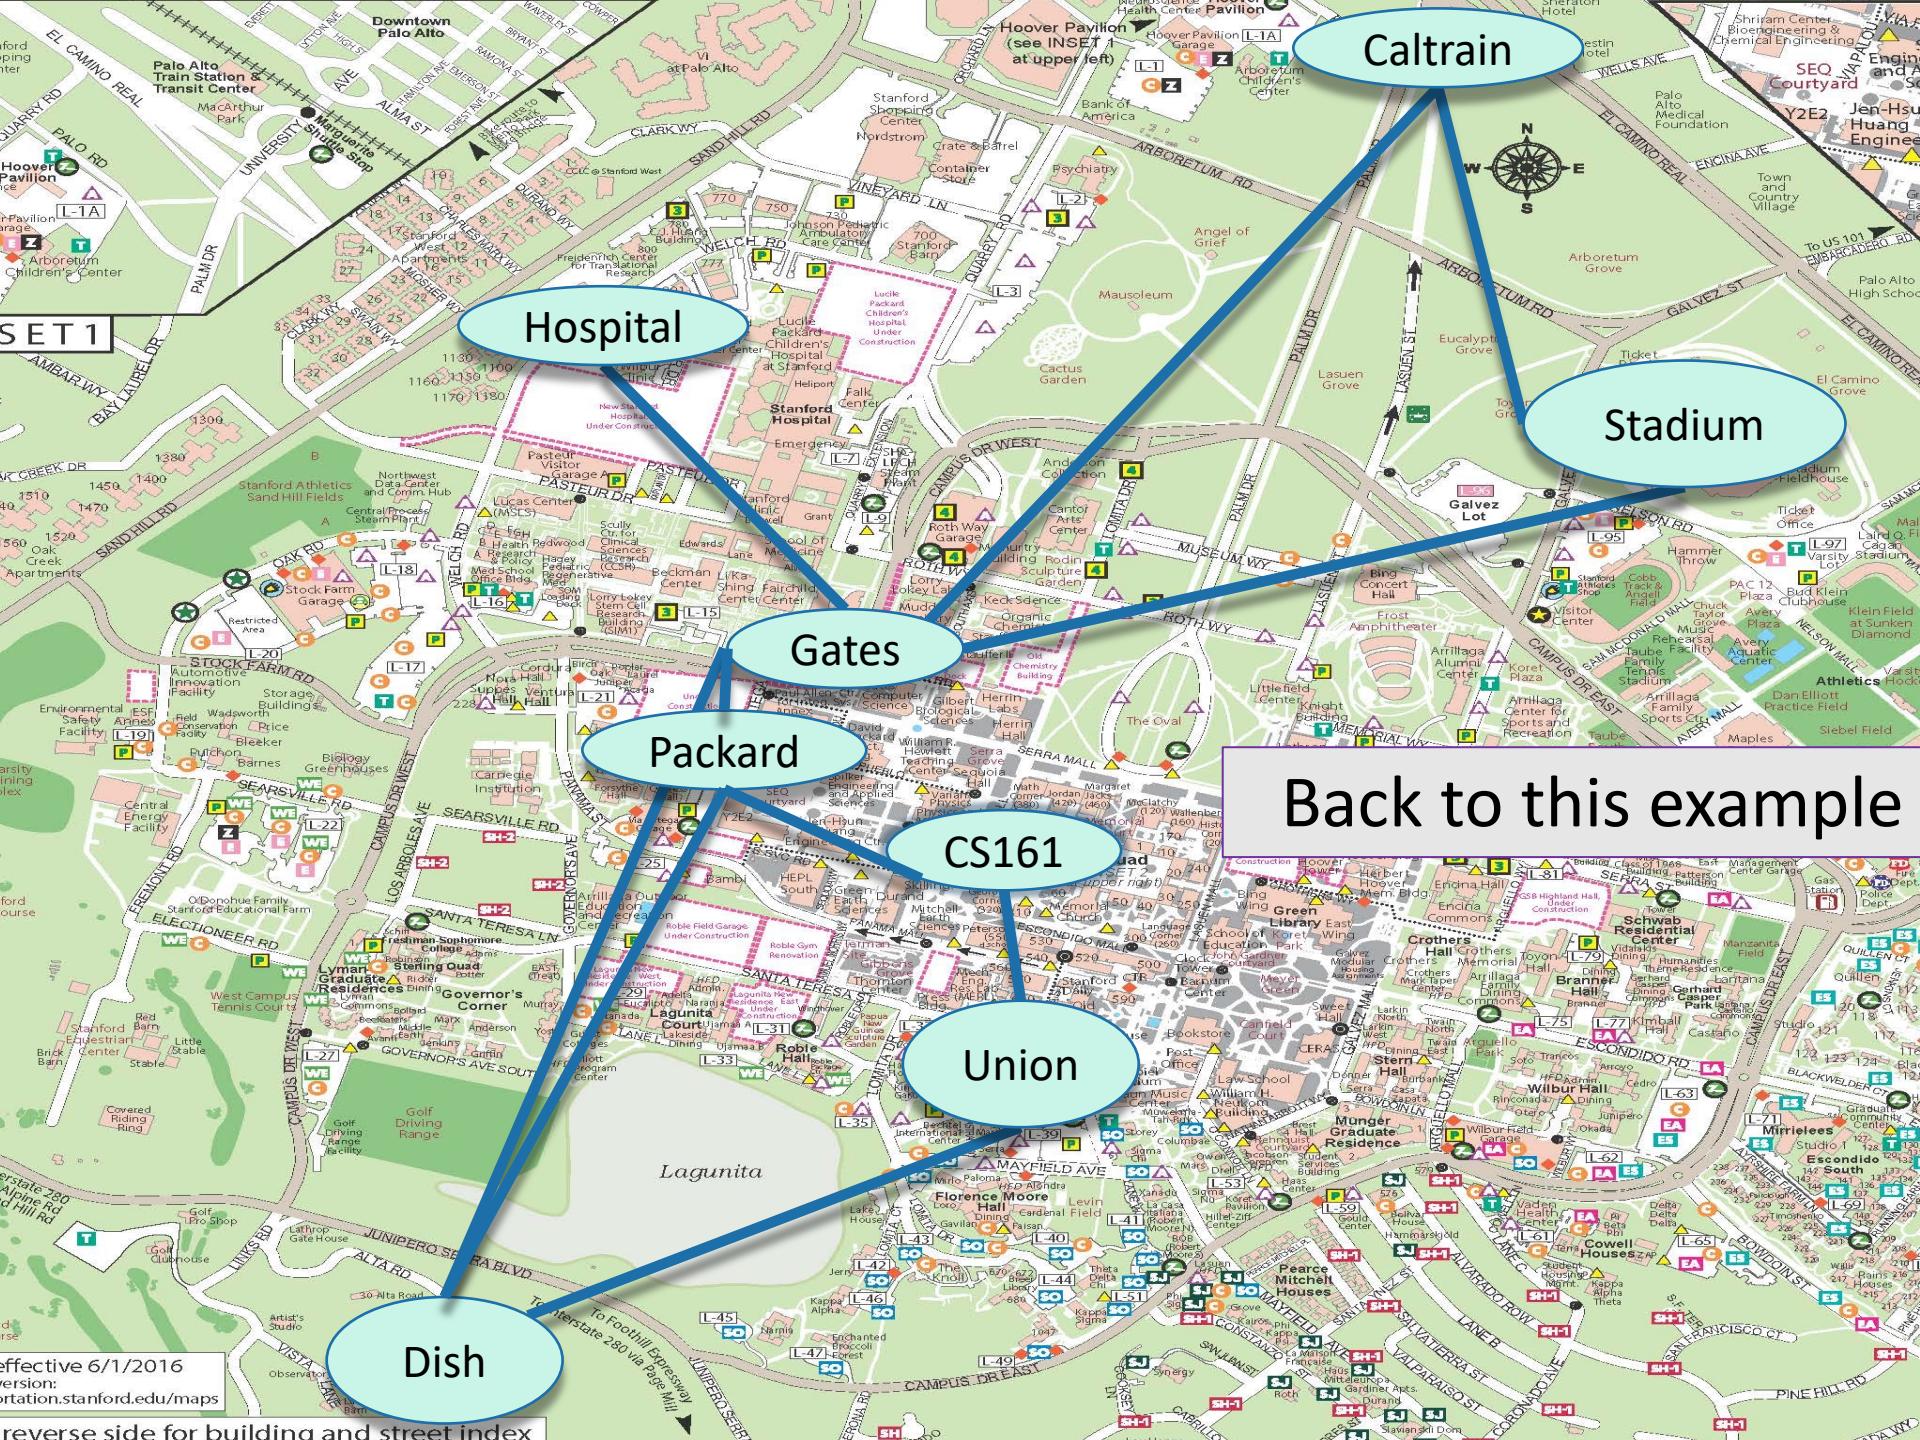
- **Network routing**
- I send information over the internet, from my computer to all over the world.
- Each path has a cost which depends on link length, traffic, other costs, etc..
- **How should we send packets?**



```
[DN0a22a0e3:~ mary$ traceroute -a www.ethz.ch
traceroute to www.ethz.ch (129.132.19.216), 64 hops max, 52 byte packets
 1 [AS0] 10.34.160.2 (10.34.160.2) 38.168 ms 31.272 ms 28.841 ms
 2 [AS0] cwa-vrtr.sunet (10.21.196.28) 33.769 ms 28.245 ms 24.373 ms
 3 [AS32] 171.66.2.229 (171.66.2.229) 24.468 ms 20.115 ms 23.223 ms
 4 [AS32] hpr-svl-rtr-vlan8.sunet (171.64.255.235) 24.644 ms 24.962 ms 11.111 ms
 5 [AS2152] hpr-svl-hpr2--stan-ge.cenic.net (137.164.27.161) 22.129 ms 4.916 ms 1.111 ms
 6 [AS2152] hpr-lax-hpr3--svl-hpr3-100ge.cenic.net (137.164.25.73) 12.125 ms 1.111 ms 1.111 ms
 7 [AS2152] hpr-i2--lax-hpr2-r&e.cenic.net (137.164.26.201) 40.174 ms 38.333 ms 1.111 ms
 8 [AS0] et-4-0-0.4079.sdn-sw.lasv.net.internet2.edu (162.252.70.28) 46.573 ms 1.111 ms 1.111 ms
 9 [AS0] et-5-1-0.4079.rtsw.salt.net.internet2.edu (162.252.70.31) 30.424 ms 1.111 ms 1.111 ms
10 [AS0] et-4-0-0.4079.sdn-sw.denv.net.internet2.edu (162.252.70.8) 47.454 ms 1.111 ms 1.111 ms
11 [AS0] et-4-1-0.4079.rtsw.kans.net.internet2.edu (162.252.70.11) 70.825 ms 1.111 ms 1.111 ms
12 [AS0] et-4-1-0.4070.rtsw.chic.net.internet2.edu (198.71.47.206) 77.937 ms 1.111 ms 1.111 ms
13 [AS0] et-0-1-0.4079.sdn-sw.ashb.net.internet2.edu (162.252.70.60) 77.682 ms 1.111 ms 1.111 ms
14 [AS0] et-4-1-0.4079.rtsw.wash.net.internet2.edu (162.252.70.65) 71.565 ms 1.111 ms 1.111 ms
15 [AS21320] internet2-gw.mx1.lon.uk.geant.net (62.40.124.44) 154.926 ms 1.111 ms 1.111 ms
16 [AS21320] ae0.mx1.lon2.uk.geant.net (62.40.98.79) 146.565 ms 146.604 ms 1.111 ms 1.111 ms
17 [AS21320] ae0.mx1.par.fr.geant.net (62.40.98.77) 153.289 ms 184.995 ms 1.111 ms 1.111 ms
18 [AS21320] ae2.mx1.gen.ch.geant.net (62.40.98.153) 160.283 ms 160.104 ms 1.111 ms 1.111 ms
19 [AS21320] swice1-100ge-0-3-0-1.switch.ch (62.40.124.22) 162.068 ms 160.104 ms 1.111 ms 1.111 ms
20 [AS559] swizh1-100ge-0-1-0-1.switch.ch (130.59.36.94) 165.824 ms 164.221 ms 1.111 ms 1.111 ms
21 [AS559] swiez3-100ge-0-1-0-4.switch.ch (130.59.38.109) 164.269 ms 164.221 ms 1.111 ms 1.111 ms
22 [AS559] rou-gw-lee-tengig-to-switch.ethz.ch (192.33.92.1) 164.082 ms 164.221 ms 1.111 ms 1.111 ms
23 [AS559] rou-fw-rz-rz-gw.ethz.ch (192.33.92.169) 164.773 ms 165.193 ms 1.111 ms 1.111 ms
```

Aside: These are difficult problems

- Costs may change
 - If it's raining the cost of biking is higher
 - If a link is congested, the cost of routing a packet along it is higher
- The network might not be known
 - My computer doesn't store a map of the internet
- We want to do these tasks really quickly

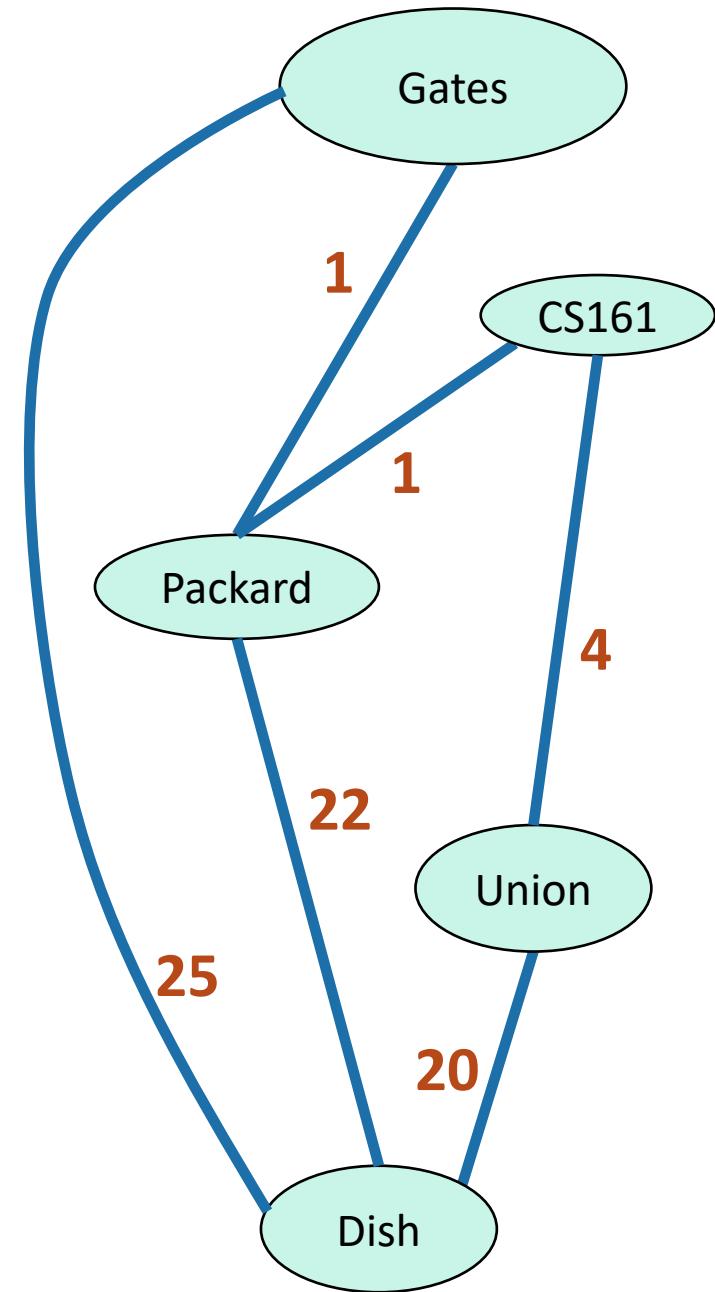


Effective 6/1/2016
Version:
version.stanford.edu/maps

reverse side for building and street index

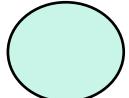
Dijkstra's algorithm

- What are the shortest paths from Gates to everywhere else?

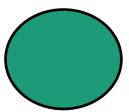


Dijkstra by example

How far is a node from Gates?



I'm not sure yet



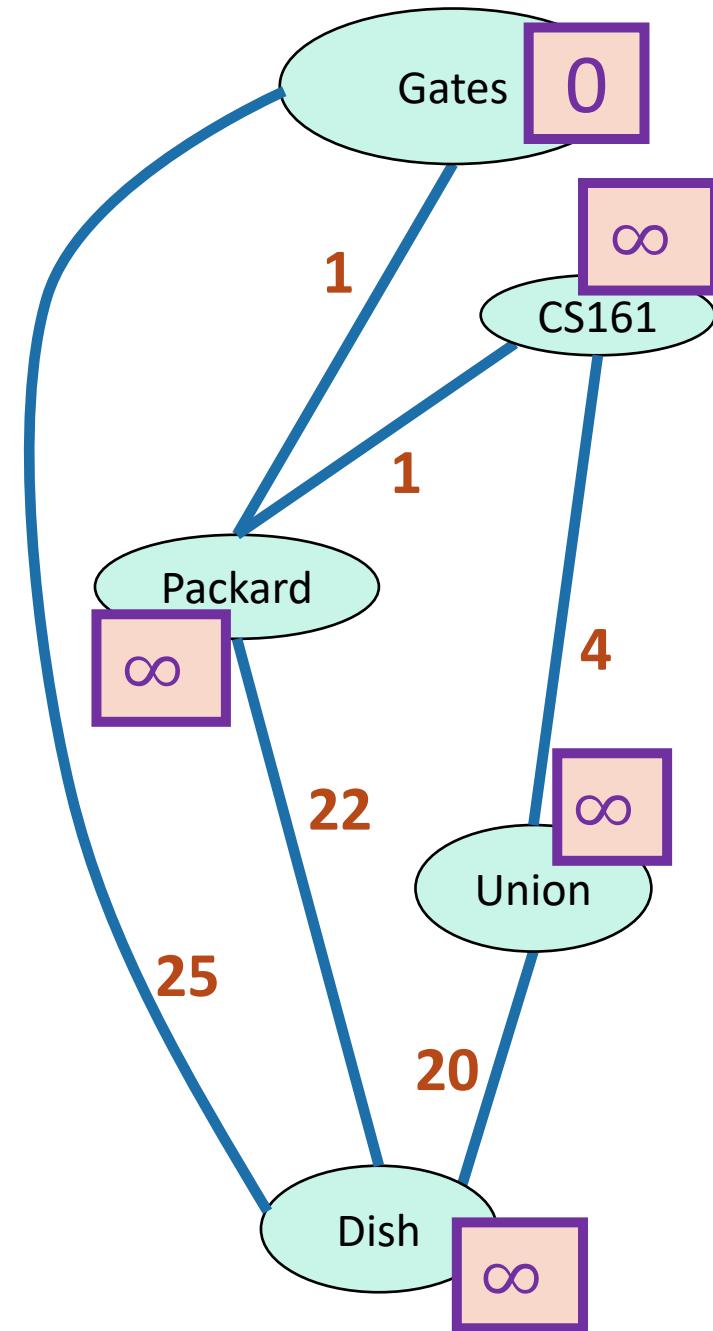
I'm sure



$x = d[v]$ is my best **over-estimate** for $\text{dist}(\text{Gates}, v)$.

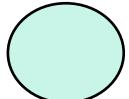
Initialize $d[v] = \infty$ for all non-starting vertices v , and $d[\text{Gates}] = 0$

- Pick the **not-sure** node u with the smallest estimate $d[u]$.

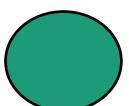


Dijkstra by example

How far is a node from Gates?



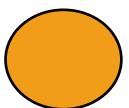
I'm not sure yet



I'm sure

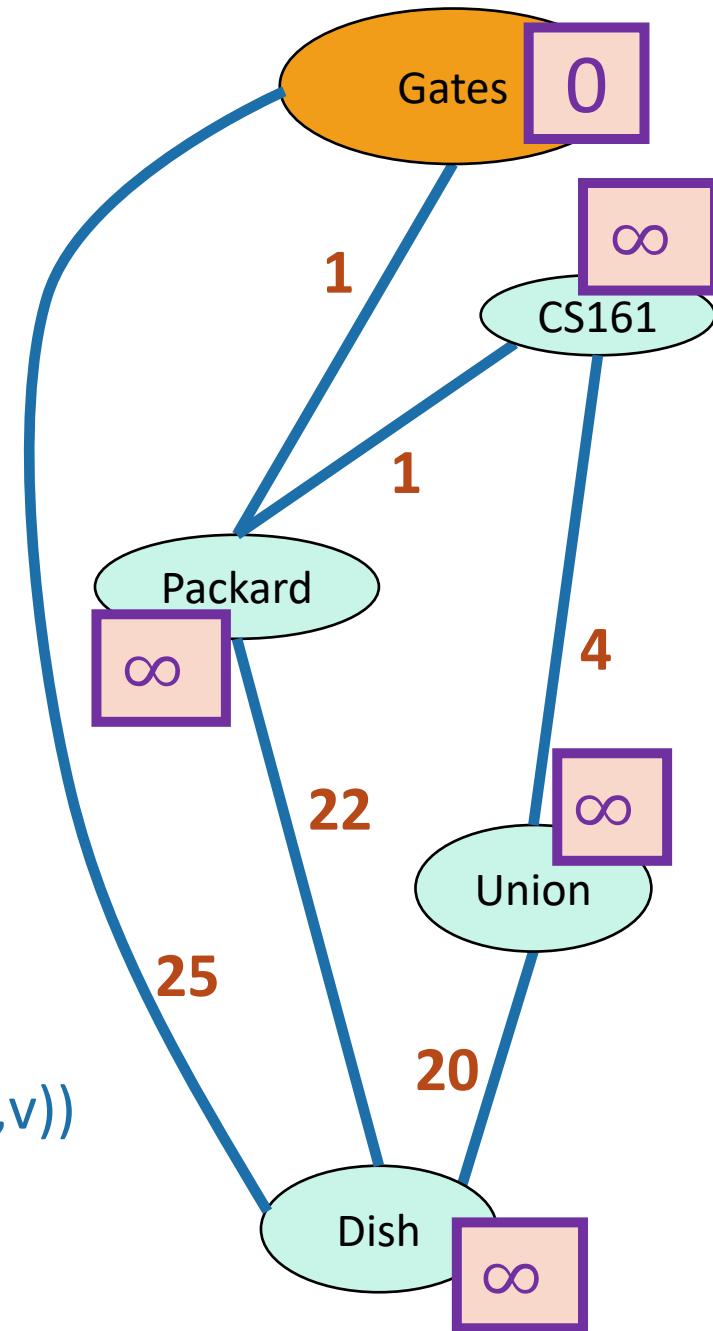


$x = d[v]$ is my best **over-estimate** for $\text{dist}(\text{Gates}, v)$.



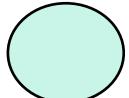
Current node u

- Pick the **not-sure** node u with the smallest estimate $d[u]$.
- Update all u's neighbors v:
 - $d[v] = \min(d[v] , d[u] + \text{edgeWeight}(u,v))$

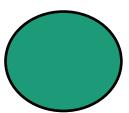


Dijkstra by example

How far is a node from Gates?



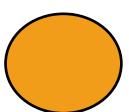
I'm not sure yet



I'm sure

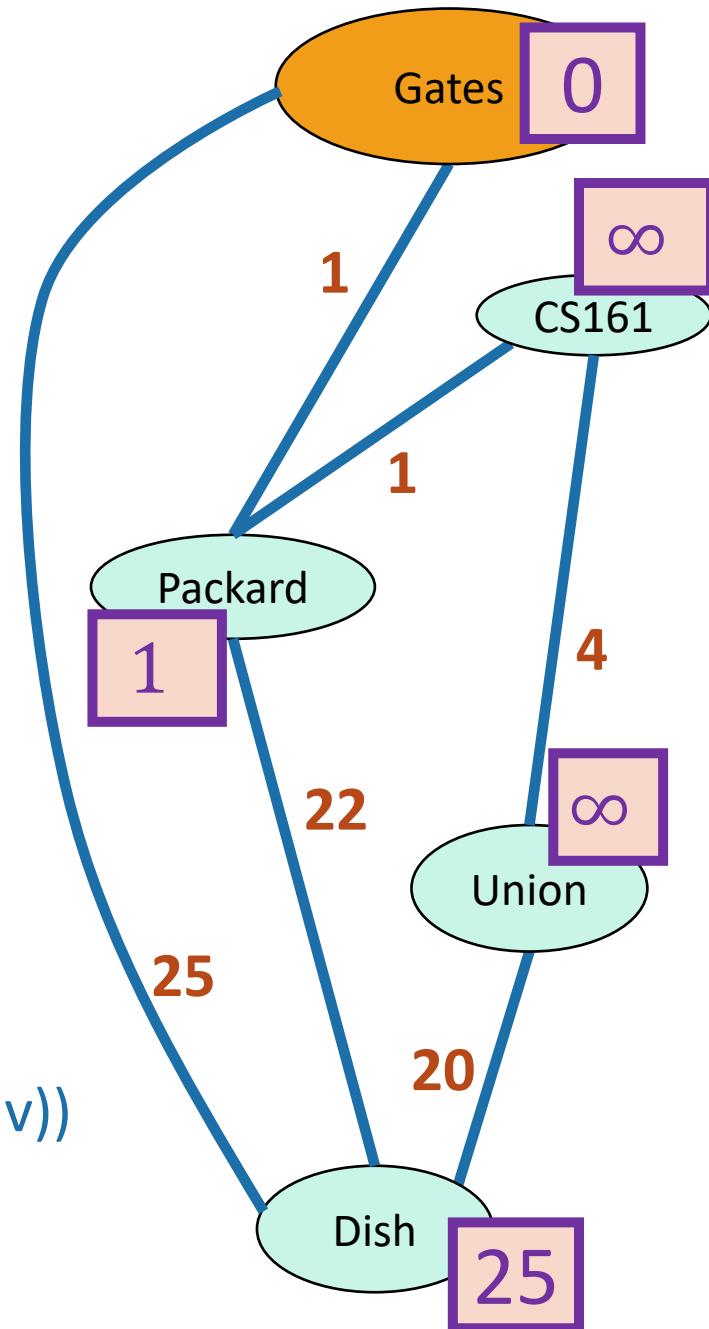


$x = d[v]$ is my best **over-estimate** for $\text{dist}(\text{Gates}, v)$.



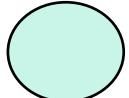
Current node u

- Pick the **not-sure** node u with the smallest estimate $d[u]$.
- Update all u's neighbors v:
 - $d[v] = \min(d[v] , d[u] + \text{edgeWeight}(u,v))$
- Mark u as **Sure**.

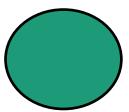


Dijkstra by example

How far is a node from Gates?



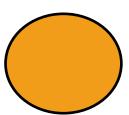
I'm not sure yet



I'm sure

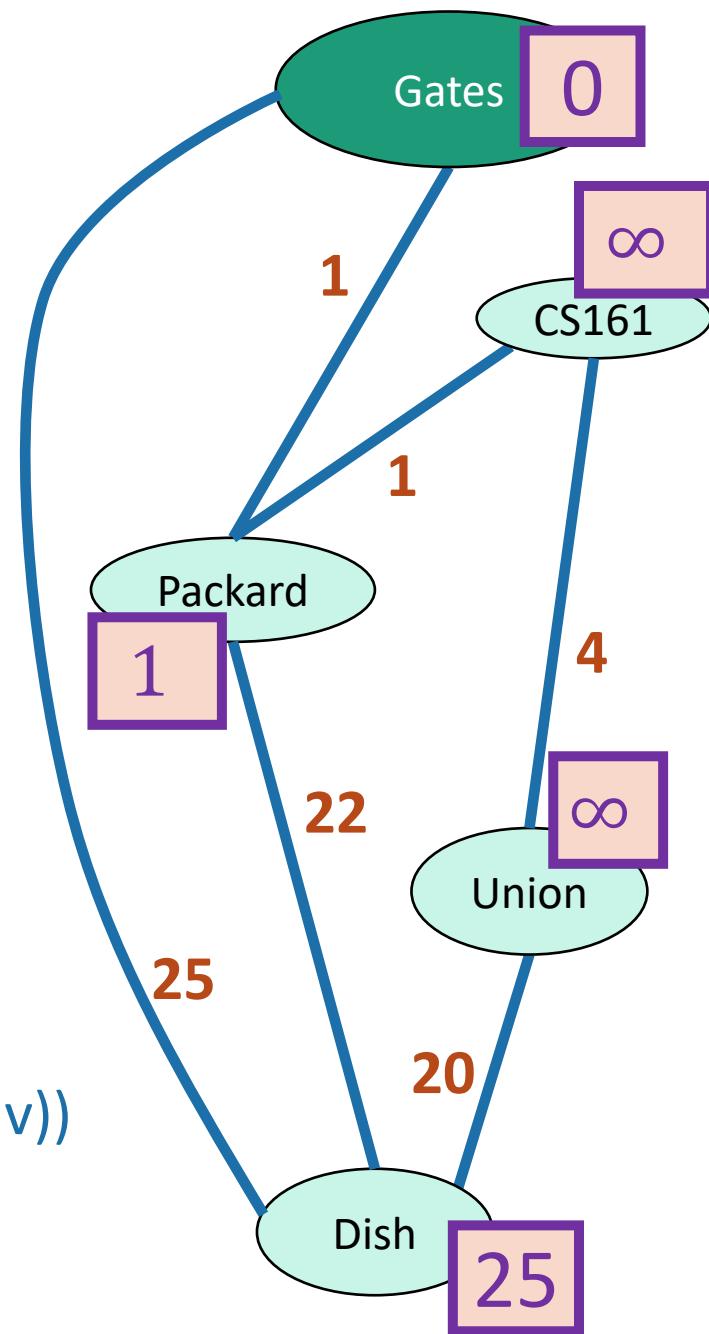


$x = d[v]$ is my best **over-estimate** for $\text{dist}(\text{Gates}, v)$.



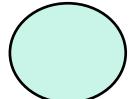
Current node u

- Pick the **not-sure** node u with the smallest estimate $d[u]$.
- Update all u's neighbors v:
 - $d[v] = \min(d[v] , d[u] + \text{edgeWeight}(u,v))$
- Mark u as **SURE**.
- Repeat

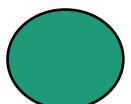


Dijkstra by example

How far is a node from Gates?



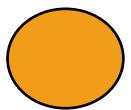
I'm not sure yet



I'm sure

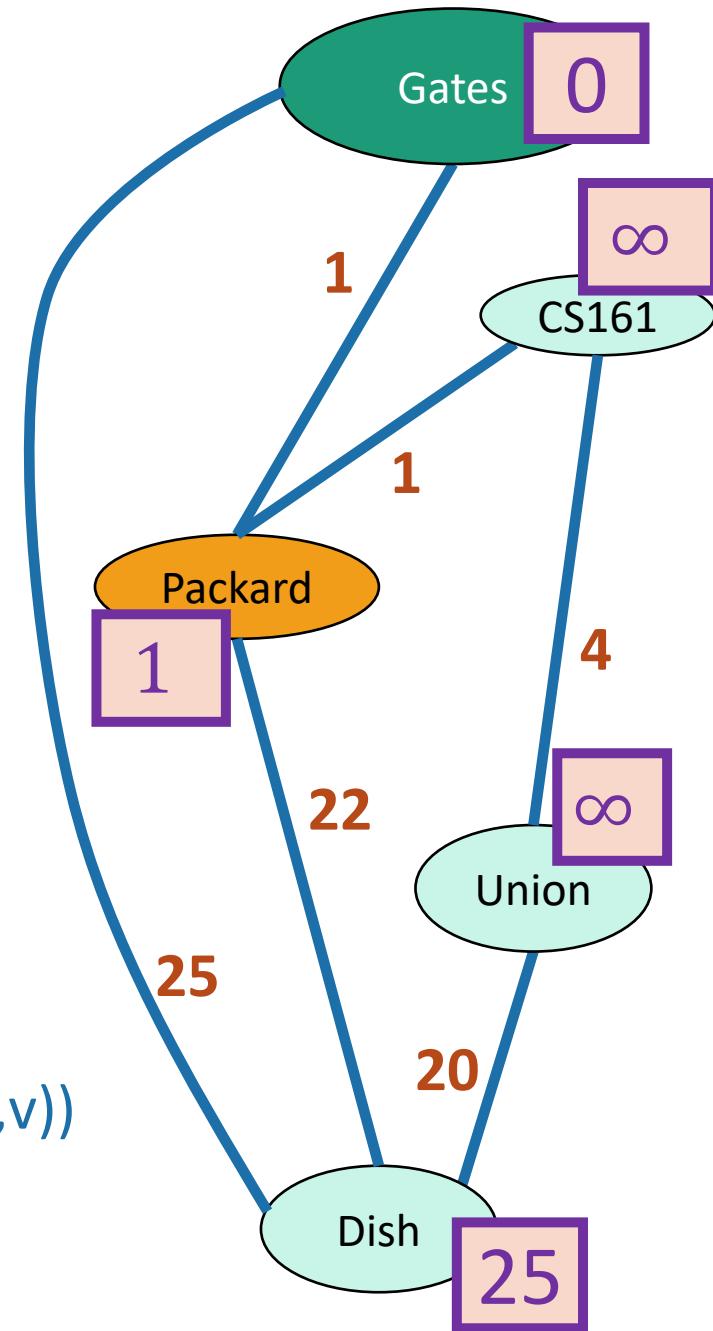


$x = d[v]$ is my best **over-estimate** for $\text{dist}(\text{Gates}, v)$.



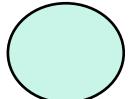
Current node u

- Pick the **not-sure** node u with the smallest estimate $d[u]$.
- Update all u's neighbors v:
 - $d[v] = \min(d[v] , d[u] + \text{edgeWeight}(u,v))$
- Mark u as **SURE**.
- Repeat

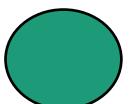


Dijkstra by example

How far is a node from Gates?



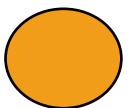
I'm not sure yet



I'm sure

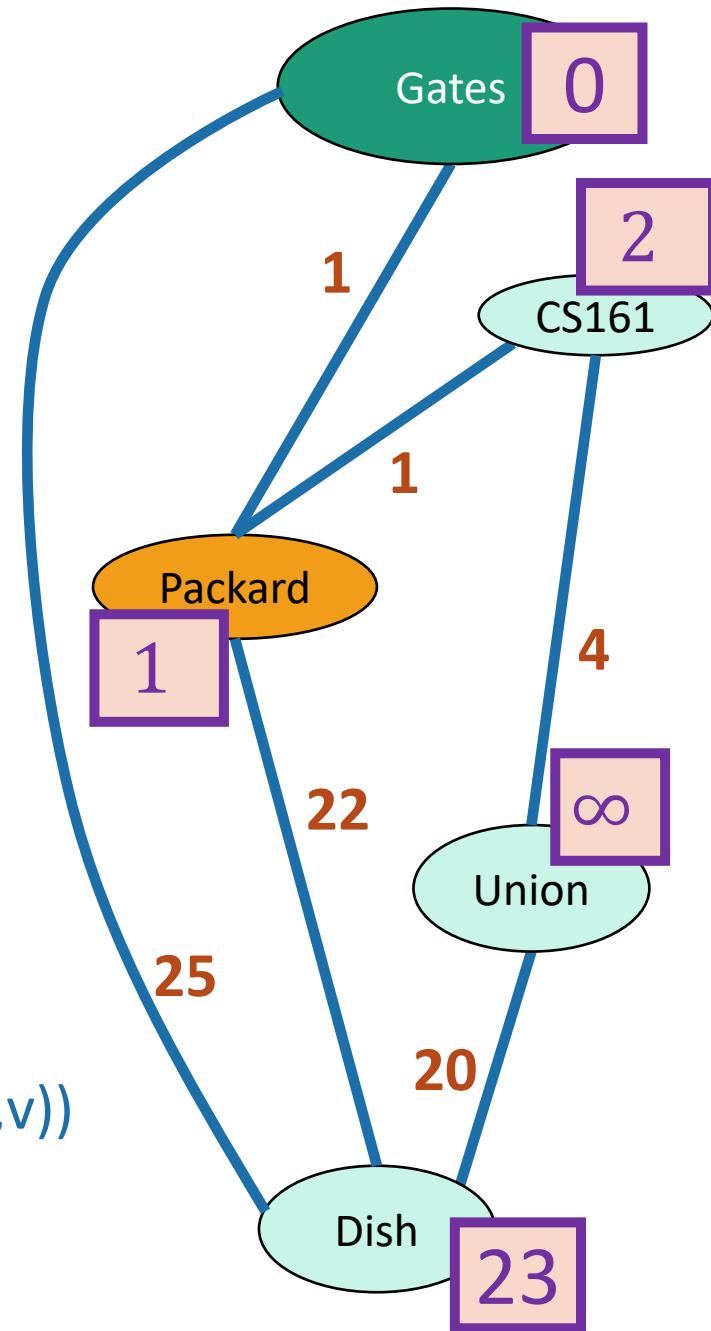


$x = d[v]$ is my best **over-estimate** for $\text{dist}(\text{Gates}, v)$.



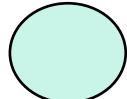
Current node u

- Pick the **not-sure** node u with the smallest estimate $d[u]$.
- Update all u's neighbors v:
 - $d[v] = \min(d[v] , d[u] + \text{edgeWeight}(u,v))$
- Mark u as **SURE**.
- Repeat

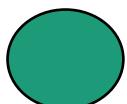


Dijkstra by example

How far is a node from Gates?



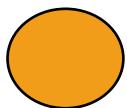
I'm not sure yet



I'm sure

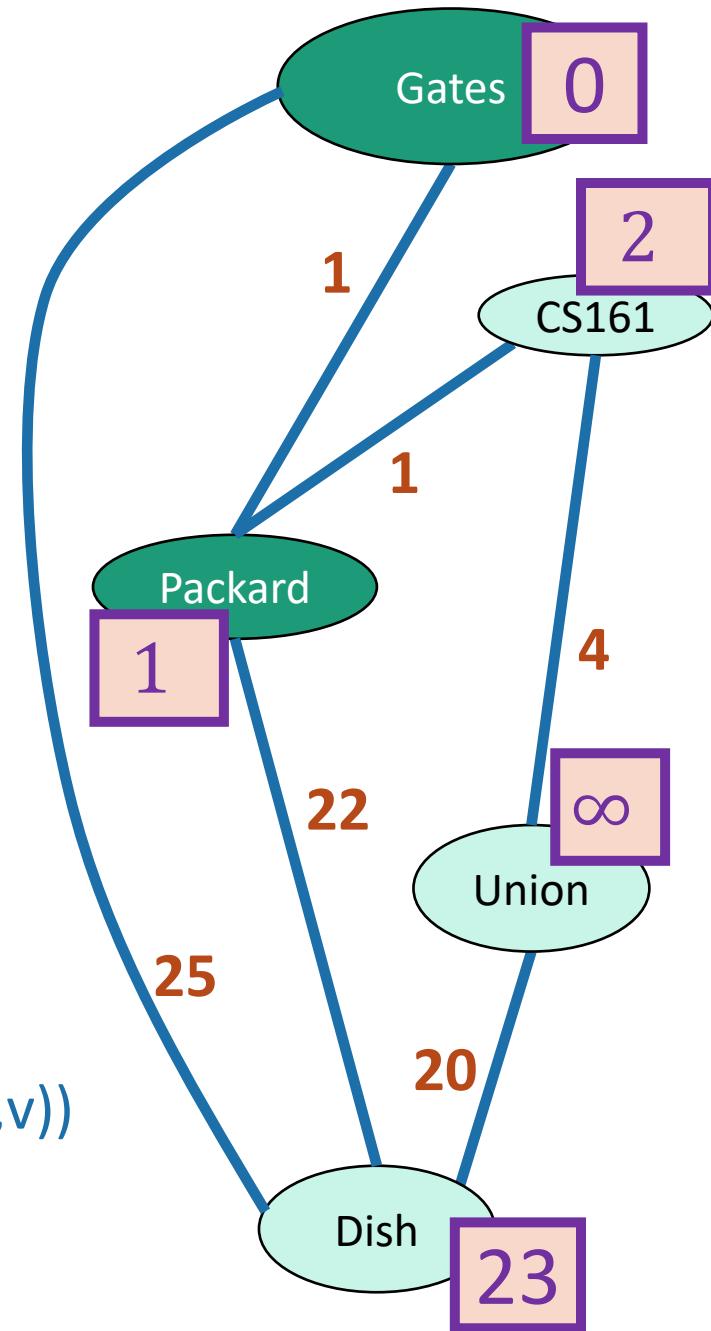


$x = d[v]$ is my best **over-estimate** for $\text{dist}(\text{Gates}, v)$.



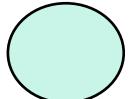
Current node u

- Pick the **not-sure** node u with the smallest estimate $d[u]$.
- Update all u's neighbors v:
 - $d[v] = \min(d[v] , d[u] + \text{edgeWeight}(u,v))$
- Mark u as **SURE**.
- Repeat

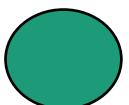


Dijkstra by example

How far is a node from Gates?



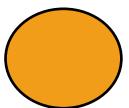
I'm not sure yet



I'm sure

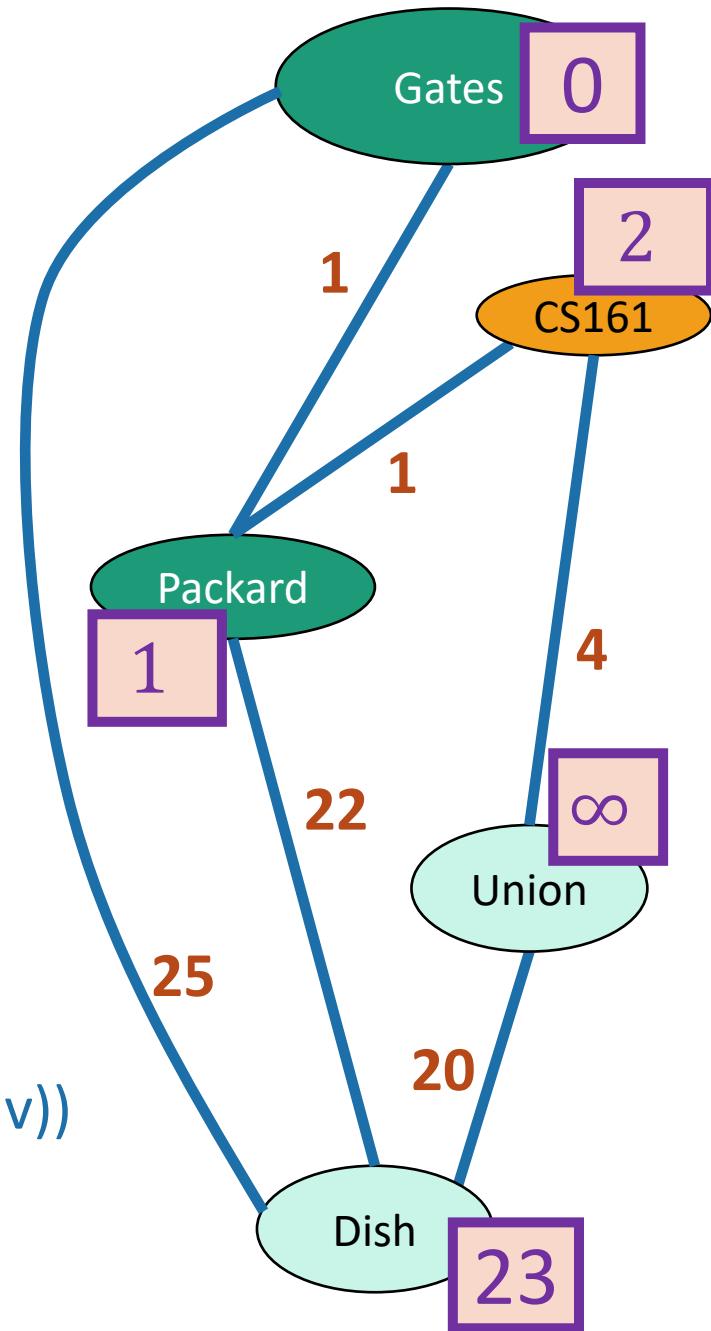


$x = d[v]$ is my best **over-estimate** for $\text{dist}(\text{Gates}, v)$.



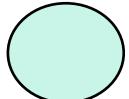
Current node u

- Pick the **not-sure** node u with the smallest estimate $d[u]$.
- Update all u 's neighbors v :
 - $d[v] = \min(d[v], d[u] + \text{edgeWeight}(u, v))$
- Mark u as **SURE**.
- Repeat

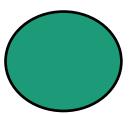


Dijkstra by example

How far is a node from Gates?



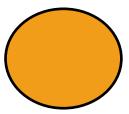
I'm not sure yet



I'm sure

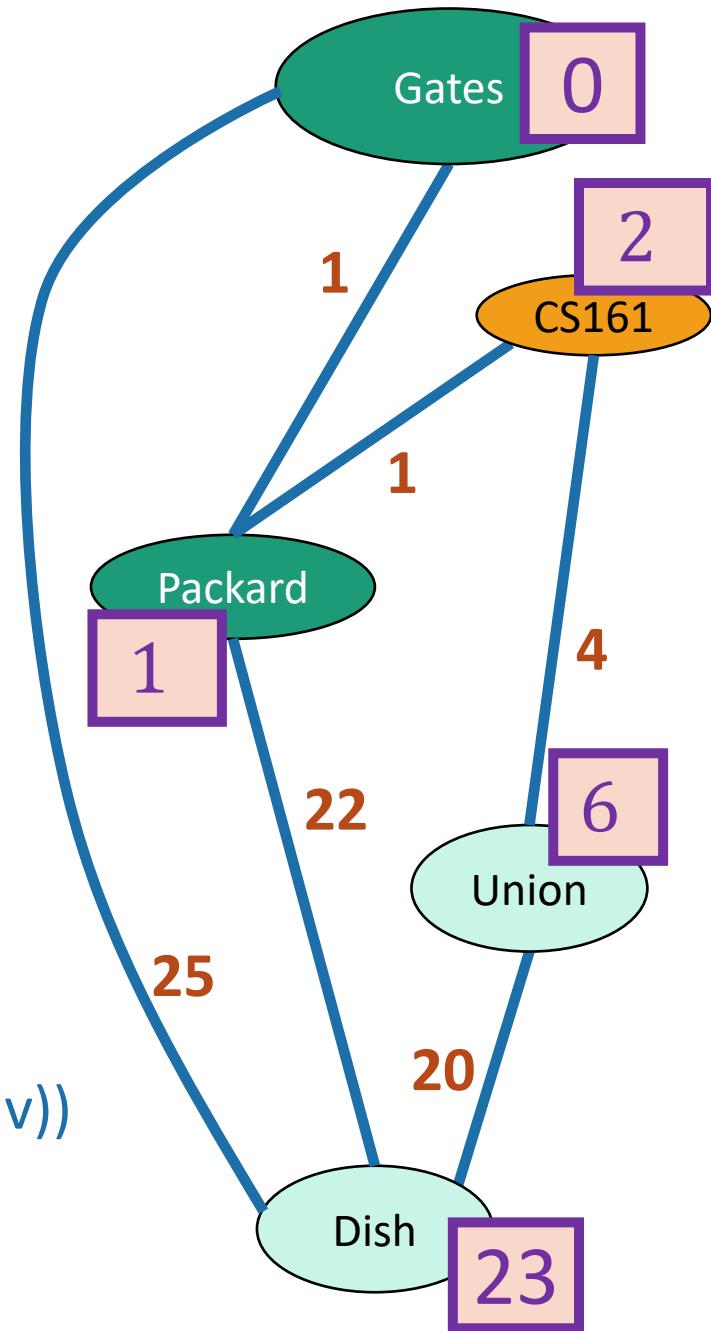


$x = d[v]$ is my best **over-estimate** for $\text{dist}(\text{Gates}, v)$.



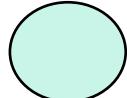
Current node u

- Pick the **not-sure** node u with the smallest estimate $d[u]$.
- Update all u 's neighbors v :
 - $d[v] = \min(d[v], d[u] + \text{edgeWeight}(u, v))$
- Mark u as **SURE**.
- Repeat

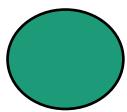


Dijkstra by example

How far is a node from Gates?



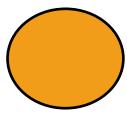
I'm not sure yet



I'm sure

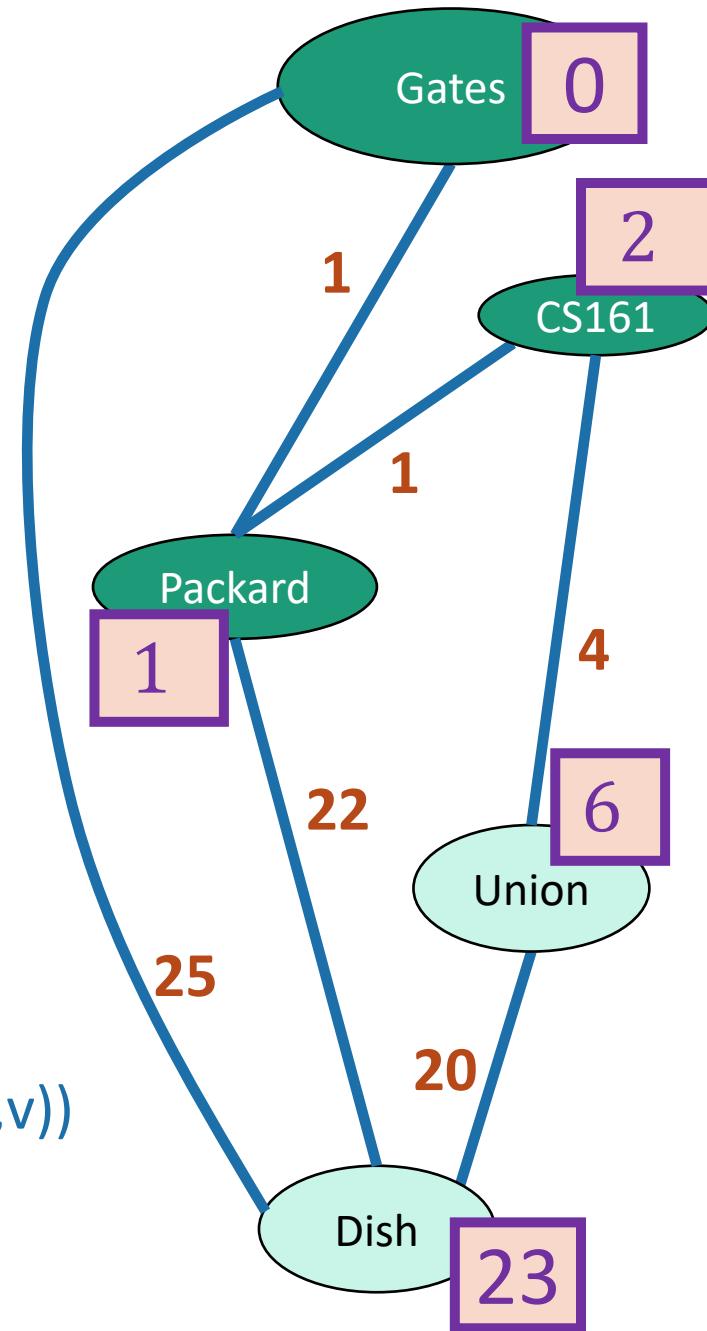


$x = d[v]$ is my best **over-estimate** for $\text{dist}(\text{Gates}, v)$.



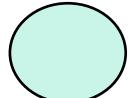
Current node u

- Pick the **not-sure** node u with the smallest estimate $d[u]$.
- Update all u 's neighbors v :
 - $d[v] = \min(d[v], d[u] + \text{edgeWeight}(u, v))$
- Mark u as **SURE**.
- Repeat

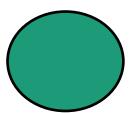


Dijkstra by example

How far is a node from Gates?



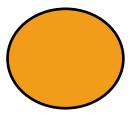
I'm not sure yet



I'm sure

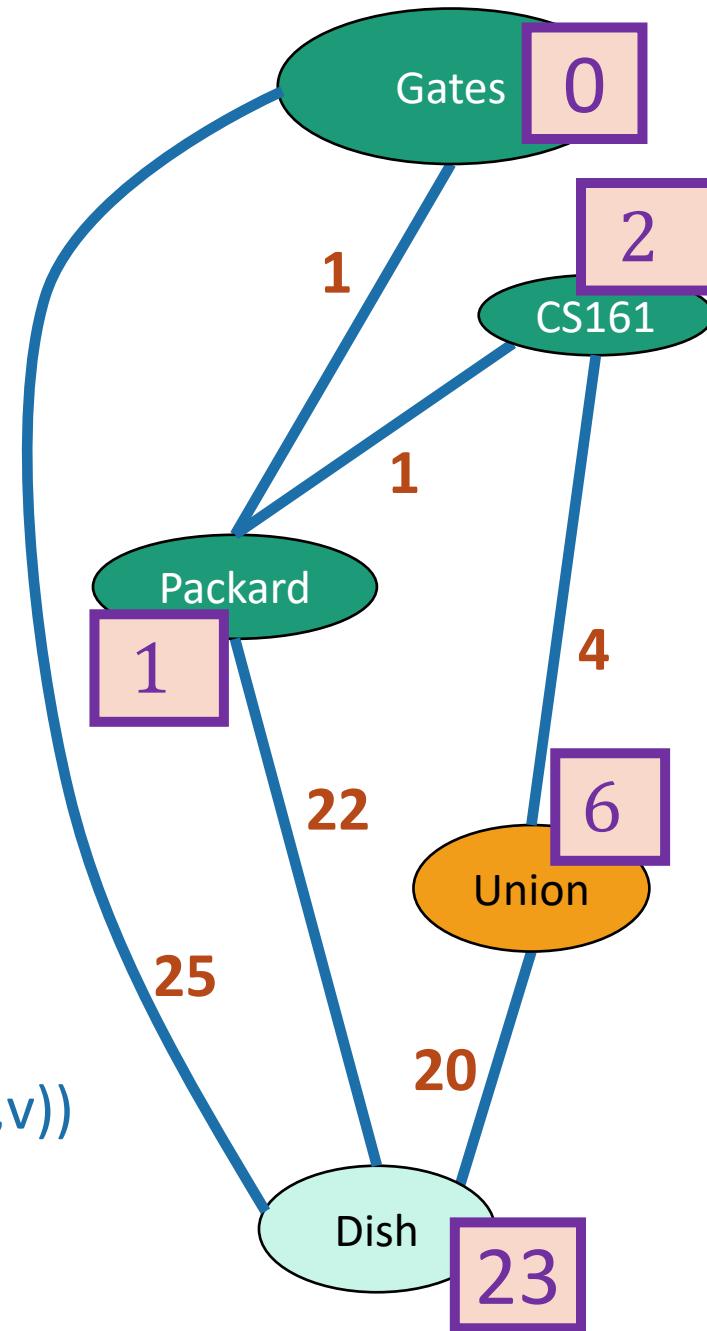


$x = d[v]$ is my best **over-estimate** for $\text{dist}(\text{Gates}, v)$.



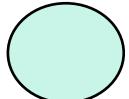
Current node u

- Pick the **not-sure** node u with the smallest estimate $d[u]$.
- Update all u's neighbors v:
 - $d[v] = \min(d[v] , d[u] + \text{edgeWeight}(u,v))$
- Mark u as **SURE**.
- Repeat

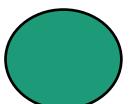


Dijkstra by example

How far is a node from Gates?



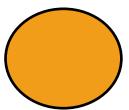
I'm not sure yet



I'm sure

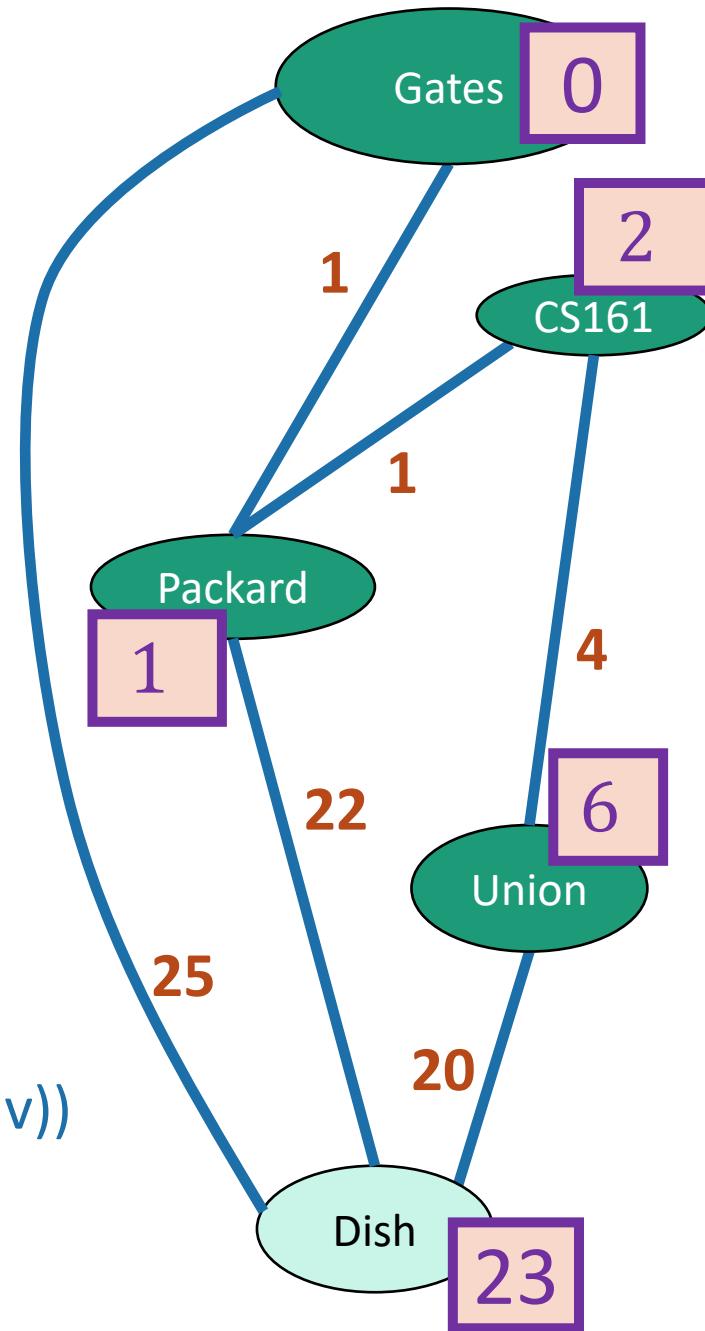


$x = d[v]$ is my best **over-estimate** for $\text{dist}(\text{Gates}, v)$.



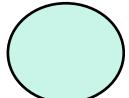
Current node u

- Pick the **not-sure** node u with the smallest estimate $d[u]$.
- Update all u's neighbors v:
 - $d[v] = \min(d[v] , d[u] + \text{edgeWeight}(u,v))$
- Mark u as **SURE**.
- Repeat

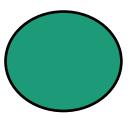


Dijkstra by example

How far is a node from Gates?



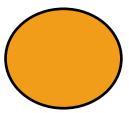
I'm not sure yet



I'm sure

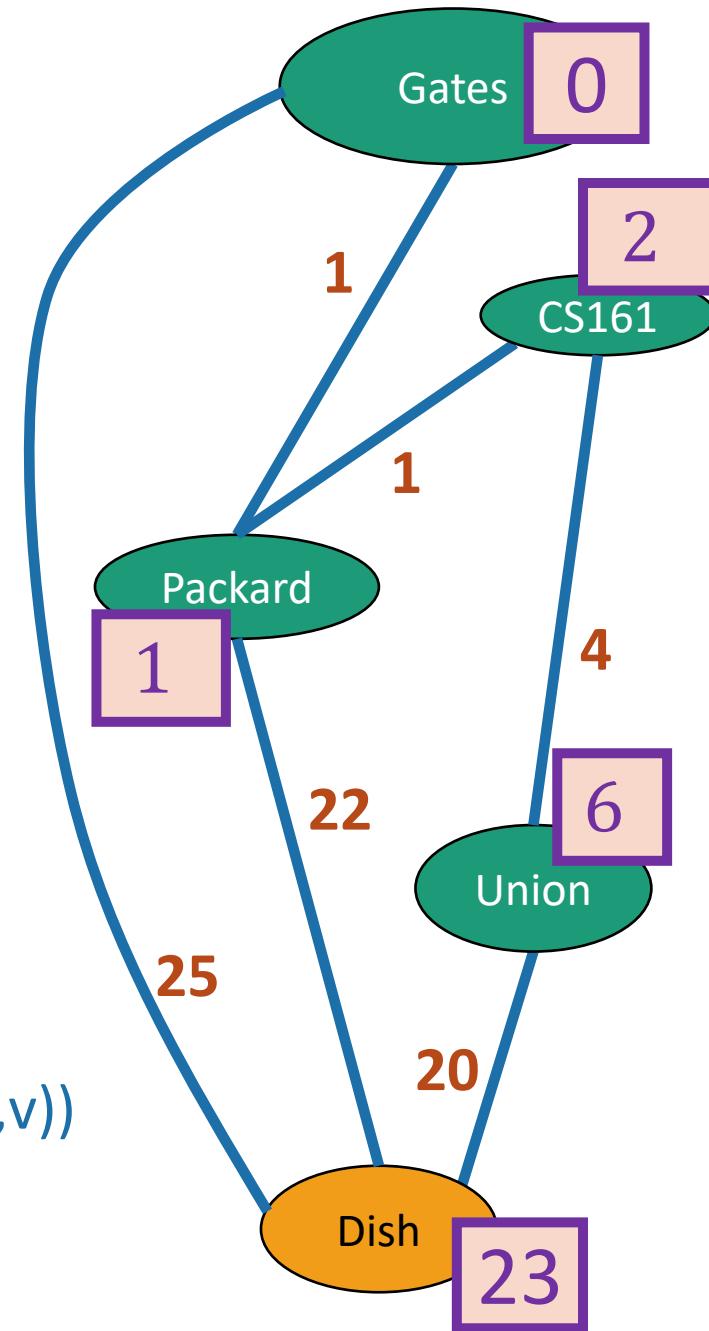


$x = d[v]$ is my best **over-estimate** for $\text{dist}(\text{Gates}, v)$.



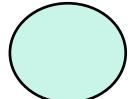
Current node u

- Pick the **not-sure** node u with the smallest estimate $d[u]$.
- Update all u 's neighbors v :
 - $d[v] = \min(d[v], d[u] + \text{edgeWeight}(u, v))$
- Mark u as **SURE**.
- Repeat

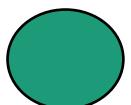


Dijkstra by example

How far is a node from Gates?



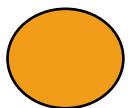
I'm not sure yet



I'm sure

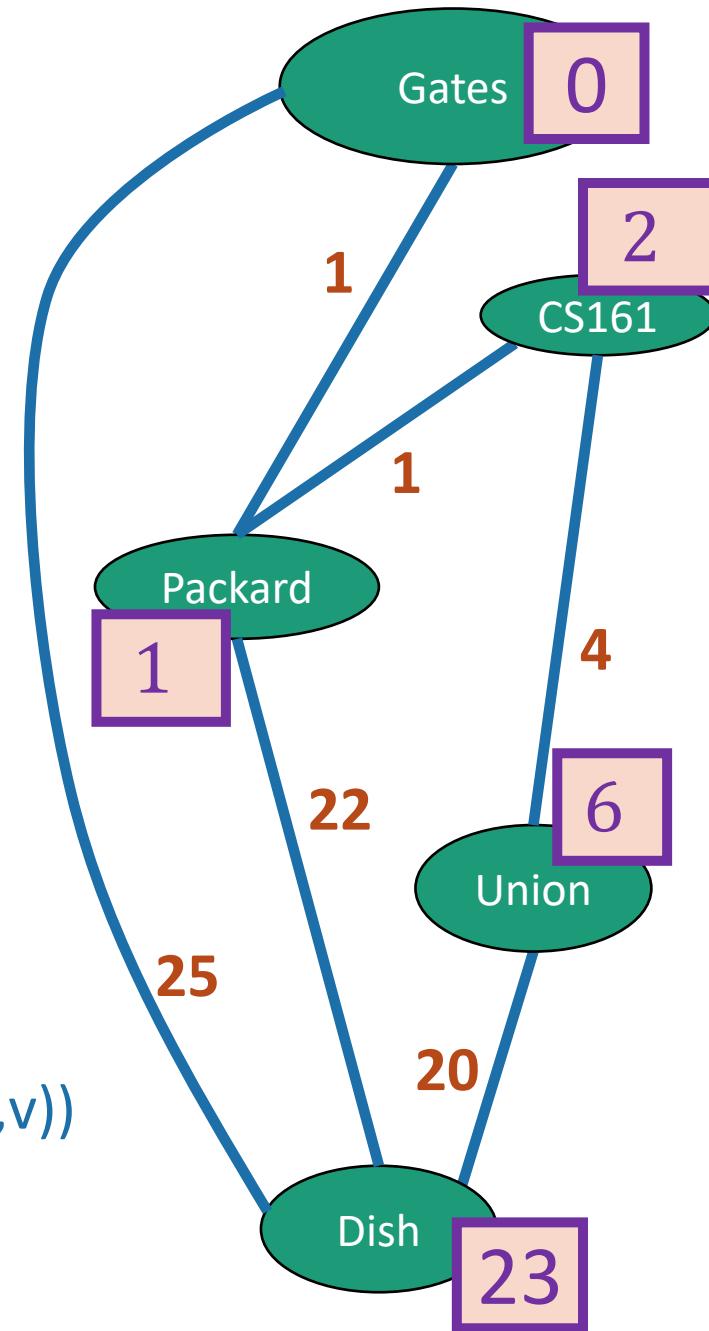


$x = d[v]$ is my best **over-estimate** for $\text{dist}(\text{Gates}, v)$.



Current node u

- Pick the **not-sure** node u with the smallest estimate $d[u]$.
- Update all u 's neighbors v :
 - $d[v] = \min(d[v], d[u] + \text{edgeWeight}(u, v))$
- Mark u as **SURE**.
- Repeat



Dijkstra's algorithm

Dijkstra(G, s):

- Set all vertices to **not-sure**
- $d[v] = \infty$ for all v in V
- $d[s] = 0$
- **While** there are **not-sure** nodes:
 - Pick the **not-sure** node u with the smallest estimate $d[u]$.
 - **For** v in $u.\text{neighbors}$:
 - $d[v] \leftarrow \min(d[v], d[u] + \text{edgeWeight}(u, v))$
 - Mark u as **sure**.
 - Now $d(s, v) = d[v]$

Running time?

Dijkstra(G, s):

- Set all vertices to **not-sure**
- $d[v] = \infty$ for all v in V
- $d[s] = 0$
- **While** there are **not-sure** nodes:
 - Pick the **not-sure** node u with the smallest estimate $d[u]$.
 - **For** v in $u.\text{neighbors}$:
 - $d[v] \leftarrow \min(d[v], d[u] + \text{edgeWeight}(u, v))$
 - Mark u as **sure**.
 - Now $\text{dist}(s, v) = d[v]$
- n iterations (one per vertex)
- How long does one iteration take?

Depends on how we implement it...

We need a data structure that:

- Stores unsure vertices v
- Keeps track of $d[v]$
- Can find u with minimum $d[u]$
 - `findMin()`
- Can remove that u
 - `removeMin(u)`
- Can update (decrease) $d[v]$
 - `updateKey(v, d)`

Just the inner loop:

- Pick the **not-sure** node u with the smallest estimate $d[u]$.
- Update all u 's neighbors v :
 - $d[v] \leftarrow \min(d[v], d[u] + \text{edgeWeight}(u,v))$
- Mark u as **sure**.

Total running time is big-oh of:

$$\sum_{u \in V} \left(T(\text{findMin}) + \left(\sum_{v \in u.\text{neighbors}} T(\text{updateKey}) \right) + T(\text{removeMin}) \right)$$

$$= n(T(\text{findMin}) + T(\text{removeMin})) + m T(\text{updateKey})$$

If we use an array

- $T(\text{findMin}) = O(n)$
- $T(\text{removeMin}) = O(n)$
- $T(\text{updateKey}) = O(1)$
- Running time of Dijkstra
 - $= O(n(T(\text{findMin}) + T(\text{removeMin})) + m T(\text{updateKey}))$
 - $= O(n^2) + O(m)$
 - $= O(n^2)$

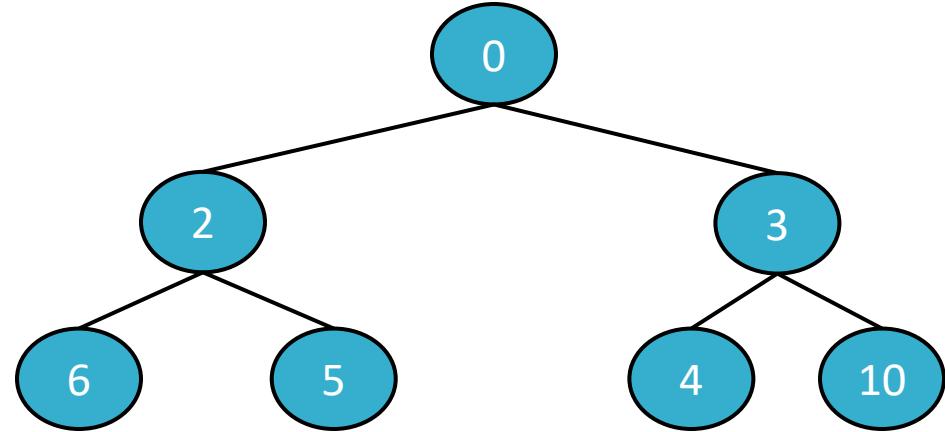
If we use a red-black tree

- $T(\text{findMin}) = O(\log(n))$
- $T(\text{removeMin}) = O(\log(n))$
- $T(\text{updateKey}) = O(\log(n))$
- Running time of Dijkstra
 - $= O(n(T(\text{findMin}) + T(\text{removeMin})) + m T(\text{updateKey}))$
 - $= O(n\log(n)) + O(m\log(n))$
 - $= O((n + m)\log(n))$

Better than an array if the graph is sparse!
aka if m is much smaller than n^2

Heaps support these operations

- T(findMin)
- T(removeMin)
- T(updateKey)



- A **heap** is a tree-based data structure that has the property that **every node has a smaller key than its children.**

Many heap implementations

Nice chart on Wikipedia:

Operation	Binary ^[7]	Leftist	Binomial ^[7]	Fibonacci ^{[7][8]}	Pairing ^[9]	Brodal ^{[10][b]}	Rank-pairing ^[12]	Strict Fibonacci ^[13]
find-min	$\Theta(1)$	$\Theta(1)$	$\Theta(\log n)$	$\Theta(1)$	$\Theta(1)$	$\Theta(1)$	$\Theta(1)$	$\Theta(1)$
delete-min	$\Theta(\log n)$	$\Theta(\log n)$	$\Theta(\log n)$	$O(\log n)^{[c]}$	$O(\log n)^{[c]}$	$O(\log n)$	$O(\log n)^{[c]}$	$O(\log n)$
insert	$O(\log n)$	$\Theta(\log n)$	$\Theta(1)^{[c]}$	$\Theta(1)$	$\Theta(1)$	$\Theta(1)$	$\Theta(1)$	$\Theta(1)$
decrease-key	$\Theta(\log n)$	$\Theta(n)$	$\Theta(\log n)$	$\Theta(1)^{[c]}$	$o(\log n)^{[c][d]}$	$\Theta(1)$	$\Theta(1)^{[c]}$	$\Theta(1)$
merge	$\Theta(n)$	$\Theta(\log n)$	$O(\log n)^{[e]}$	$\Theta(1)$	$\Theta(1)$	$\Theta(1)$	$\Theta(1)$	$\Theta(1)$

Say we use a Fibonacci Heap

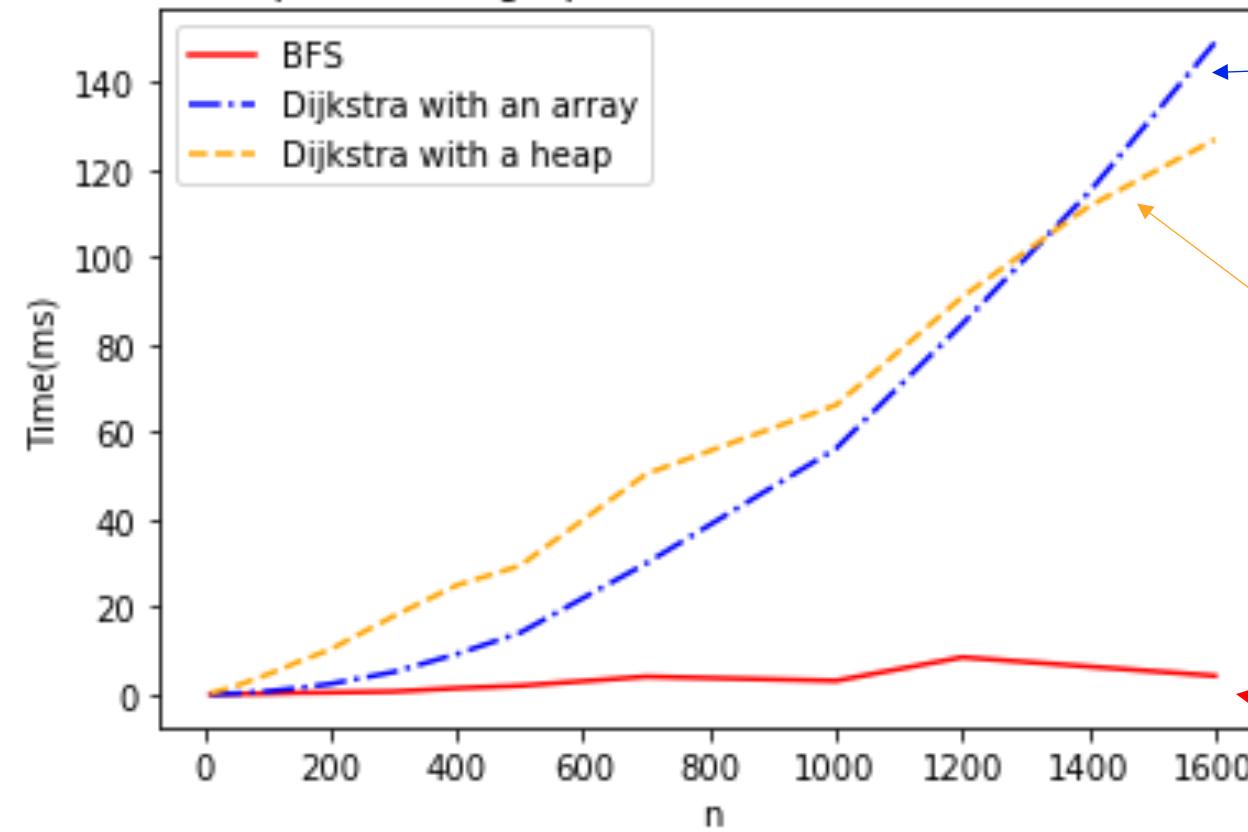
- $T(\text{findMin}) = O(1)$ (amortized time*)
- $T(\text{removeMin}) = O(\log(n))$ (amortized time*)
- $T(\text{updateKey}) = O(1)$ (amortized time*)
- Running time of Dijkstra
 - $= O(n(T(\text{findMin}) + T(\text{removeMin})) + m T(\text{updateKey}))$
 - $= O(n \log(n) + m)$ (amortized time)

*This means that any sequence of d `removeMin` calls takes time at most $O(d \log(n))$.
But a few of the d may take longer than $O(\log(n))$ and some may take less time..

See IPython Notebook for Lecture 11
The heap is implemented using `heapdict`

In practice

Shortest paths on a graph with n vertices and about $5n$ edges



Dijkstra using a Python list to keep track of vertices has quadratic runtime.

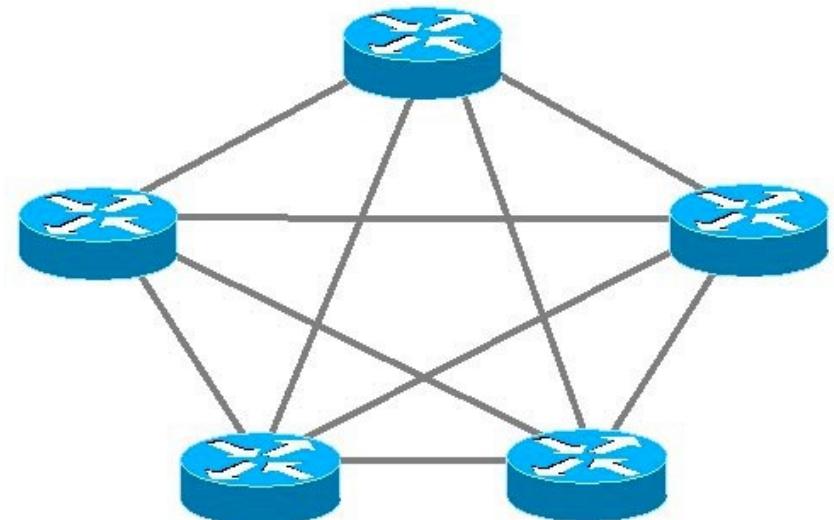
Dijkstra using a heap looks a bit more linear (actually $n\log(n)$)

BFS is really fast by comparison! But it doesn't work on weighted graphs.

Dijkstra is used in practice

- eg, [OSPF \(Open Shortest Path First\)](#), a routing protocol for IP networks, uses Dijkstra.

But there are
some things it's
not so good at.

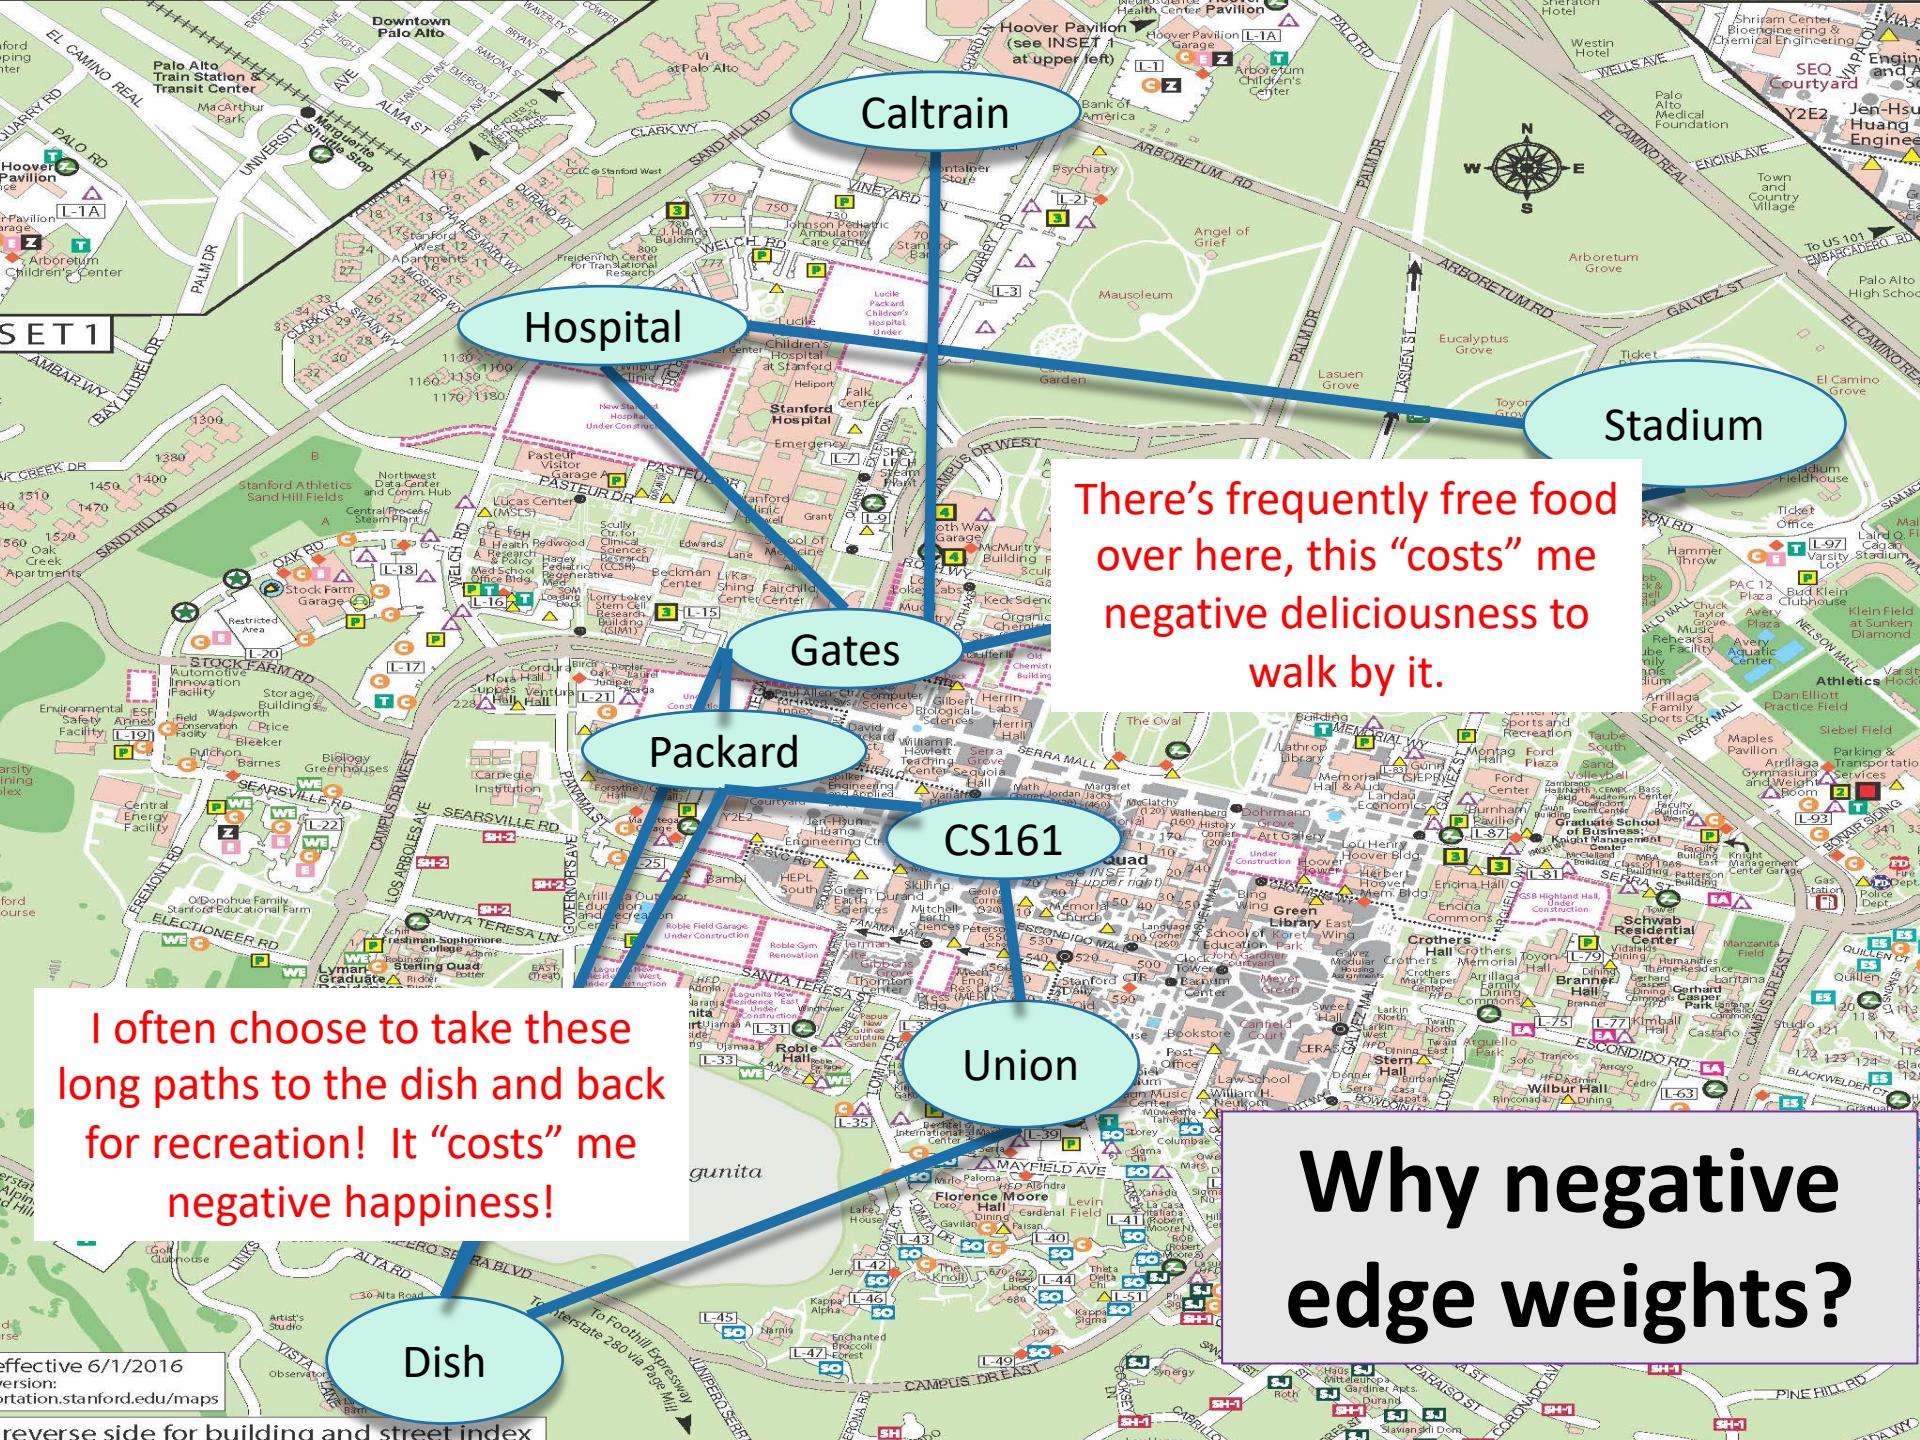


Dijkstra Drawbacks

- Needs **non-negative edge weights**.
- If the weights change, we need to re-run the whole thing.
 - in OSPF, a vertex broadcasts any changes to the network, and then every vertex re-runs Dijkstra's algorithm from scratch.

Bellman-Ford algorithm

- (-) Slower than Dijkstra's algorithm
- (+) Can handle negative edge weights.
- (+) Allows for some flexibility if the weights change.
 - We'll see what this means later



I often choose to take these long paths to the dish and back for recreation! It “costs” me negative happiness!

There's frequently free food over here, this “costs” me negative deliciousness to walk by it.

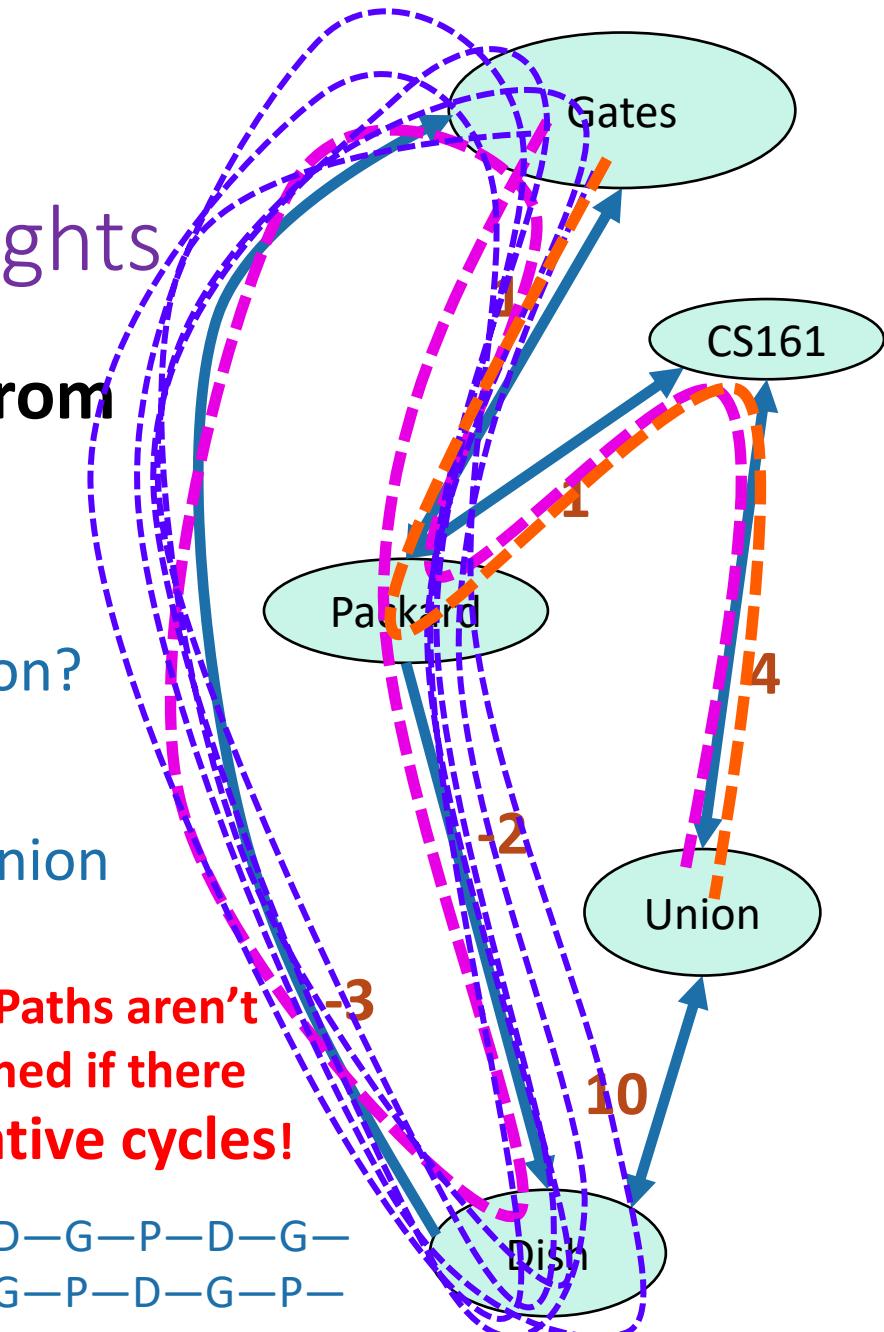
Why negative edge weights?

One problem with negative edge weights

- What is the shortest path from Gates to the Union?
- Should it still be Gates—Packard—CS161—Union?
- But what about
 - G—P—D—G—P—CS161—Union
- That costs
 - $1-2-3+1+1+4 = 2.$
- And why not

G—P—D—G—P—D—G—P—D—G—P—D—G—
P—D—G—P—D—G—P—D—G—P—D—G—P—D—G—
D—G—P—D—G—P—D—G—P—D—G—P—D—etc....

Shortest Paths aren't
well-defined if there
are negative cycles!



Bellman-Ford algorithm

Bellman-Ford(G, s):

- $d[v] = \infty$ for all v in V
 - $d[s] = 0$
 - **For** $i=0, \dots, n-1$:
 - **For** u in V :
 - **For** v in $u.\text{neighbors}$:
 - $d[v] \leftarrow \min(d[v], d[u] + \text{edgeWeight}(u,v))$
- Instead of picking u cleverly,
just update for all of the u 's.

Compare to Dijkstra:

- **While** there are **not-sure** nodes:
 - Pick the **not-sure** node u with the smallest estimate $d[u]$.
 - **For** v in $u.\text{neighbors}$:
 - $d[v] \leftarrow \min(d[v], d[u] + \text{edgeWeight}(u,v))$
 - Mark u as **sure**.

For pedagogical reasons which we will see next week

- We are actually going to change this to be dumber.
- Keep n arrays: $d^{(0)}, d^{(1)}, \dots, d^{(n-1)}$

Bellman-Ford*(G,s):

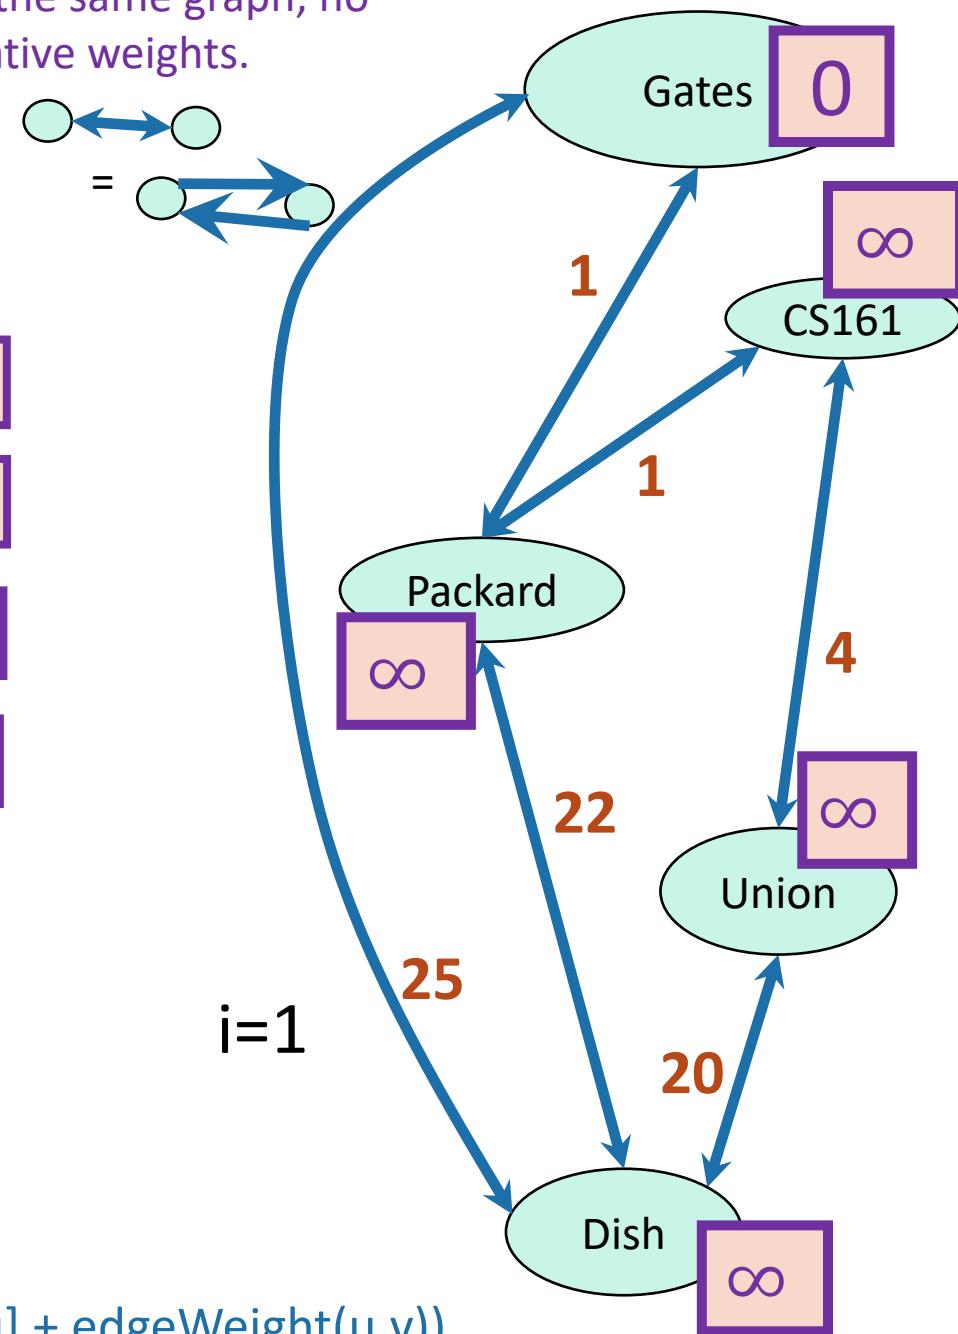
- $d^{(0)}[v] = \infty$ for all v in V
- $d^{(0)}[s] = 0$
- **For** $i=0, \dots, n-1$:
 - **For** u in V :
 - **For** v in $u.\text{neighbors}$:
 - $d^{(i+1)}[v] \leftarrow \min(d^{(i)}[v], d^{(i)}[u] + \text{edgeWeight}(u,v))$
 - Then $\text{dist}(s,v) = d^{(n-1)}[v]$

Bellman-Ford

How far is a node from Gates?

	Gates	Packard	CS161	Union	Dish
$d^{(0)}$	0	∞	∞	∞	∞
$d^{(1)}$	∞	∞	∞	∞	∞
$d^{(2)}$	∞	∞	∞	∞	∞
$d^{(3)}$	∞	∞	∞	∞	∞
$d^{(4)}$	∞	∞	∞	∞	∞

Start with the same graph, no negative weights.



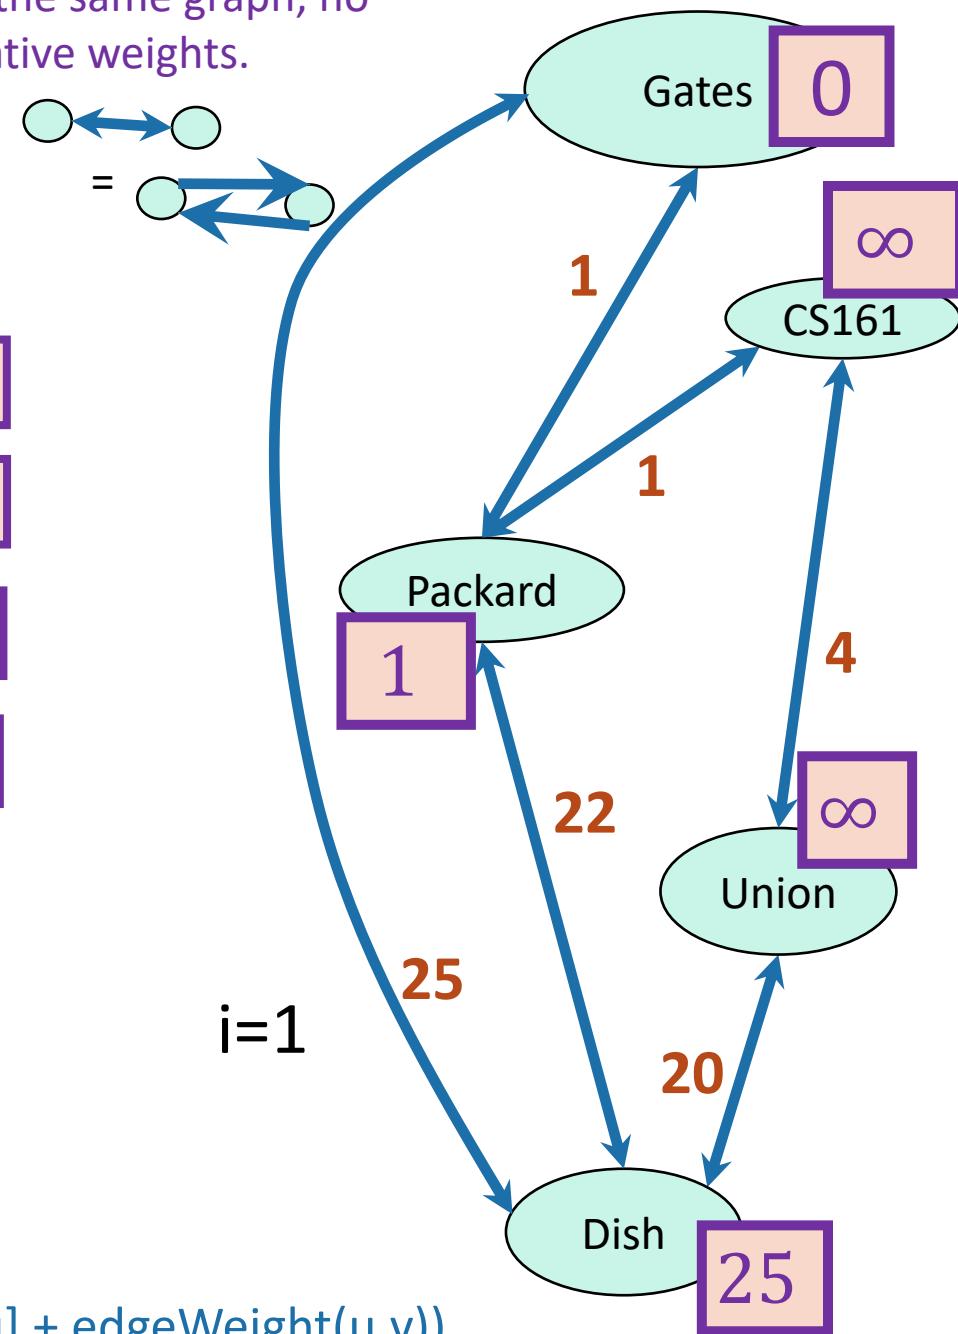
- For $i=0, \dots, n-2$:
 - For u in V :
 - For v in $u.\text{neighbors}$:
 - $d^{(i+1)}[v] \leftarrow \min(d^{(i)}[v], d^{(i)}[u] + \text{edgeWeight}(u,v))$

Bellman-Ford

How far is a node from Gates?

	Gates	Packard	CS161	Union	Dish
$d^{(0)}$	0	∞	∞	∞	∞
$d^{(1)}$	0	1	∞	∞	25
$d^{(2)}$					
$d^{(3)}$					
$d^{(4)}$					

Start with the same graph, no negative weights.



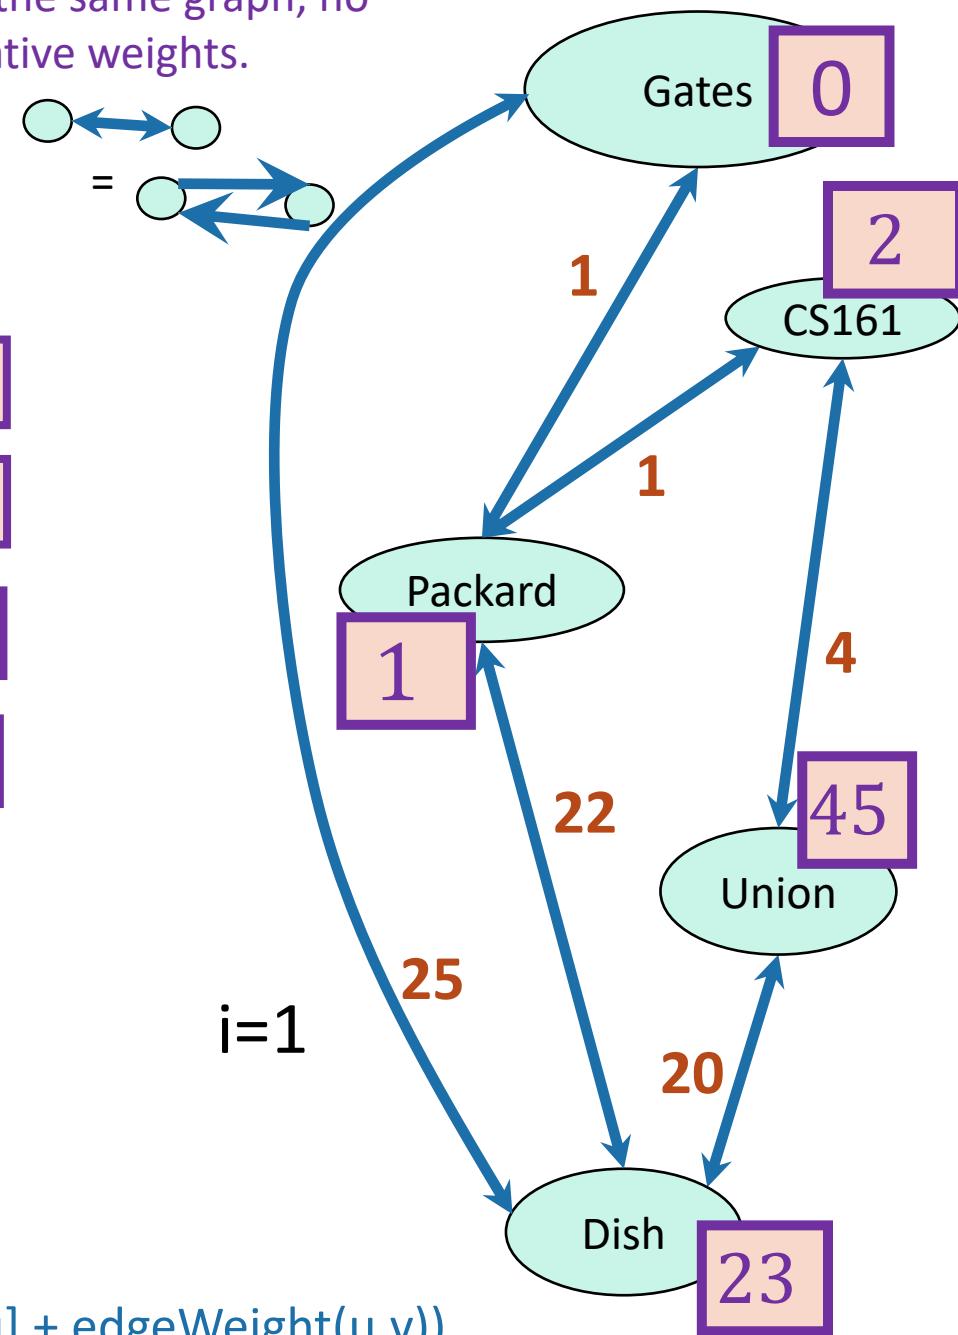
- For $i=0, \dots, n-2$:
 - For u in V :
 - For v in $u.\text{neighbors}$:
 - $d^{(i+1)}[v] \leftarrow \min(d^{(i)}[v], d^{(i)}[u] + \text{edgeWeight}(u,v))$

Bellman-Ford

How far is a node from Gates?

	Gates	Packard	CS161	Union	Dish
$d^{(0)}$	0	∞	∞	∞	∞
$d^{(1)}$	0	1	∞	∞	25
$d^{(2)}$	0	1	2	45	23
$d^{(3)}$					
$d^{(4)}$					

Start with the same graph, no negative weights.



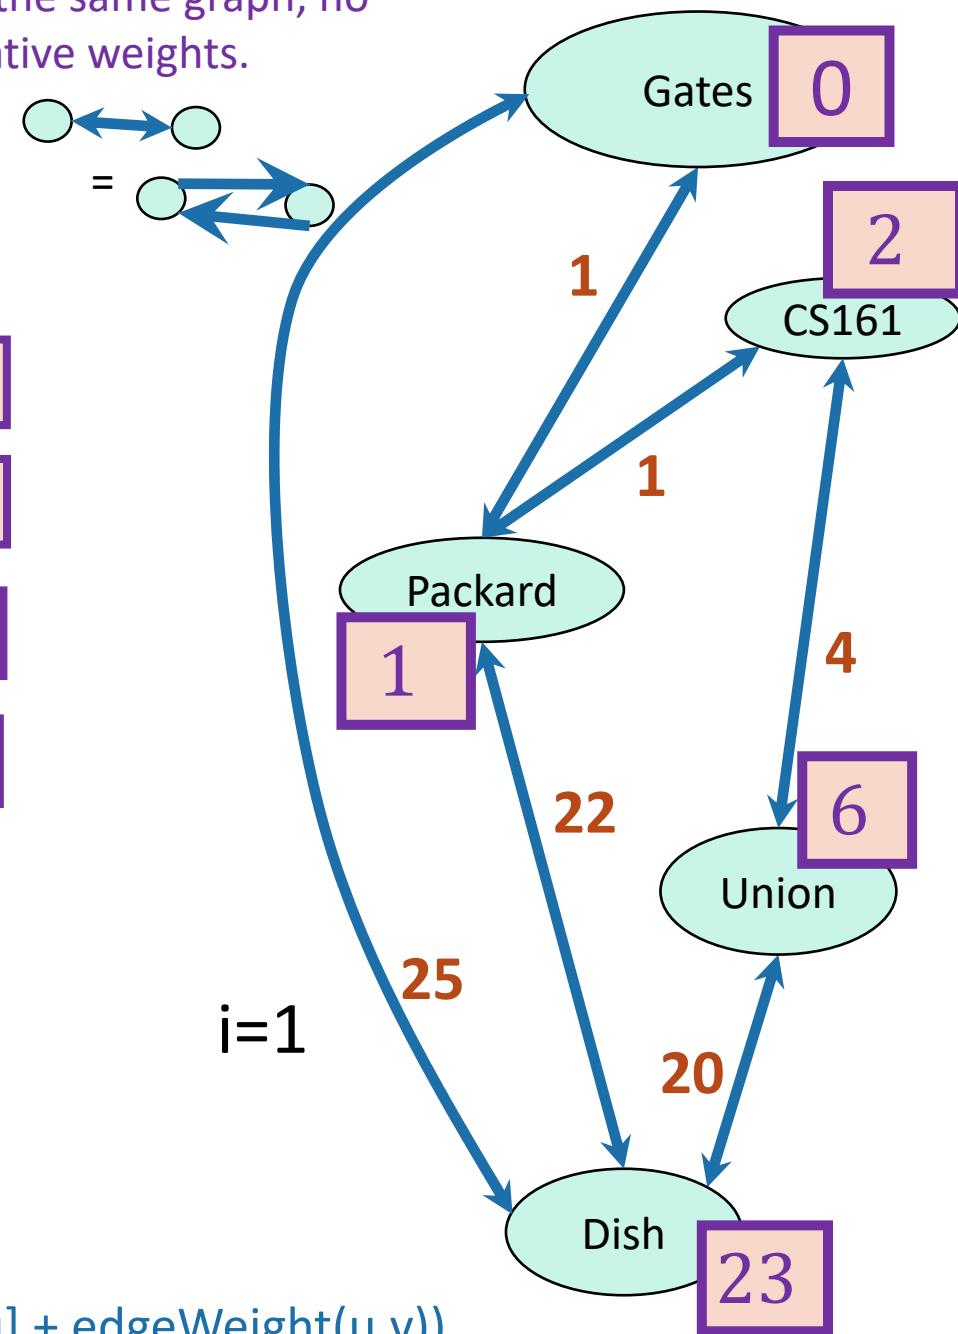
- For $i=0, \dots, n-2$:
 - For u in V :
 - For v in $u.\text{neighbors}$:
 - $d^{(i+1)}[v] \leftarrow \min(d^{(i)}[v], d^{(i)}[u] + \text{edgeWeight}(u,v))$

Bellman-Ford

How far is a node from Gates?

	Gates	Packard	CS161	Union	Dish
$d^{(0)}$	0	∞	∞	∞	∞
$d^{(1)}$	0	1	∞	∞	25
$d^{(2)}$	0	1	2	45	23
$d^{(3)}$	0	1	2	6	23
$d^{(4)}$					

Start with the same graph, no negative weights.



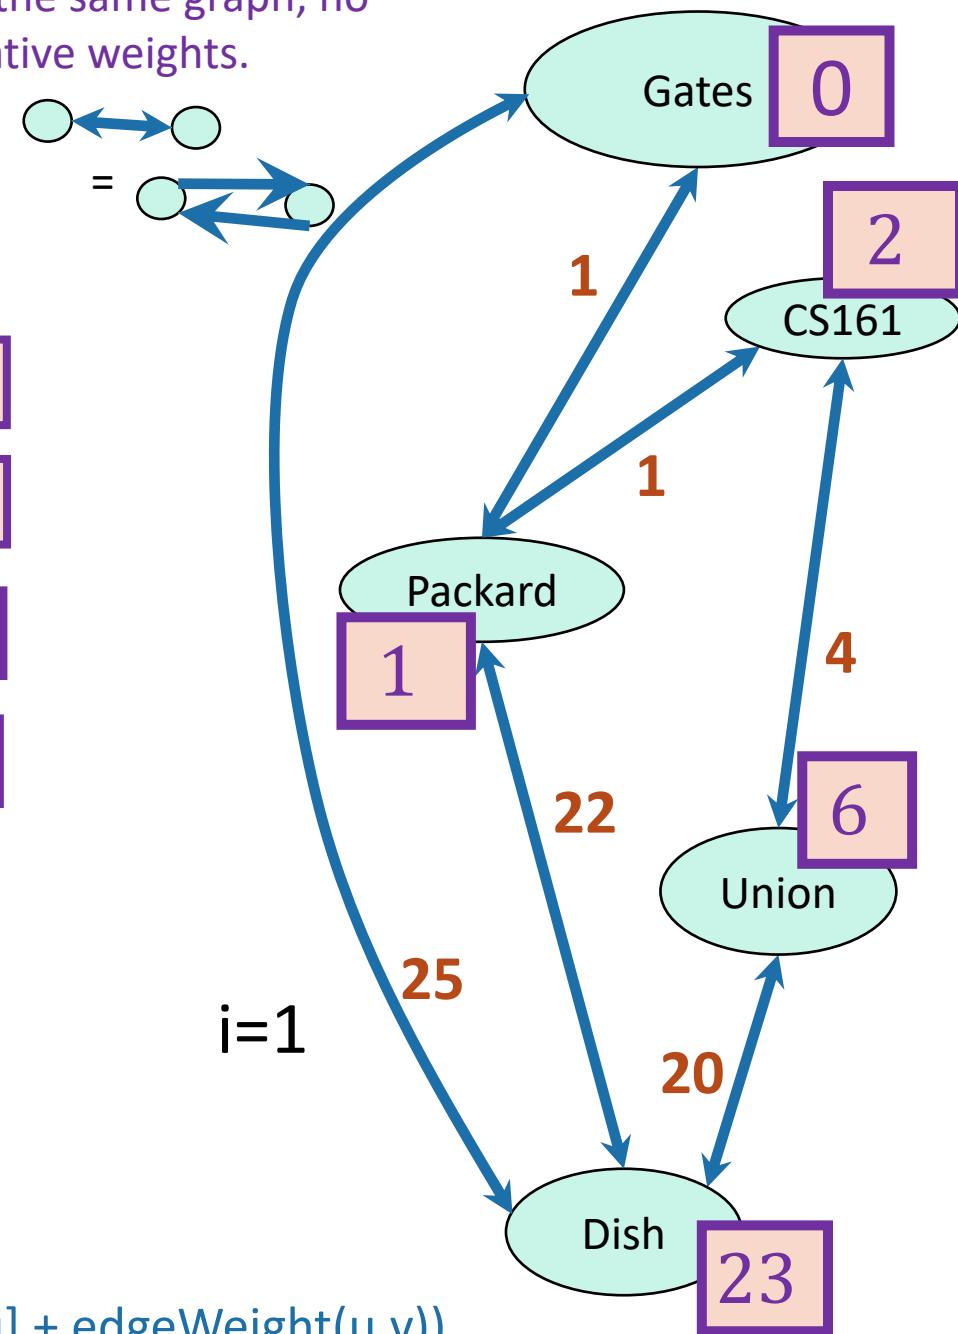
- For $i=0, \dots, n-2$:
 - For u in V :
 - For v in $u.\text{neighbors}$:
 - $d^{(i+1)}[v] \leftarrow \min(d^{(i)}[v], d^{(i)}[u] + \text{edgeWeight}(u,v))$

Bellman-Ford

How far is a node from Gates?

	Gates	Packard	CS161	Union	Dish
$d^{(0)}$	0	∞	∞	∞	∞
$d^{(1)}$	0	1	∞	∞	25
$d^{(2)}$	0	1	2	45	23
$d^{(3)}$	0	1	2	6	23
$d^{(4)}$	0	1	2	6	23

Start with the same graph, no negative weights.



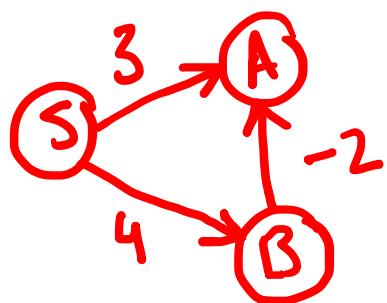
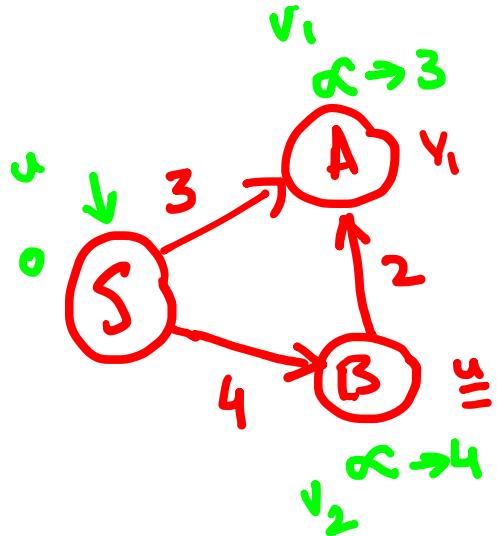
- For $i=0, \dots, n-2$:
 - For u in V :
 - For v in $u.\text{neighbors}$:
 - $d^{(i+1)}[v] \leftarrow \min(d^{(i)}[v], d^{(i)}[u] + \text{edgeWeight}(u,v))$

S	S	A	B
0	0	∞	∞
0	3	4	

* $\hookrightarrow d[v] = \min(d[v], d[u] + w(u, v))$

$\infty, 0+3$

$3, 4+2$



This seems much slower than Dijkstra

- And it is:

Running time $O(mn)$

- However, it's also more flexible in a few ways.
 - Can handle negative edges
 - If we keep on doing these iterations, then changes in the network will propagate through.

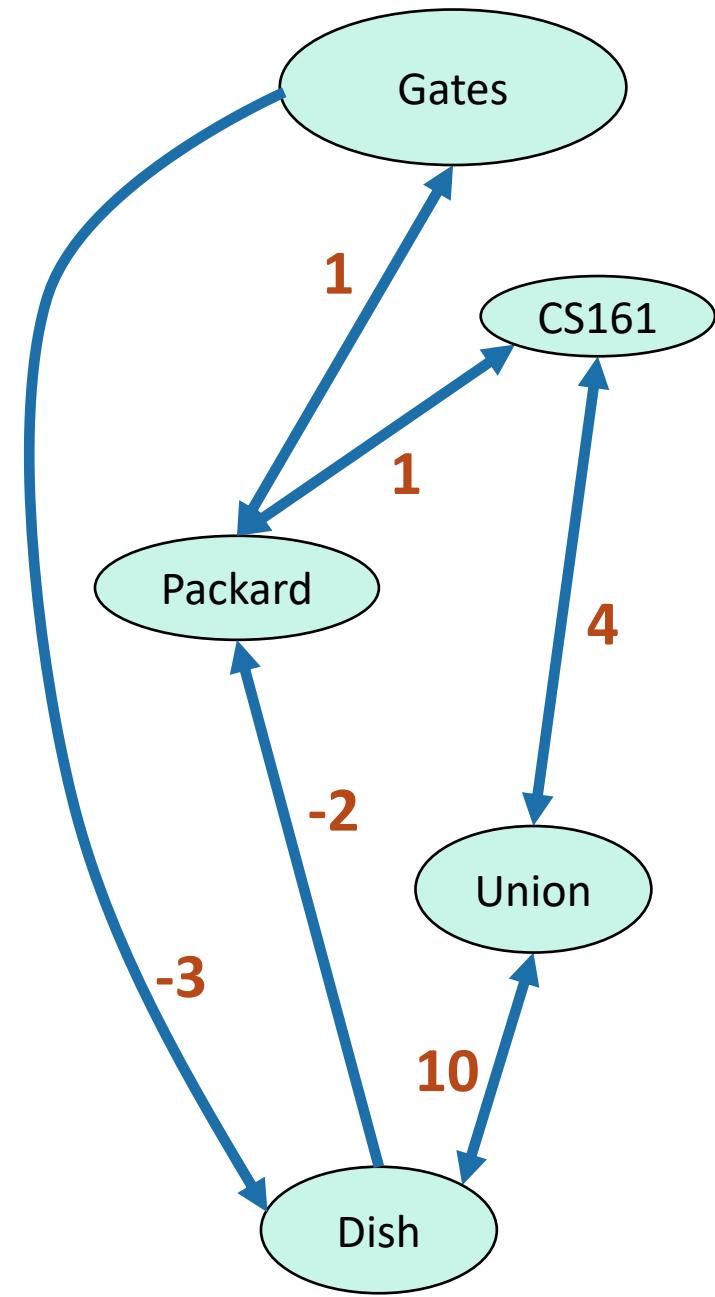
- **For** $i=0, \dots, n-1$:
 - **For** u in V :
 - **For** v in $u.\text{neighbors}$:
 - $d^{(i+1)}[v] \leftarrow \min(d^{(i)}[v], d^{(i)}[u] + \text{edgeWeight}(u,v))$

Negative edge weights

	Gates	Packard	CS161	Union	Dish
$d^{(0)}$	0	∞	∞	∞	∞
$d^{(1)}$	0	1	∞	∞	-3
$d^{(2)}$	0	-5	2	7	-3
$d^{(3)}$	-4	-5	-4	6	-3

This is not looking good!

- For $i=0, \dots, n-2$:
 - For u in V :
 - For v in $u.\text{neighbors}$:
 - $d^{(i+1)}[v] \leftarrow \min(d^{(i)}[v], d^{(i)}[u] + \text{edgeWeight}(u,v))$



Negative edge weights

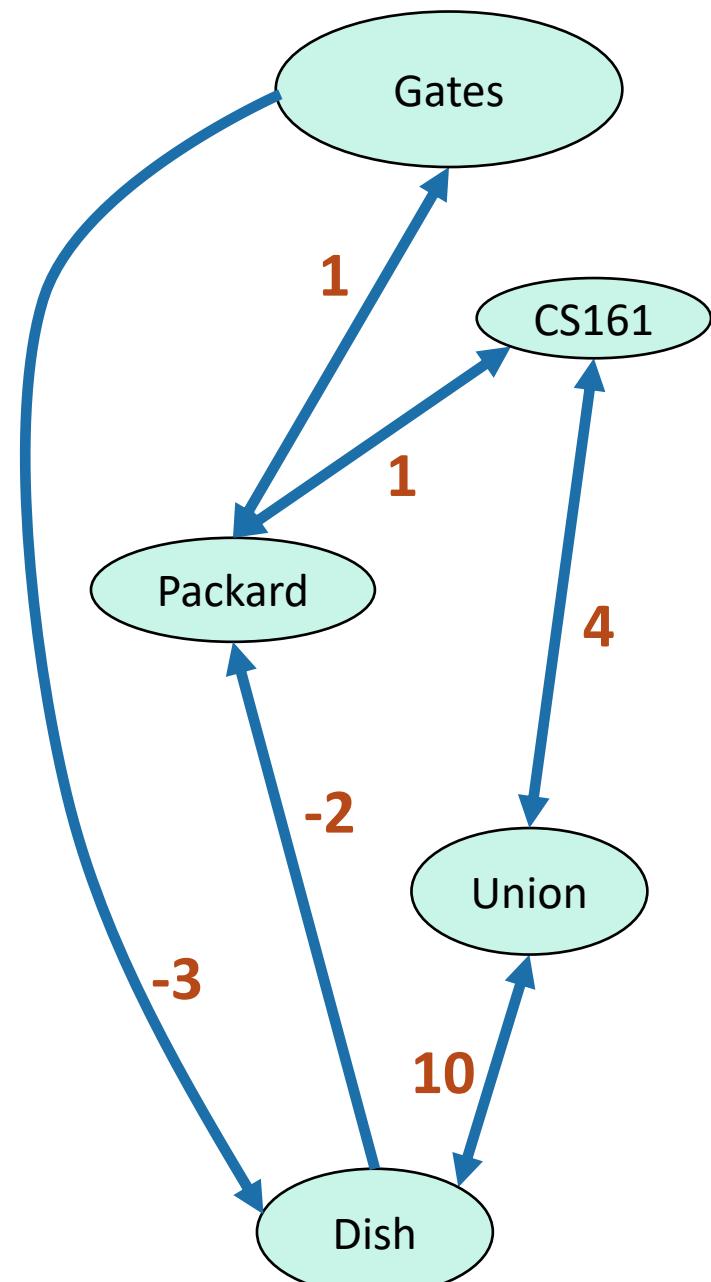
	Gates	Packard	CS161	Union	Dish
$d^{(0)}$	0	∞	∞	∞	∞
$d^{(1)}$	0	1	∞	∞	-3
$d^{(2)}$	0	-5	2	7	-3
$d^{(3)}$	-4	-5	-4	6	-3
$d^{(4)}$	-4	-5	-4	6	-7

But we can tell that it's not looking good:

$d^{(5)}$	-4	-9	-4	3	-7

Some stuff changed!

- For $i=0, \dots, n-1$:
- For u in V :
 - For v in $u.\text{neighbors}$:
 - $d^{(i+1)}[v] \leftarrow \min(d^{(i)}[v], d^{(i)}[u] + \text{edgeWeight}(u,v))$



How Bellman-Ford deals with negative cycles

- If there are no negative cycles:
 - Everything works as it should.
 - The algorithm stabilizes after $n-1$ rounds.
 - Note: Negative *edges* are okay!!
- If there are negative cycles:
 - Not everything works as it should...
 - Note: it couldn't possibly work, since shortest paths aren't well-defined if there are negative cycles.
 - The $d[v]$ values will keep changing.
- Solution:
 - Go one round more and see if things change.

Bellman-Ford algorithm

Bellman-Ford*(G,s):

- $d^{(0)}[v] = \infty$ for all v in V
- $d^{(0)}[s] = 0$
- **For** $i=0, \dots, n-1$:
 - **For** u in V :
 - **For** v in $u.\text{neighbors}$:
 - $d^{(i+1)}[v] \leftarrow \min(d^{(i)}[v], d^{(i)}[u] + \text{edgeWeight}(u,v))$
 - If $d^{(n-1)} \neq d^{(n)}$:
 - **Return NEGATIVE CYCLE** ☹
 - Otherwise, $\text{dist}(s,v) = d^{(n-1)}[v]$

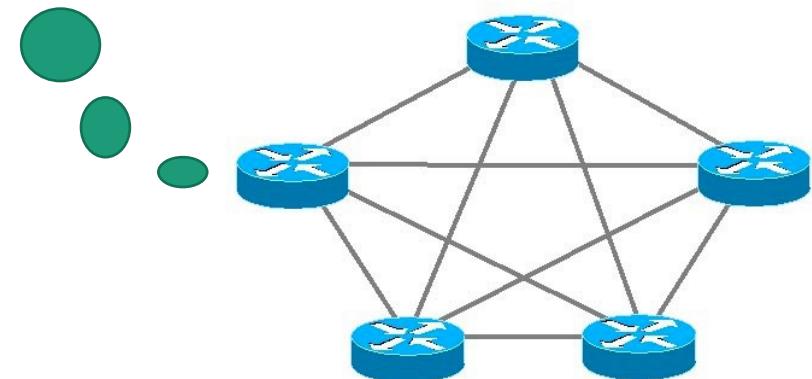
What have we learned?

- The Bellman-Ford algorithm:
 - Finds shortest paths in weighted graphs with negative edge weights
 - runs in time $O(nm)$ on a graph G with n vertices and m edges.
- If there are no negative cycles in G :
 - the BF algorithm terminates with $d^{(n-1)}[v] = d(s,v)$.
- If there are negative cycles in G :
 - the BF algorithm returns **negative cycle**.

Bellman-Ford is also used in practice.

- eg, **Routing Information Protocol (RIP)** uses something like Bellman-Ford.
 - Older protocol, not used as much anymore.
- Each router keeps a **table** of distances to every other router.
- Periodically we do a Bellman-Ford update.
- This means that if there are changes in the network, this will propagate. (maybe slowly...)

Destination	Cost to get there	Send to whom?
172.16.1.0	34	172.16.1.1
10.20.40.1	10	192.168.1.2
10.155.120.1	9	10.13.50.0



Recap: shortest paths

- **BFS:**
 - (+) $O(n+m)$
 - (-) only unweighted graphs
- **Dijkstra's algorithm:**
 - (+) weighted graphs
 - (+) $O(n\log(n) + m)$ if you implement it right.
 - (-) no negative edge weights
 - (-) very “centralized” (need to keep track of all the vertices to know which to update).
- **The Bellman-Ford algorithm:**
 - (+) weighted graphs, even with negative weights
 - (+) can be done in a distributed fashion, every vertex using only information from its neighbors.
 - (-) $O(nm)$

Bellman-Ford is an example of...

Dynamic Programming!

Today:

- Example of Dynamic programming:
 - Fibonacci numbers.
 - (And Bellman-Ford)
- What is dynamic programming, exactly?
 - And why is it called “dynamic programming”?
- Another example: Floyd-Warshall algorithm
 - An “all-pairs” shortest path algorithm

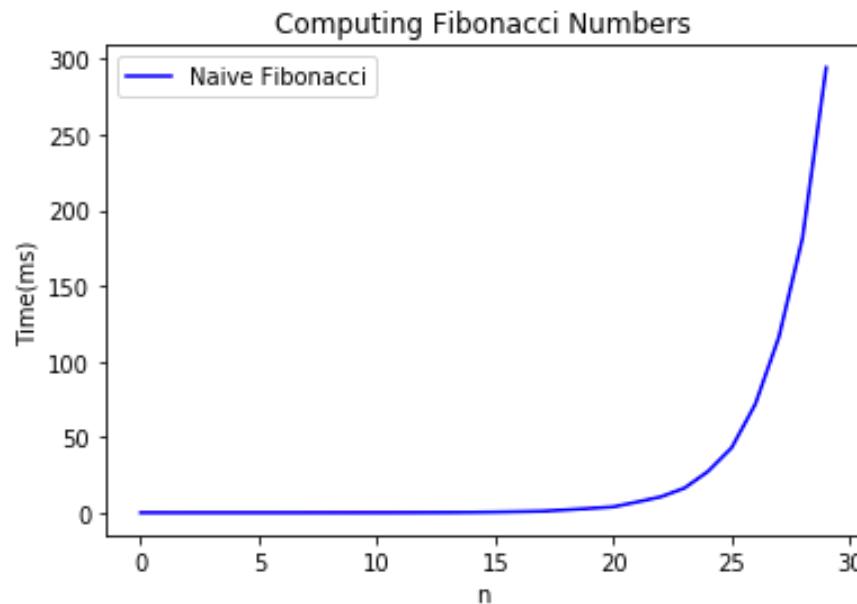


How not to compute Fibonacci Numbers

- Definition:
 - $F(n) = F(n-1) + F(n-2)$, with $F(0) = F(1) = 1$.
 - The first several are:
 $1, 1, 2, 3, 5, 8, 13, 21, 34, 55, 89, 144, \dots$
- Question:
 - Given n , what is $F(n)$?

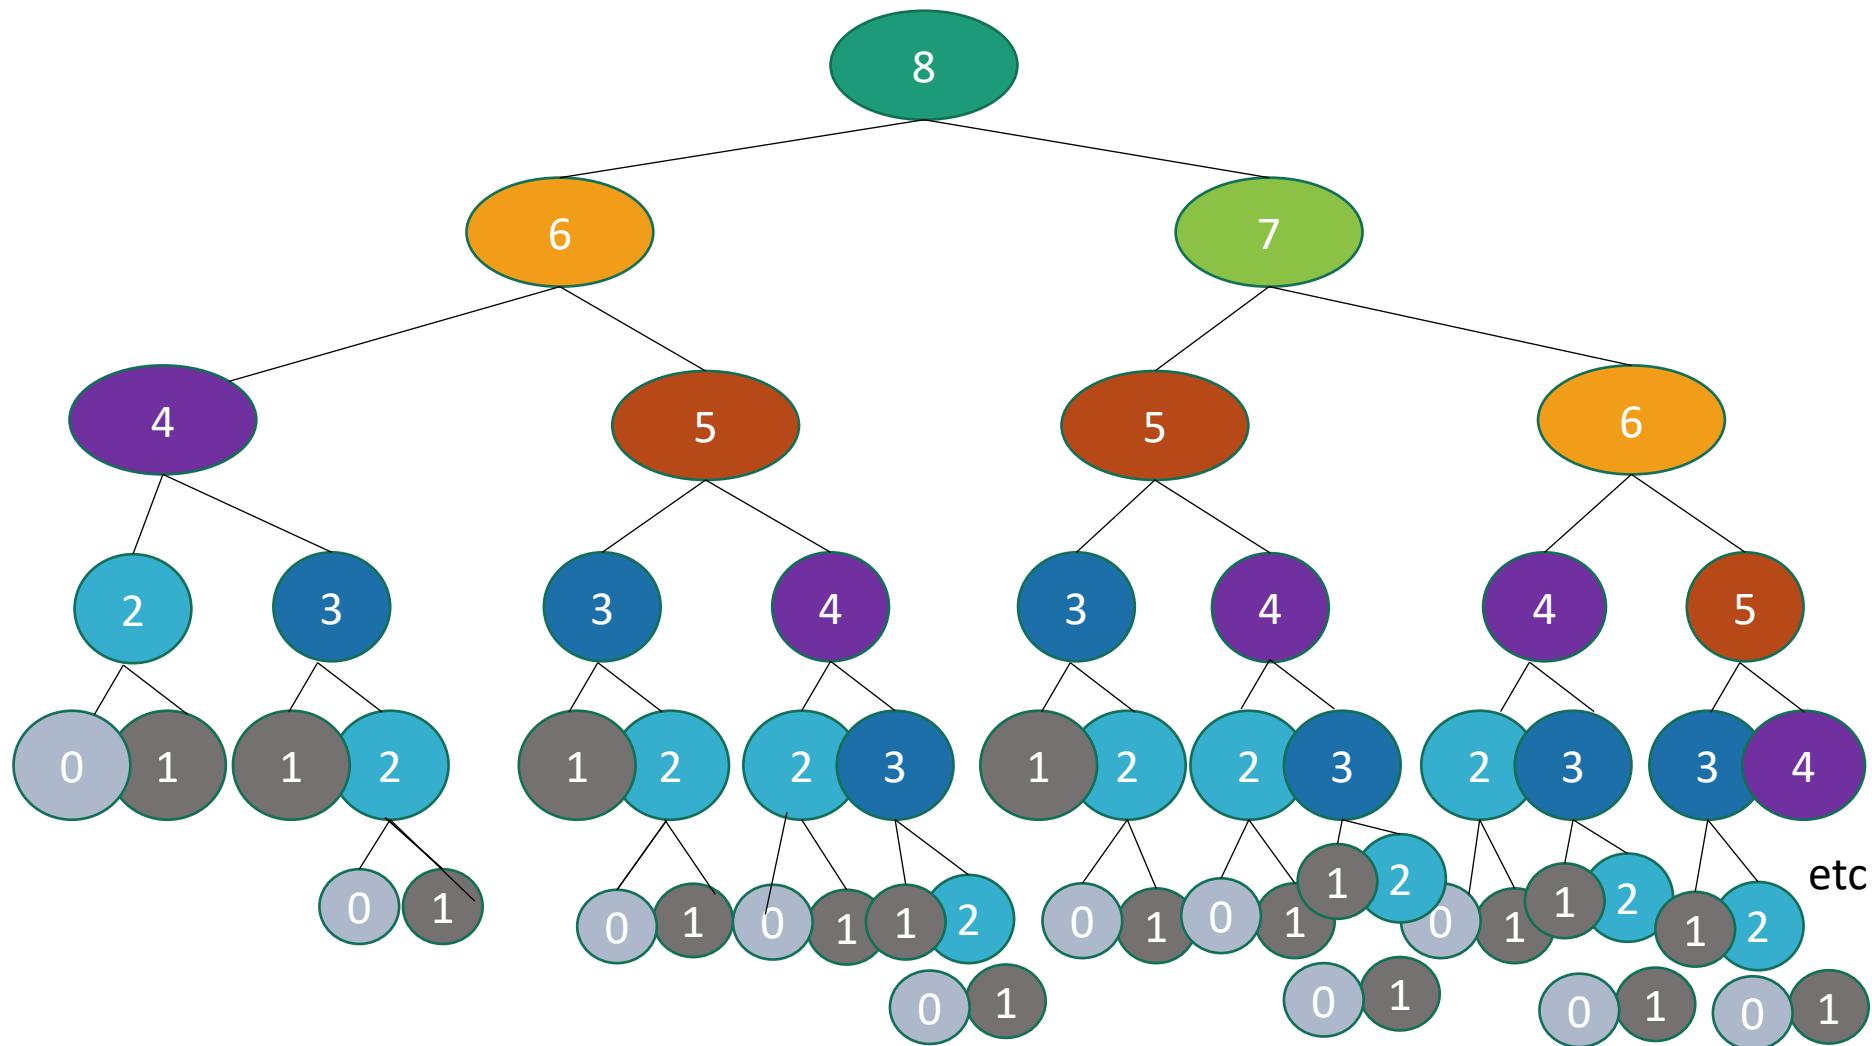
Candidate algorithm

- **def** Fibonacci(*n*):
 - **if** *n* == 0 or *n* == 1:
 - **return** 1
 - **return** Fibonacci(*n*-1) + Fibonacci(*n*-2)

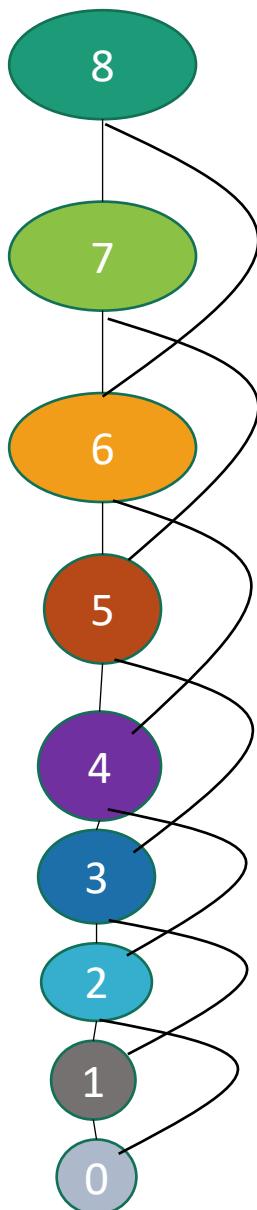


What's going on?
Consider Fib(8)

That's a lot of
repeated
computation!

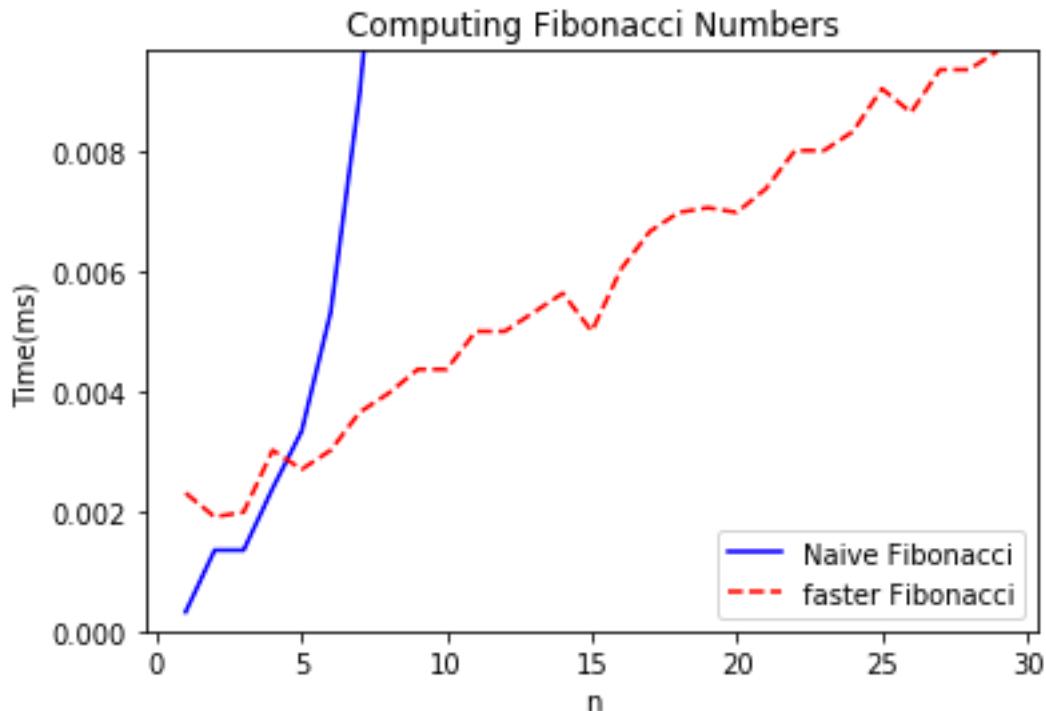


Maybe this would be better:



```
def fasterFibonacci(n):
    • F = [1, 1, None, None, ..., None ]
        • \\\ F has length n
    • for i = 2, ..., n:
        • F[i] = F[i-1] + F[i-2]
    • return F[n]
```

Much better running time!



This was an example of...

*Dynamic
Programming!*

What is *dynamic programming*?

- It is an algorithm design paradigm
 - like divide-and-conquer is an algorithm design paradigm.
- Usually it is for solving **optimization problems**
 - eg, *shortest* path
 - (Fibonacci numbers aren't an optimization problem, but they are a good example...)

Elements of dynamic programming

1. Optimal sub-structure:

- Big problems break up into sub-problems.
 - Fibonacci: $F(i)$ for $i \leq n$
 - Bellman-Ford: Shortest paths with at most i edges for $i \leq n$
- The solution to a problem can be expressed in terms of solutions to smaller sub-problems.
 - Fibonacci:

$$F(i+1) = F(i) + F(i-1)$$

- Bellman-Ford:

$$d^{(i+1)}[v] \leftarrow \min\{ d^{(i)}[v], \min_u \{d^{(i)}[u] + \text{weight}(u,v)\} \}$$

Shortest path with at
most i edges from s to v

Shortest path with at most
 i edges from s to u .

Elements of dynamic programming

2. Overlapping sub-problems:

- The sub-problems overlap a lot.
 - Fibonacci:
 - Lots of different $F[j]$ will use $F[i]$.
 - Bellman-Ford:
 - Lots of different entries of $d^{(i+1)}$ will use $d^{(i)}[v]$.
- This means that we can save time by solving a sub-problem just once and storing the answer.

Elements of dynamic programming

- Optimal substructure.
 - Optimal solutions to sub-problems are sub-solutions to the optimal solution of the original problem.
- Overlapping subproblems.
 - The subproblems show up again and again
- Using these properties, we can design a ***dynamic programming*** algorithm:
 - Keep a table of solutions to the smaller problems.
 - Use the solutions in the table to solve bigger problems.
 - At the end we can use information we collected along the way to find the solution to the whole thing.

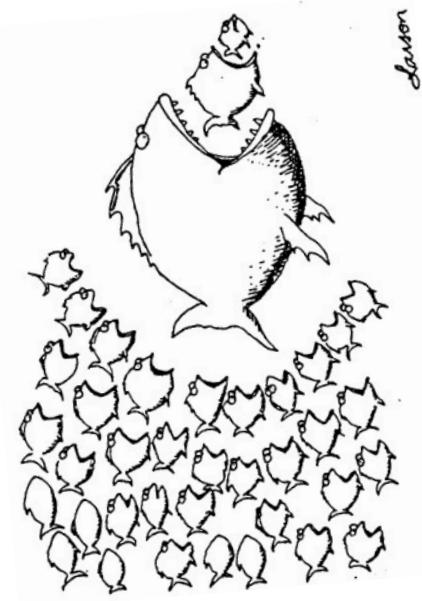
Two ways to think about and/or implement DP algorithms

- Top down
- Bottom up

Bottom up approach

what we just saw.

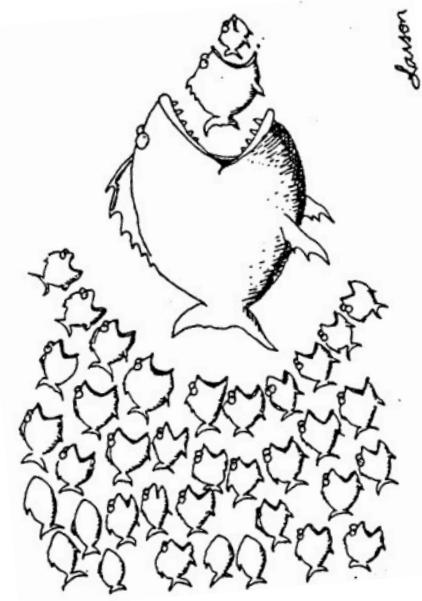
- For Fibonacci:
- Solve the small problems first
 - fill in $F[0], F[1]$
- Then bigger problems
 - fill in $F[2]$
- ...
- Then bigger problems
 - fill in $F[n-1]$
- Then finally solve the real problem.
 - fill in $F[n]$



Bottom up approach

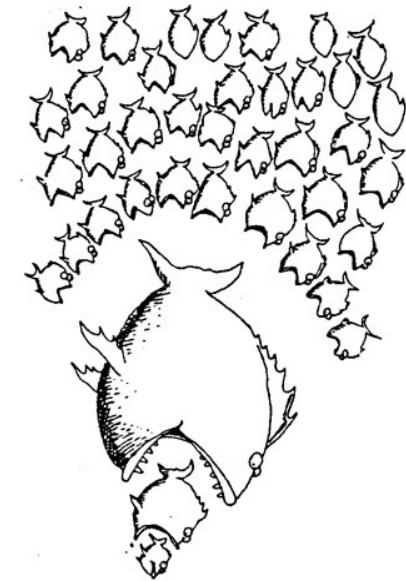
what we just saw.

- For Bellman-Ford:
- Solve the small problems first
 - fill in $d^{(0)}$
- Then bigger problems
 - fill in $d^{(1)}$
- ...
- Then bigger problems
 - fill in $d^{(n-2)}$
- Then finally solve the real problem.
 - fill in $d^{(n-1)}$



Top down approach

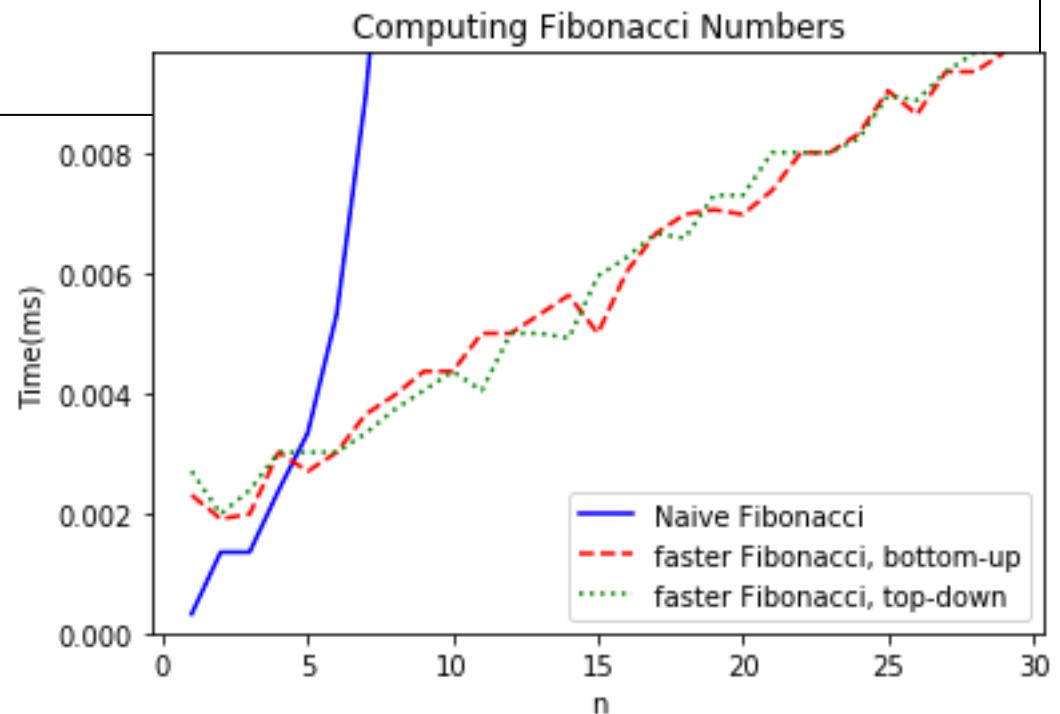
- Think of it like a recursive algorithm.
- To solve the big problem:
 - Recurse to solve smaller problems
 - Those recurse to solve smaller problems
 - etc..
- The difference from divide and conquer:
 - **Memo-ization**
 - Keep track of what small problems you've already solved to prevent re-solving the same problem twice.



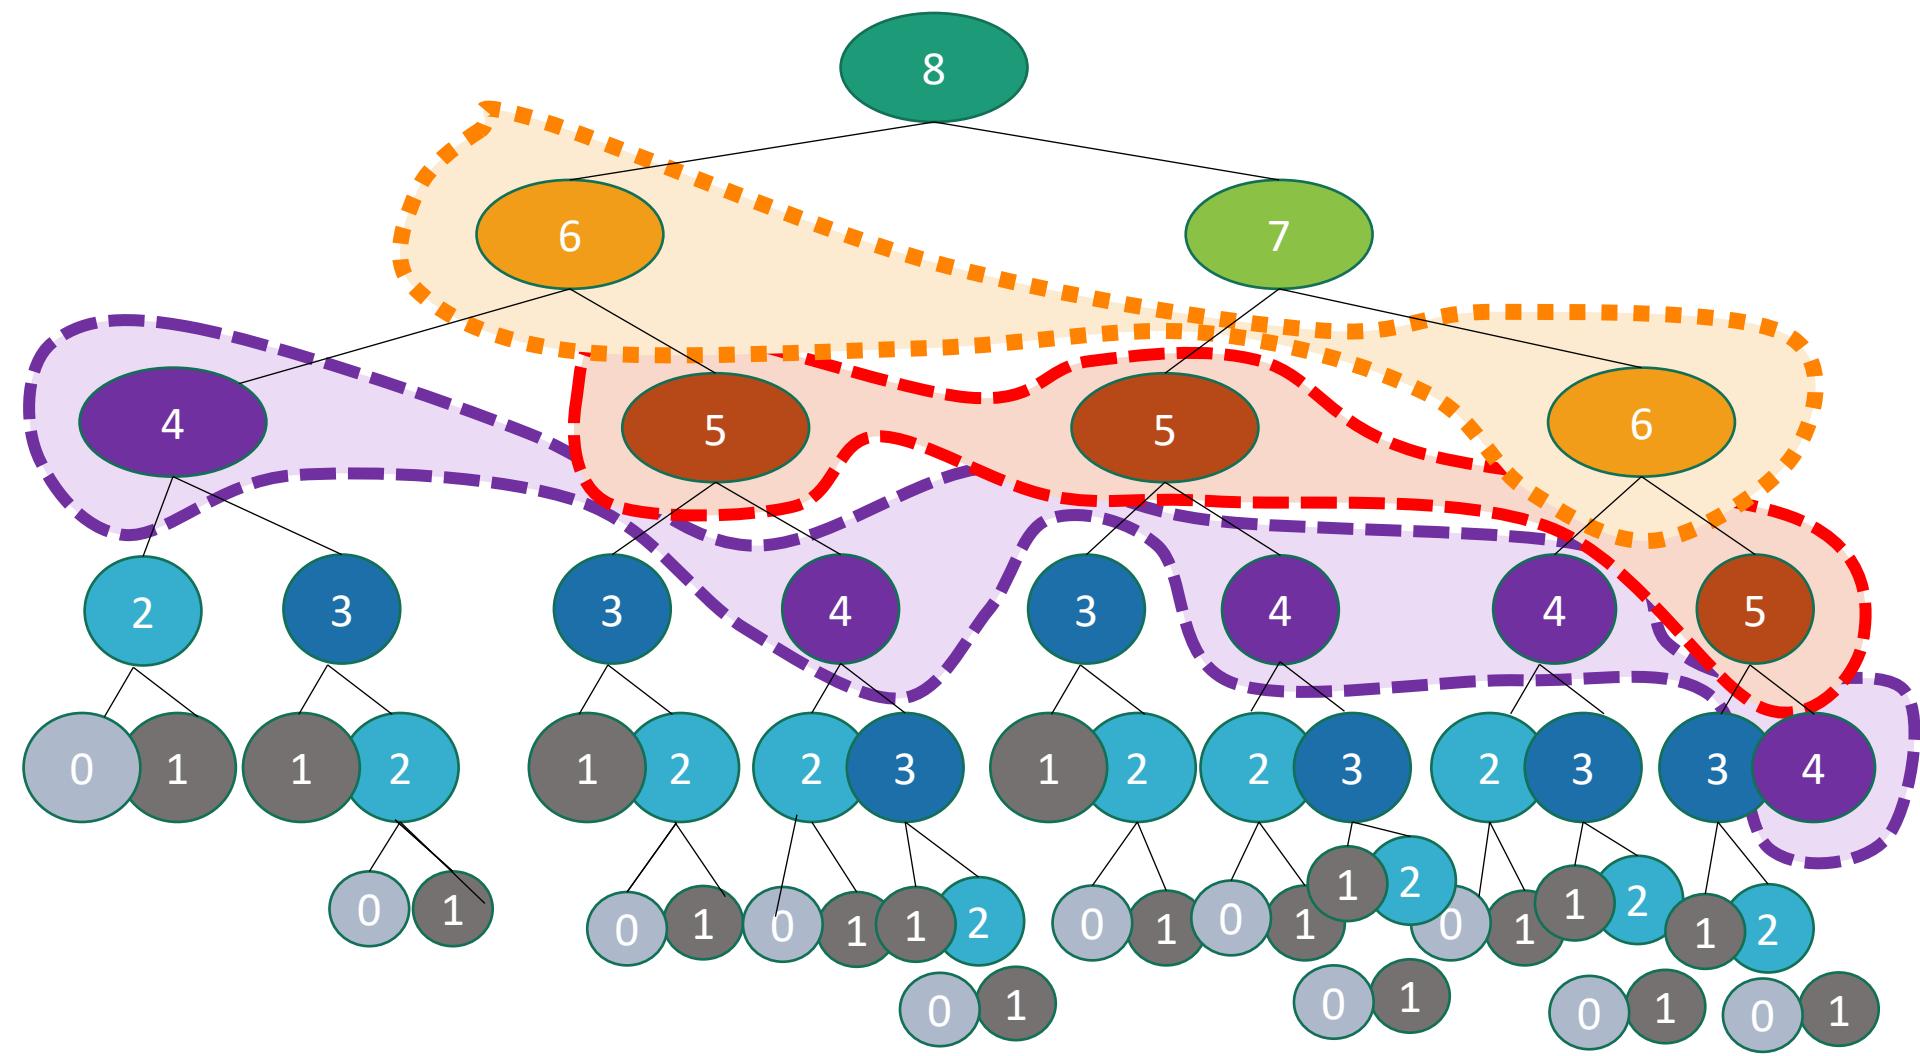
Example of top-down Fibonacci

- define a global list $F = [1, 1, \text{None}, \text{None}, \dots, \text{None}]$
- **def** Fibonacci(n):
 - **if** $F[n] \neq \text{None}$:
 - **return** $F[n]$
 - **else**:
 - $F[n] = \text{Fibonacci}(n-1) + \text{Fibonacci}(n-2)$
 - **return** $F[n]$

Memo-ization:
Keeps track (in F)
of the stuff you've
already done.



Memo-ization visualization



repe
ar
the

What have we learned?

- ***Dynamic programming:***

- Paradigm in algorithm design.
- Uses **optimal substructure**
- Uses **overlapping subproblems**
- Can be implemented **bottom-up** or **top-down**.
- It's a fancy name for a pretty common-sense idea:



Why “*dynamic programming*” ?

- Programming refers to finding the optimal “program.”
 - as in, a shortest route is a *plan* aka a *program*.
- Dynamic refers to the fact that it’s multi-stage.
- But also it’s just a fancy-sounding name.

Why “*dynamic programming*” ?

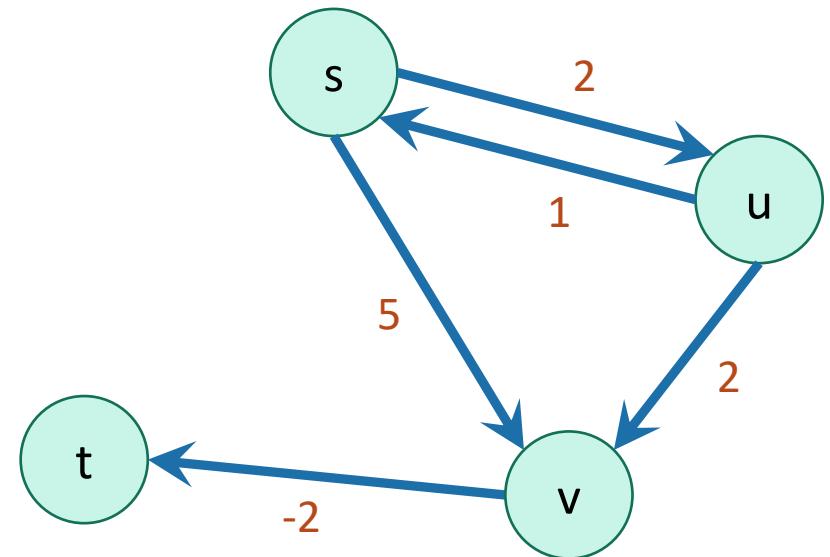
- Richard Bellman invented the name in the 1950's.
- At the time, he was working for the RAND Corporation, which was basically working for the Air Force, and government projects needed flashy names to get funded.
- From Bellman's autobiography:
 - “It's impossible to use the word, dynamic, in the pejorative sense...I thought dynamic programming was a good name. It was something not even a Congressman could object to.”

Floyd-Warshall Algorithm

Another example of DP

- This is an algorithm for **All-Pairs Shortest Paths (APSP)**
 - That is, I want to know the shortest path from u to v for **ALL pairs** u,v of vertices in the graph.
 - Not just from a special single source s .

Destination		s	u	v	t
Source	s	0	2	4	2
u	1	0	2	0	
v	∞	∞	0	-2	
t	∞	∞	∞	0	



Floyd-Warshall Algorithm

Another example of DP

- This is an algorithm for **All-Pairs Shortest Paths (APSP)**
 - That is, I want to know the shortest path from u to v for **ALL pairs** u,v of vertices in the graph.
 - Not just from a special single source s .
- Naïve solution (if we want to handle negative edge weights):
 - For all s in G :
 - Run Bellman-Ford on G starting at s .
 - Time $O(n \cdot nm) = O(n^2m)$,
 - may be as bad as n^4 if $m=n^2$

Can we do better?

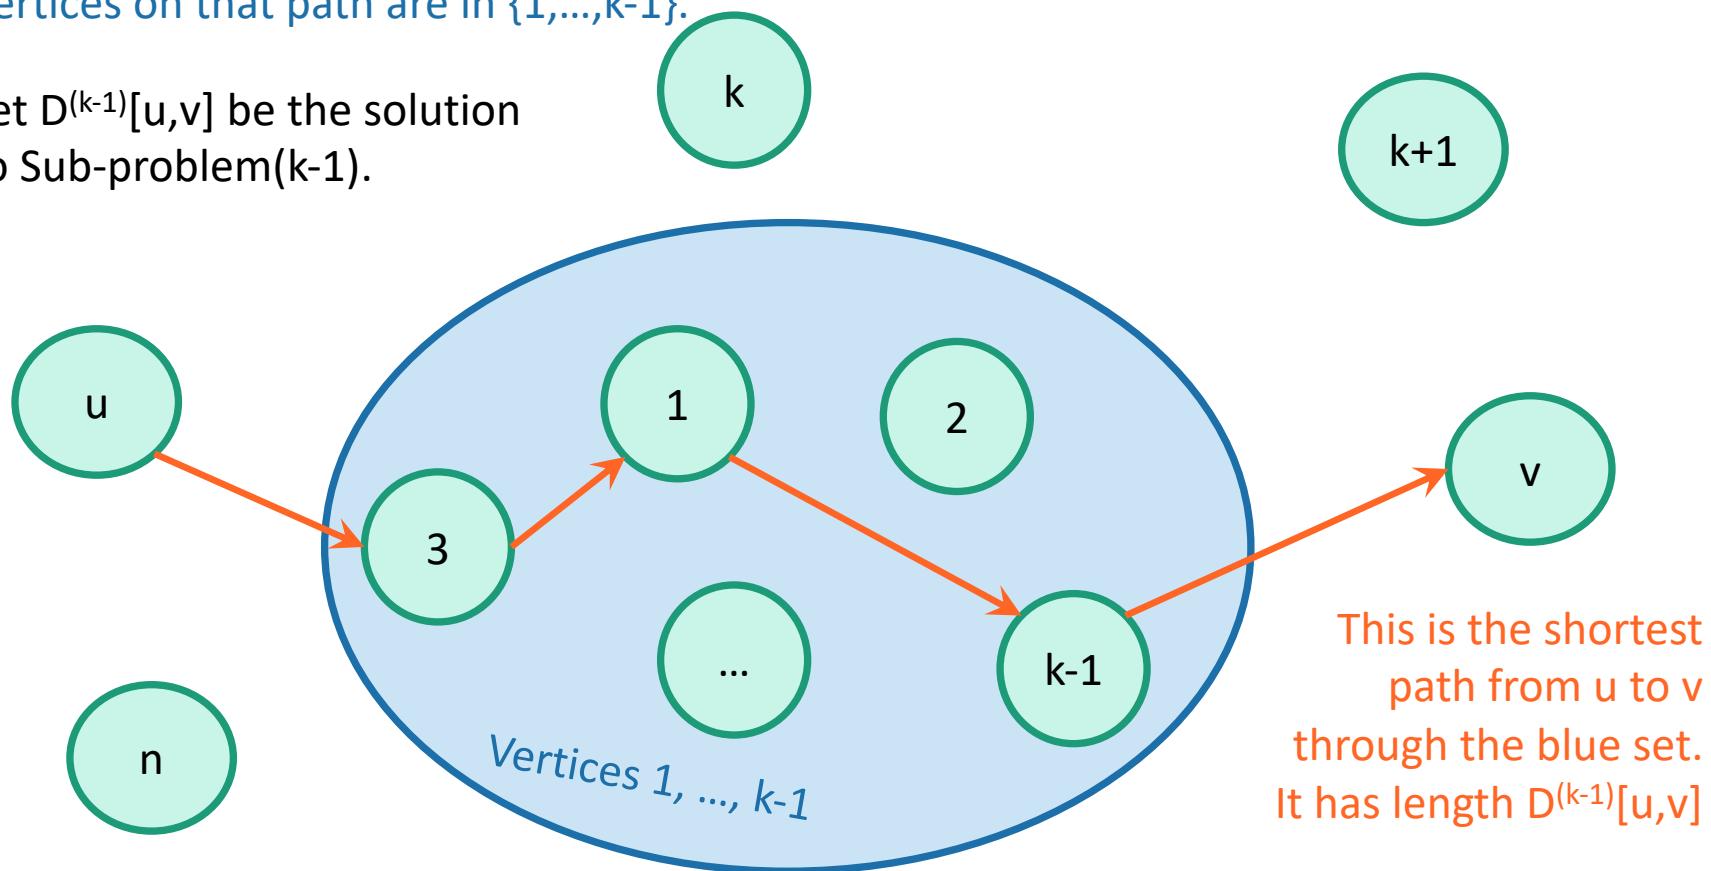
Optimal substructure

Label the vertices $1, 2, \dots, n$
(We omit some edges in the picture below).

Sub-problem(k-1):

For all pairs, u, v , find the cost of the shortest path from u to v , so that all the internal vertices on that path are in $\{1, \dots, k-1\}$.

Let $D^{(k-1)}[u, v]$ be the solution to Sub-problem(k-1).



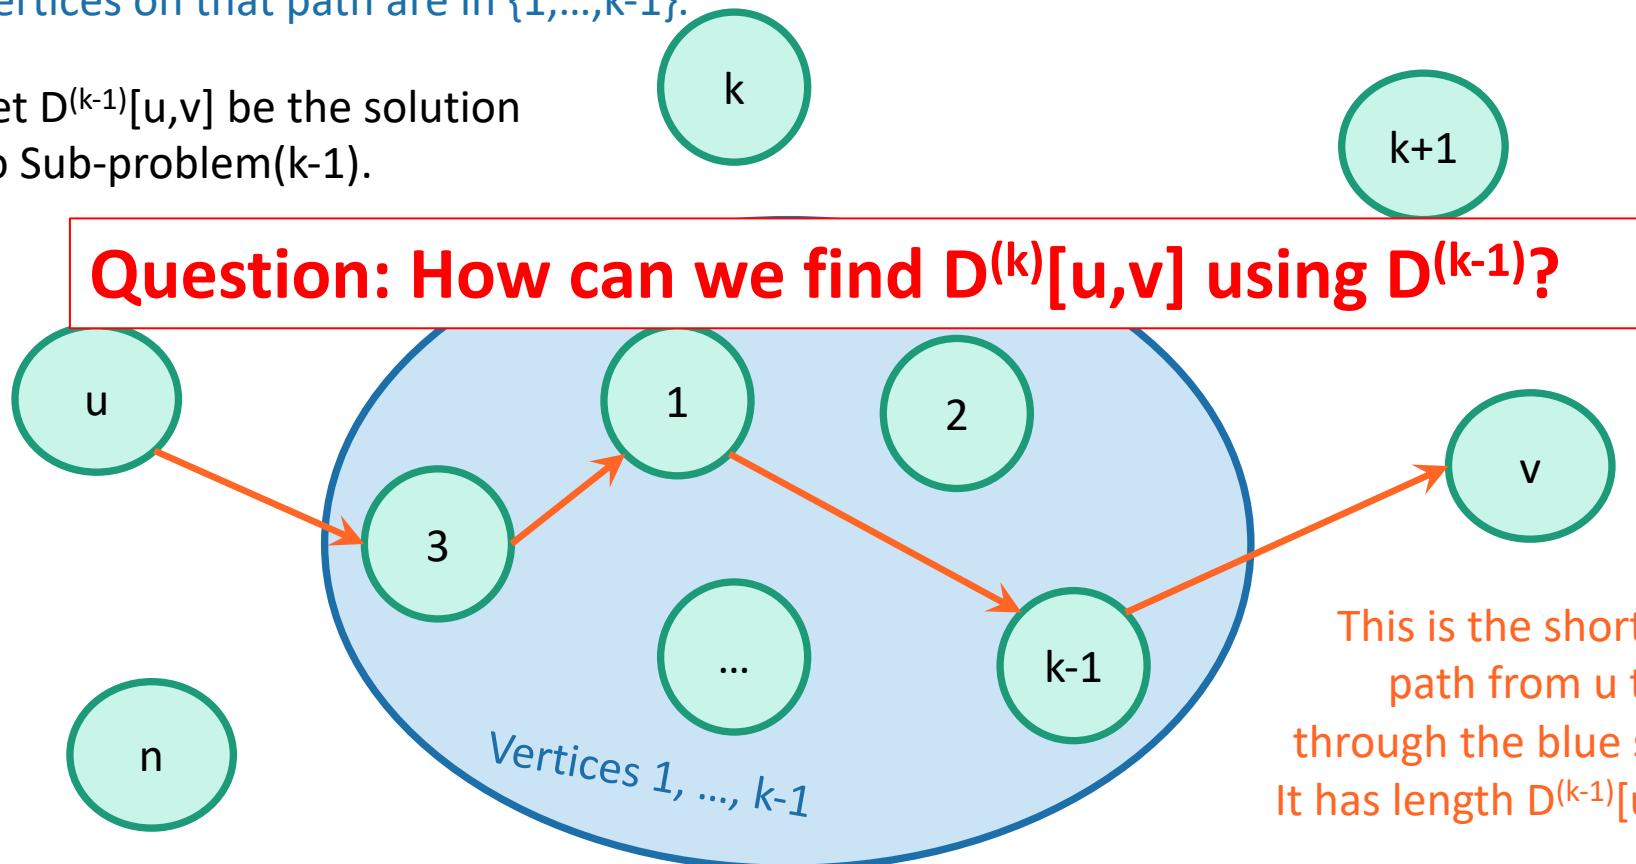
Optimal substructure

Label the vertices $1, 2, \dots, n$
(We omit some edges in the picture below).

Sub-problem($k-1$):

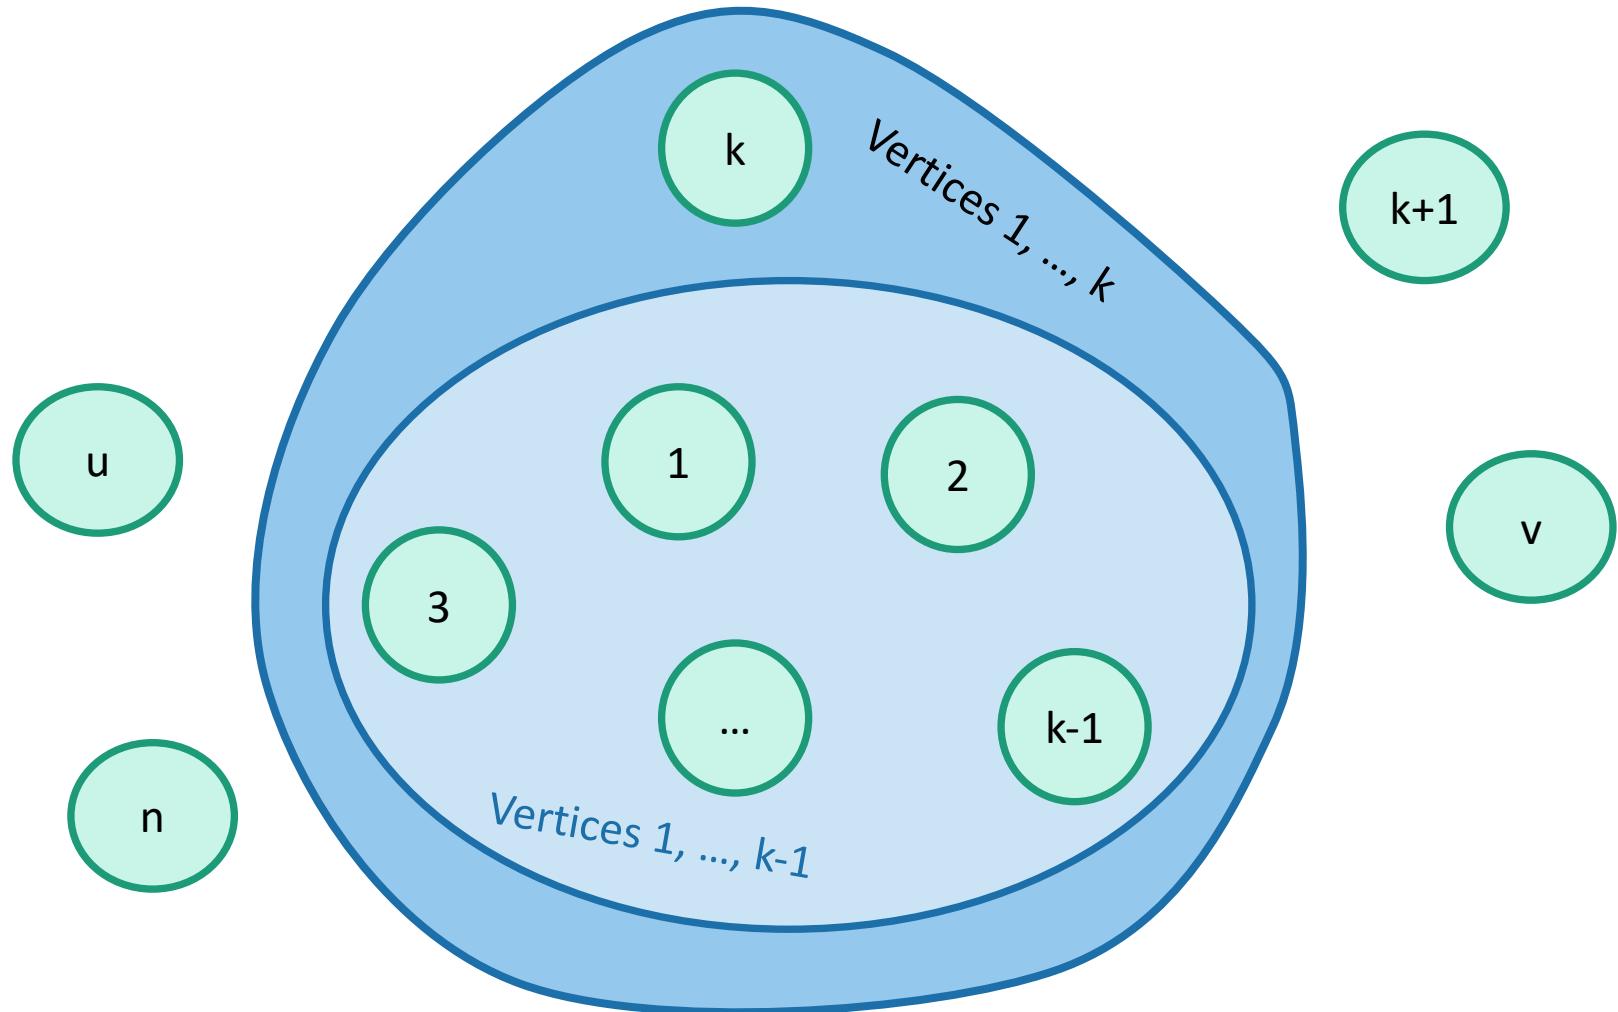
For all pairs, u, v , find the cost of the shortest path from u to v , so that all the internal vertices on that path are in $\{1, \dots, k-1\}$.

Let $D^{(k-1)}[u, v]$ be the solution to Sub-problem($k-1$).



How can we find $D^{(k)}[u,v]$ using $D^{(k-1)}$?

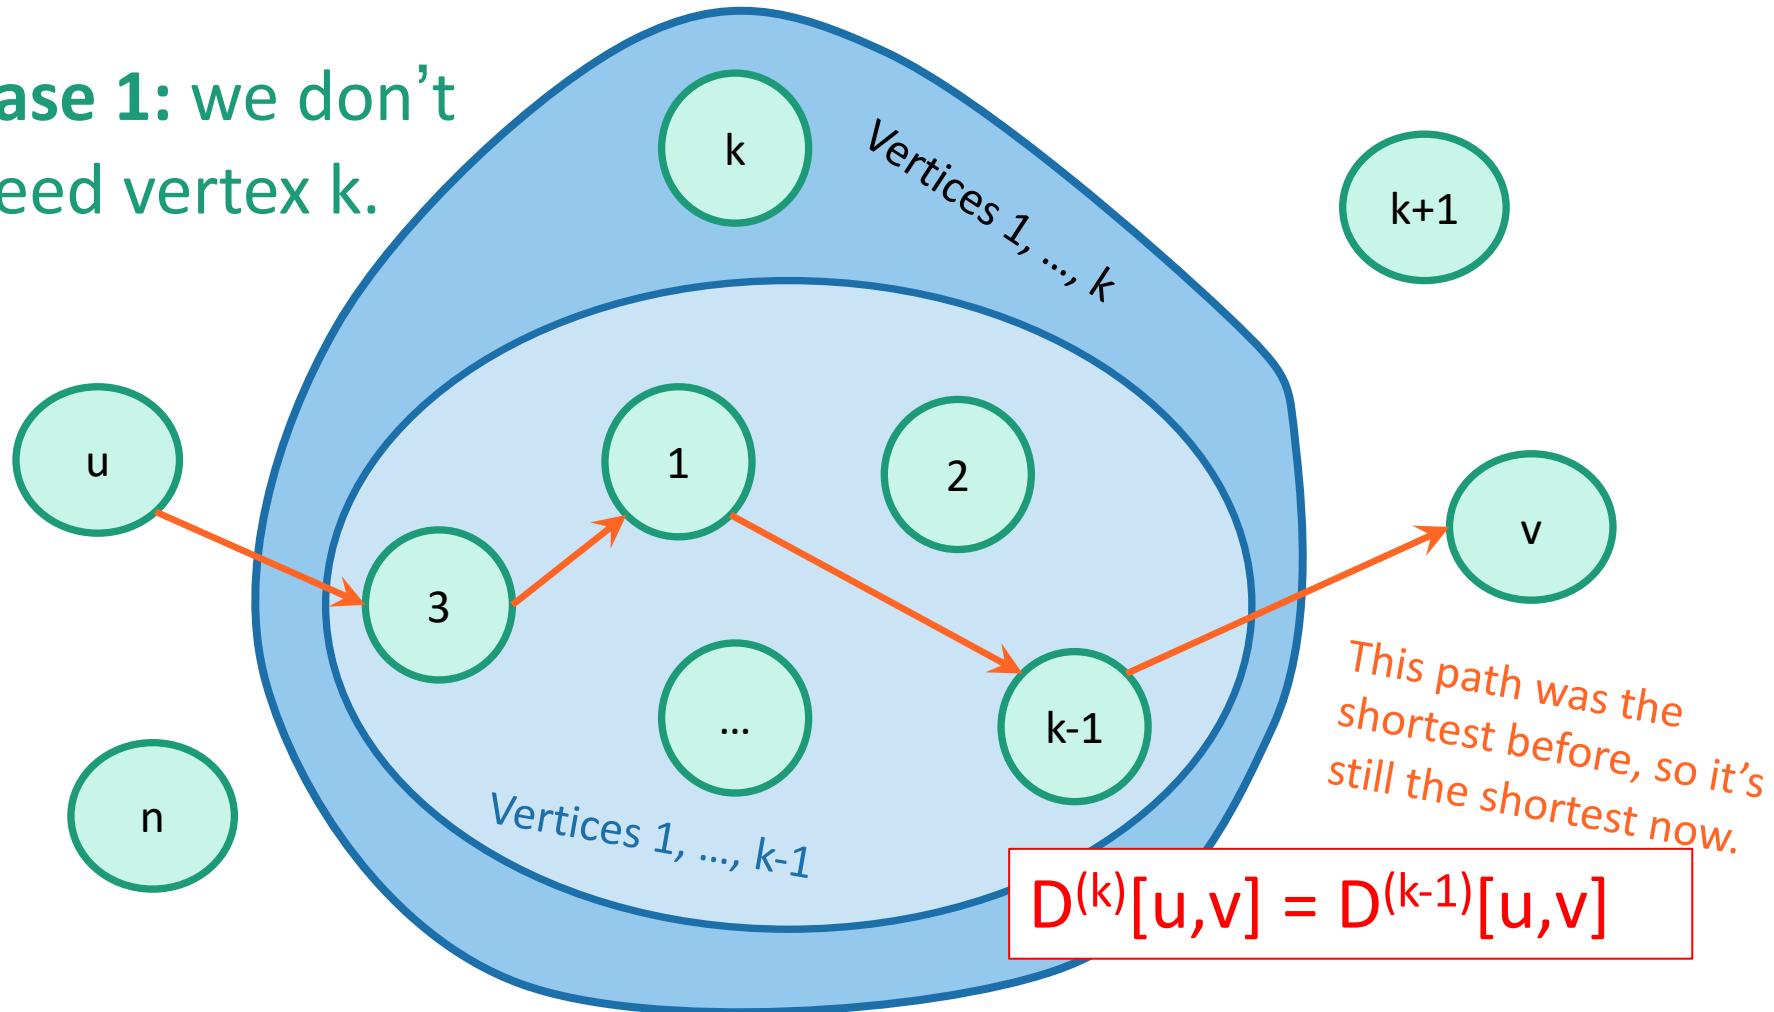
$D^{(k)}[u,v]$ is the cost of the shortest path from u to v so that all internal vertices on that path are in $\{1, \dots, k\}$.



How can we find $D^{(k)}[u,v]$ using $D^{(k-1)}$?

$D^{(k)}[u,v]$ is the cost of the shortest path from u to v so that all internal vertices on that path are in $\{1, \dots, k\}$.

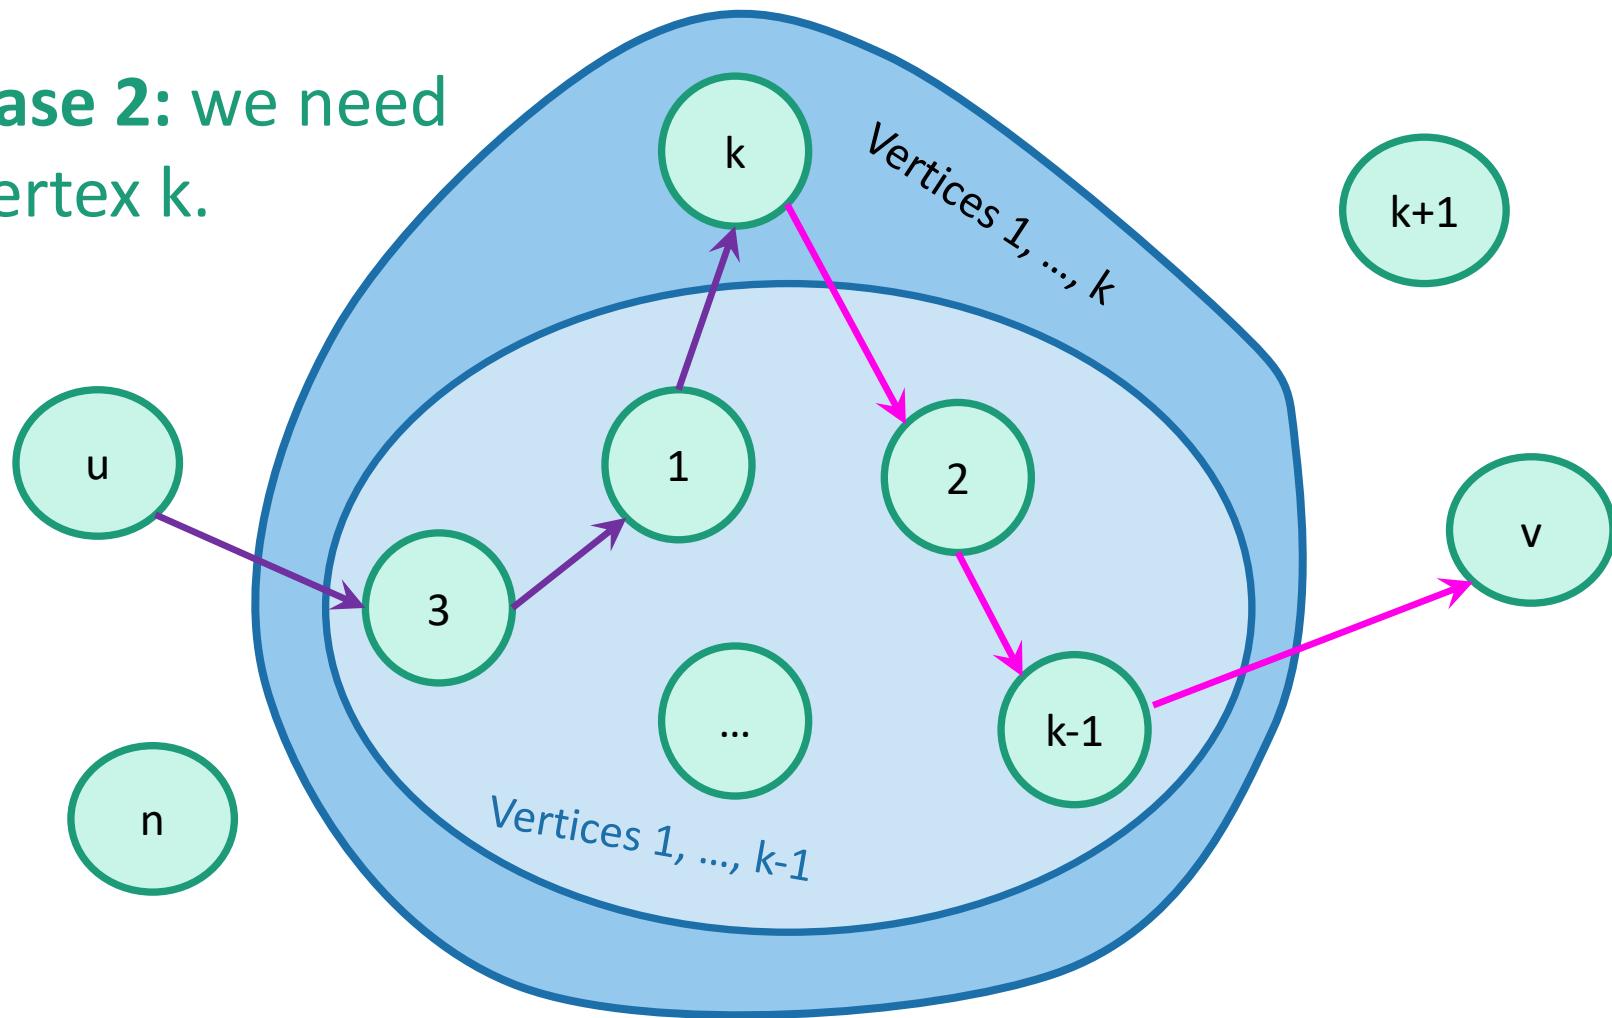
Case 1: we don't need vertex k .



How can we find $D^{(k)}[u,v]$ using $D^{(k-1)}$?

$D^{(k)}[u,v]$ is the cost of the shortest path from u to v so that all internal vertices on that path are in $\{1, \dots, k\}$.

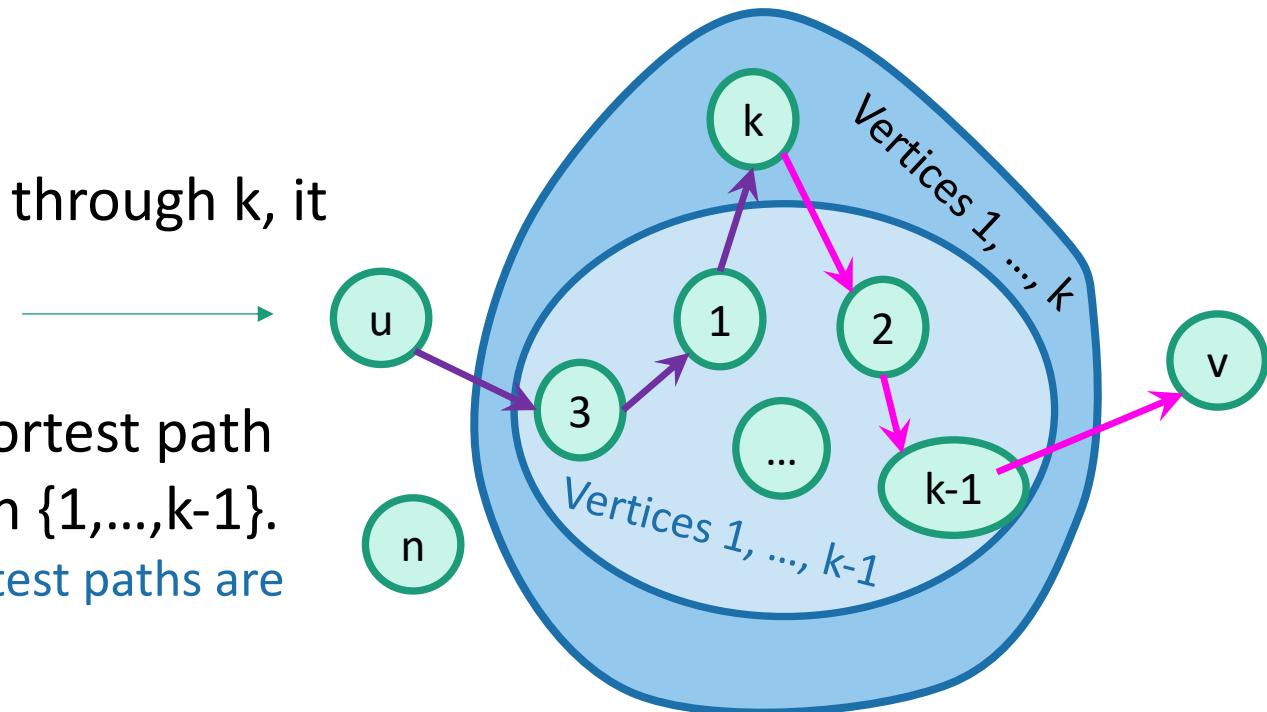
Case 2: we need vertex k .



Case 2 continued

Case 2: we need vertex k.

- If that path passes through k, it must look like this:
- This path is the shortest path from u to k through {1,...,k-1}.
 - sub-paths of shortest paths are shortest paths
- Same for this path.



$$D^{(k)}[u, v] = D^{(k-1)}[u, k] + D^{(k-1)}[k, v]$$

How can we find $D^{(k)}[u,v]$ using $D^{(k-1)}$?

- $D^{(k)}[u,v] = \min\{ D^{(k-1)}[u,v], D^{(k-1)}[u,k] + D^{(k-1)}[k,v] \}$

Case 1: Cost of
shortest path
through $\{1,\dots,k-1\}$

Case 2: Cost of shortest path
from **u to k** and then from **k to v**
through $\{1,\dots,k-1\}$

- Optimal substructure:
 - We can solve the big problem using smaller problems.
- Overlapping sub-problems:
 - $D^{(k-1)}[k,v]$ can be used to help compute $D^{(k)}[u,v]$ for lots of different u's.

How can we find $D^{(k)}[u,v]$ using $D^{(k-1)}$?

- $D^{(k)}[u,v] = \min\{ D^{(k-1)}[u,v], D^{(k-1)}[u,k] + D^{(k-1)}[k,v] \}$

Case 1: Cost of
shortest path
through $\{1,\dots,k-1\}$

Case 2: Cost of shortest path
from **u to k** and then from **k to v**
through $\{1,\dots,k-1\}$

- Using our *Dynamic programming* paradigm, this immediately gives us an algorithm!

Floyd-Warshall algorithm

- Initialize n -by- n arrays $D^{(k)}$ for $k = 0, \dots, n$
 - $D^{(k)}[u,u] = 0$ for all u , for all k
 - $D^{(k)}[u,v] = \infty$ for all $u \neq v$, for all k
 - $D^{(0)}[u,v] = \text{weight}(u,v)$ for all (u,v) in E .
- **For** $k = 1, \dots, n$:
 - **For** pairs u,v in V^2 :
 - $D^{(k)}[u,v] = \min\{ D^{(k-1)}[u,v], D^{(k-1)}[u,k] + D^{(k-1)}[k,v] \}$
- **Return** $D^{(n)}$



This is a bottom-up *Dynamic programming* algorithm.

We've basically just shown

- Theorem:
 - If there are **no negative cycles** in a weighted directed graph G , then the Floyd-Warshall algorithm, running on G , returns a matrix $D^{(n)}$ so that:
$$D^{(n)}[u,v] = \text{distance between } u \text{ and } v \text{ in } G.$$
- Running time: $O(n^3)$
 - Better than running BF n times!
 - Not really better than running Dijkstra n times.
 - But it's simpler to implement and handles negative weights.
- Storage:
 - Need to store **two** n -by- n arrays, and the original graph.
As with Bellman-Ford, we don't really need to store all n of the $D^{(k)}$.

What if there *are* negative cycles?

- Just like Bellman-Ford, Floyd-Warshall can detect negative cycles:
 - Negative cycle $\Leftrightarrow \exists v$ s.t. there is a path from v to v that goes through all n vertices that has cost < 0 .
 - Negative cycle $\Leftrightarrow \exists v$ s.t. $D^{(n)}[v,v] < 0$.
- Algorithm:
 - Run Floyd-Warshall as before.
 - If there is some v so that $D^{(n)}[v,v] < 0$:
 - **return negative cycle.**

What have we learned?

- The Floyd-Warshall algorithm is another example of *dynamic programming*.
- It computes All Pairs Shortest Paths in a directed weighted graph in time $O(n^3)$.